

LARGE-WAVE AMPLIFICATION ON PERIODIC DISCRETE LATTICES: LINEAR THEORY AND A NONLINEAR EXTENSION

ILIA GRADINA

ABSTRACT. We study the large-wave effect for free oscillations of the periodic discrete chain (the ring of N beads) and for the discrete Schrödinger equation on the same ring. Writing $A_N = \sup_t \max_j |u_j(t)|$ for the unit-impulse Green's function, we show that the supremum is, by Bohr's theorem, the ℓ^1 mass of the Fourier coefficients whenever the frequencies $\{2 \sin(\pi r/N)\}$ are linearly independent over $\mathbb{Q}$, and we prove this independence for N an odd prime and for $N = 2^m$ by a cyclotomic argument. Consequently $A_N = U_N := \frac{1}{2N} \sum_{r=1}^{N-1} \csc(\pi r/N) = \frac{1}{\pi} \ln N + \frac{1}{\pi}(\gamma + \ln \frac{2}{\pi}) + o(1)$ unconditionally on these classes, and the same sharp constant holds for the composite family $N = 2p$. For every N the upper bound $A_N \leq U_N = \frac{1}{\pi} \ln N + O(1)$ is rigorous; the orbit-closure subtorus has dimension exactly $\frac{1}{2}\varphi(2N)$, and $A_N = U_N$ holds *if and only if* every integer relation among the frequencies has coefficient-sum divisible by four (with an antipodal variant for even N)—an arithmetic criterion met *exactly* by the primes and the powers of two (Theorem 16), and by no composite N . Moreover the *order* is now settled for *all* N : $A_N \sim \frac{1}{\pi} \ln N$ unconditionally (the lower bound from a palindromic argument showing the first $\frac{1}{2}\varphi(2N)$ frequencies are $\mathbb{Q}$ -independent, so a $\ln N$ fraction of the ceiling can always be aligned). The composite deficit $U_N - A_N$ is $O(\ln \ln \ln N)$, and the order law extends to the Dirichlet segment and to every ring product $C_N \square H$. For the discrete Schrödinger propagator e^{-itL_N} , unitarity forbids amplification of a localized state, while the $\ell^\infty \rightarrow \ell^\infty$ amplification satisfies $B_N = \Theta(\sqrt{N})$; here $\liminf B_N/\sqrt{N} \geq c_0/\sqrt{2}$ is proved and the constant splits by parity: numerically the even subsequence approaches $c_0/\sqrt{2}$ while the odd one approaches $\beta_{\text{odd}} = 0.928 \dots$ (a Bessel-envelope integral for which no elementary closed form was found), with the rigorous bound $\limsup \geq \beta_{\text{odd}}$. For the focusing discrete NLS the linear large wave persists under weak nonlinearity, the modulational-instability band is read off the L_N spectrum, and the on-site breather exists and is stable. Proved results are separated throughout from computational evidence and conjectures; all quantitative claims are reproduced by the accompanying scripts, with a GAP/PARI/FRICAS/ROCQ exact-arithmetic layer.

1. INTRODUCTION

A localized disturbance on a homogeneous one-dimensional medium does not amplify: for the continuous wave equation the classical bound of Zhukovsky gives $|\sigma(x, t)| \leq F$ [49]. On the *discrete* chain this is false [42, 41, 47]. The discreteness of the spectrum, together with the finiteness of the number of modes, allows many modes to come into phase and to produce a local peak that grows with the number of nodes N . This is the *large-wave effect*; related phenomena include the splash effect [46] and its Schrödinger analogue [48].

The effect is genuinely discrete: it disappears in the continuum limit. The mechanism is the *finite, discrete spectrum*—finitely many discrete modes can be brought into near-coincidence of phase (arbitrarily closely when their frequencies are rationally *independent*, by Kronecker, and otherwise as far as the relation lattice allows); on the infinite chain the spectrum is continuous, no exact alignment occurs, and the amplification is absent. It is thus a finite-size spectral resonance rather than a feature of the limiting partial differential equation. (For the discrete Schrödinger propagator the absence of a continuum large wave is explicit: its single-site kernel is the Bessel function $J_n(2t)$, whose peak over sites decays like

$t^{-1/3}$, Section 11.) This discrete–continuous gap, together with the appearance of discrete Laplacians in quantum walks, photonic lattices, non-Hermitian lattice optics [39], and discrete nonlinear models of anomalous (rogue) waves, motivates studying the lattice in its own right.

Contributions. (i) We identify $\sup_t G_N(0, t) = U_N$ through Bohr’s theorem under rational independence of the frequencies (Section 3), and prove this independence for N prime and $N = 2^m$ by a cyclotomic argument (Theorems 3, 4), giving the unconditional asymptotics $A_N \sim \frac{1}{\pi} \ln N$ with the explicit constant of Lemma 2; the same sharp constant extends to the composite family $N = 2p$ (Proposition 17). (ii) For every N the upper bound $A_N \leq U_N$ is rigorous, and we reduce A_N *exactly* to an optimization over the orbit-closure subtorus, whose dimension we compute in closed form as $\frac{1}{2}\varphi(2N)$ (Theorem 18); a sharp arithmetic criterion—coefficient-sums divisible by four (plus an antipodal variant for even N)—decides exactly when $A_N = U_N$ (Theorem 13), and an explicit root-of-unity obstruction settles the saturation locus completely: $A_N = U_N$ *iff* N is prime or a power of two (Theorem 16). Separately, the *order* is settled for *all* N : $A_N \sim \frac{1}{\pi} \ln N$ (Theorem 6), the lower bound coming from a long alignable prefix ($M(N) = \frac{1}{2}\varphi(2N)$ frequencies are $\mathbb{Q}$ -independent, Lemma 5). (iii) For composite N the value is algebraic with increasing arithmetic complexity (Proposition 12); we further separate the infinite-time ceiling from the finite-time amplitude and show both the finite-horizon recurrence time and the effect of mass disorder (Section 10). (iv) For the discrete Schrödinger equation a localized state is never amplified (unitarity, rigorous), and the $\ell^\infty \rightarrow \ell^\infty$ norm satisfies $B_N = \Theta(\sqrt{N})$, both bounds rigorous (Theorems 20, 22). (v) As a nonlinear outlook we connect L_N to the rogue-wave problem (Section 18): the discrete NLS $i\dot{z} = -L_N z + \gamma|z|^2 z$ has L_N as its linear part and the focusing NLS as its continuum limit, and we verify the Peregrine rogue solution, the modulational-instability gain, and the emergence of heavy-tailed amplitude statistics in sharp contrast with the platykurtic linear value law. Every quantitative claim is reproduced by a standalone, version-controlled script; the arithmetic of Theorems 18 and 13 is additionally checked in exact cyclotomic arithmetic by two computer-algebra systems and machine-checked by the Rocq proof assistant (Section 14).

Relation to prior work. The large-wave *phenomenon*—that on the discrete bead string a localized disturbance amplifies, growing without bound at bounded energy on suitable N —is classical, due to the Kurchanov–Myshkis–Filimonov line of work [42, 41, 47], with the $\ln$ -type asymptotics on the Dirichlet segment first in [45] and sharpened in [41], and the Schrödinger analogue in [48]. The linear independence over $\mathbb{Q}$ of the frequencies for N prime or $N = 2^m$ is itself due to [42]; our Theorems 3, 4 reprove it cyclotomically, and Theorem 18 subsumes it in a closed-form $\mathbb{Q}$ -rank valid for *every* N . The lineage also includes two early C. R. Acad. Sci. notes [43, 44], of which Filimonov’s 1992 note [44] is the closest precedent: it is itself on the *circular* string, and by its published summary (Zbl 0761.73053) it already establishes that for N prime or $N = 2^m$ the *ring* vibrations exhibit peaks, addresses the degenerate “multiple-frequency” (composite) case that our mod 4 criterion governs, and proposes a continuum model finer than the wave equation. A recent monograph [38] reproduces its splash table (its Table 10.1, with $P_N \rightarrow \infty$ as $N \rightarrow \infty$), which grows at the *segment* constant $4/\pi^2$ and quantifies the splash *magnitude*—distinct from our *ring* law $\frac{1}{\pi} \ln N$ and from the arithmetic of *which* N saturate. *It is therefore possible that the qualitative ring effect and part of the composite analysis below are already Filimonov’s.* We have consulted only its abstract and this secondary reproduction (the full note is not digitized in the sources available to us); [44] must be obtained and compared against Theorems 3–4 and 13 before submission. The structural and quantitative claims below are accordingly stated as novel *to our knowledge*. Our contribution is *quantitative and structural*: the order theorem $A_N \sim \frac{1}{\pi} \ln N$ for *every* N , via the palindromic prefix-independence lemma (Theorem 6, Lemma 5); the saturation classification $A_N = U_N \Leftrightarrow N$ prime or 2^m (Theorem 16); the explicit ring constant $\frac{1}{\pi}$ and the next term $-\frac{\pi}{72}N^{-2}$

(Lemma 2); the closed-form subtorus dimension $\frac{1}{2}\varphi(2N)$ for every N , with a cyclotomic reproof of the prime/ 2^m independence of [42] as its full-rank case (Theorems 18, 3, 4); the exact mod 4 criterion for reaching the ceiling (Theorem 13); the new composite family $N = 2p$ (Proposition 17); the spectral-zeta systematization; the finite-time/disorder dichotomy (Section 10); the sharp $B_N = \Theta(\sqrt{N})$ growth of the discrete Schrödinger propagator (Theorem 22); the segment rank and classification (Theorem 10); and the order criterion with its ring-product corollary (Proposition 33, Corollary 35). We claim neither the large-wave phenomenon nor its appearance on the ring for N prime or 2^m as new (the latter is already in [44]; see Remark 14). The same transient overshoot—coupling forces exceeding the applied load by a spatial redistribution of energy absent in homogeneous media—is independently recognized in the elastic-metamaterials and homogenization literature, where it is traced to the Filimonov line and treated by two-point Padé and gradient-elasticity continualization [35, 36, 37, 38]; our contribution is orthogonal, namely the arithmetic that governs *which* N and *how large*.

Main results at a glance.

- T1** $A_N = U_N = \frac{1}{\pi} \ln N + \frac{1}{\pi}(\gamma + \ln \frac{2}{\pi}) - \frac{\pi}{72N^2} + \dots$ for N prime and $N = 2^m$ (independence); $A_N = U_N - O(1/N)$ with the same constant $\frac{1}{\pi}$ for $N = 2p$ (Thms. 3, 4, Prop. 17).
- T2** The orbit-closure subtorus has dimension $\frac{1}{2}\varphi(2N)$ for *every* N (Thm. 18).
- T3** *Ceiling criterion and classification:* $A_N = U_N$ iff every integer frequency relation has coefficient-sum $\equiv 0 \pmod{4}$ (or, N even, an antipodal variant; Thm. 13); this holds *exactly* for N prime or $N = 2^m$, so $\{N : A_N = U_N\} = \{\text{prime}\} \cup \{2^m\}$ (Thm. 16, via an explicit root-of-unity obstruction for every composite N).
- T4** Discrete Schrödinger: localized states never amplify (unitarity), yet $B_N = \Theta(\sqrt{N})$ in ℓ^∞ (Thms. 20, 22).
- T5** *Order, all N :* $A_N \sim \frac{1}{\pi} \ln N$ for *every* N (Thm. 6)—the lower bound from an alignable independent prefix of length $\frac{1}{2}\varphi(2N)$ (Lemma 5, by a palindromic argument).
- T6** The incoherent *ceiling* has critical spectral dimension $d_s = 1$ (incl. quasi-1D); the amplitude is an infinite-time, Diophantine-limited quantity (§17, §10).
- T7** Focusing DNLS: the linear large wave persists under weak nonlinearity, the modulational-instability band is read off the L_N spectrum, and the on-site breather exists and is stable (Props. 39, 40, 43, 44).

We treat the periodic chain (the ring) with *free* oscillations, and contrast it with the discrete Schrödinger dynamics on the same ring. The central quantity is

$$A_N = \sup_{t \geq 0} \max_{0 \leq j < N} |u_j(t)|,$$

where $u_j(t)$ is the response to a unit velocity impulse. Our tools are the theory of almost-periodic functions (Bohr), equidistribution (Kronecker–Weyl), and the arithmetic of cyclotomic fields. The main message is that A_N grows like $\ln N$, that the growth constant is exactly $1/\pi$ on an explicit family of N , and that the obstruction on the remaining N is a purely number-theoretic one (rational relations among the frequencies).

Part 1. Linear theory

2. TOEPLITZ AND CIRCULANT MATRICES: A UNIFIED LANGUAGE

The bead chain is a spring–mass lattice whose stiffness operator is shift invariant—hence Toeplitz—and, under periodic boundary conditions, circulant (Figure 1). We collect the structures and the classical algorithms that organize the entire analysis.




FIGURE 1. The two geometries. (a) The ring of N beads coupled by identical springs gives a *circulant* Laplacian $L_N = \text{circ}(2, -1, 0, \dots, 0, -1)$. (b) The Dirichlet segment gives the *tridiagonal Toeplitz* $T_N = \text{tridiag}(-1, 2, -1)$. Both have the same symbol $f(\theta) = 4 \sin^2(\theta/2)$.

2.1. Toeplitz, Hankel, circulant. A matrix is *Toeplitz* if it is constant along diagonals, $A_{ij} = a_{i-j}$ ($2n - 1$ scalars, $O(n)$ storage); *Hankel* if constant along anti-diagonals, $H_{ij} = h_{i+j}$; and *circulant* if Toeplitz with a periodic wrap, $C_{ij} = c_{(i-j) \bmod n}$. For $n = 4$,

$$T = \begin{pmatrix} a_0 & a_{-1} & a_{-2} & a_{-3} \\ a_1 & a_0 & a_{-1} & a_{-2} \\ a_2 & a_1 & a_0 & a_{-1} \\ a_3 & a_2 & a_1 & a_0 \end{pmatrix}, \quad H = \begin{pmatrix} h_0 & h_1 & h_2 & h_3 \\ h_1 & h_2 & h_3 & h_4 \\ h_2 & h_3 & h_4 & h_5 \\ h_3 & h_4 & h_5 & h_6 \end{pmatrix}, \quad C = \begin{pmatrix} c_0 & c_3 & c_2 & c_1 \\ c_1 & c_0 & c_3 & c_2 \\ c_2 & c_1 & c_0 & c_3 \\ c_3 & c_2 & c_1 & c_0 \end{pmatrix}.$$

Reversing rows turns Toeplitz into Hankel, $H = JT$ with J the exchange matrix; a Toeplitz matrix is the finite compression of a convolution (Laurent) operator. The ring Laplacian $L_N = \text{circ}(2, -1, 0, \dots, 0, -1)$ is the periodic case of the segment Toeplitz $T_N = \text{tridiag}(-1, 2, -1)$.

2.2. The symbol and circulant diagonalization. A banded Toeplitz/circulant matrix with entries a_k has symbol $f(\theta) = \sum_k a_k e^{ik\theta}$; for both L_N and T_N ,

$$f(\theta) = 2 - 2 \cos \theta = 4 \sin^2(\theta/2).$$

Every circulant is a polynomial $C = \sum_k c_k P^k$ in the cyclic shift P , and the Fourier modes $\varphi_j^{(r)} = N^{-1/2} e^{2\pi i r j / N}$ are common eigenvectors of P ($P \varphi_j^{(r)} = e^{2\pi i r / N} \varphi_j^{(r)}$). Hence

$$L_N = F_N^* \Lambda F_N, \quad \lambda_r = f\left(\frac{2\pi r}{N}\right) = 4 \sin^2 \frac{\pi r}{N}, \quad \omega_r = \sqrt{\lambda_r} = 2 \sin \frac{\pi r}{N}, \quad (1)$$

$r = 0, \dots, N - 1$. The discrete dispersion relation $\omega(\theta) = 2|\sin(\theta/2)|$ is literally the square root of the symbol, sampled at $\theta = 2\pi r/N$ on the ring and at $\theta = \pi k/(N+1)$ on the segment ($\lambda_k^{\text{Dir}} = 4 \sin^2 \frac{\pi k}{2(N+1)}$). The twofold degeneracy $\omega_r = \omega_{N-r}$ on the ring—absent on the segment—is the source of the arithmetic difficulty in Section 8.

2.3. Classical Toeplitz algorithms. Determined by $2n - 1$ scalars, a Toeplitz system admits operations far below the dense $O(n^3)$ (Table 1); three enter the analysis.

Levinson–Durbin. For symmetric Toeplitz $R = (r_{|i-j|})$ the order- i system $R^{(i)} a^{(i)} = (r_1, \dots, r_i)^\top$ is built recursively through the *reflection coefficients* k_i (PARCOR): with $E_0 = r_0$,

$$k_i = \frac{r_i - \sum_{j=1}^{i-1} a_j^{(i-1)} r_{i-j}}{E_{i-1}}, \quad a_j^{(i)} = a_j^{(i-1)} - k_i a_{i-j}^{(i-1)}, \quad a_i^{(i)} = k_i, \quad E_i = (1 - k_i^2) E_{i-1}. \quad (2)$$

Task	Algorithm	Cost
Solve $Tx = b$ (symmetric)	Levinson (1947)	$O(n^2)$
AR / Yule–Walker	Durbin (1960)	$O(n^2)$
Inverse T^{-1} (explicit)	Trench (1964) / Gohberg–Semencul (1972)	$O(n^2)$
Multiply Tx , circulant solve, $f(L_N)$	FFT / circulant embedding	$O(n \log n)$
Semi-infinite chain	Wiener–Hopf factorization	—

TABLE 1. Structured solvers used here; all beat the dense $O(n^3)$.

Recursion (2) is due to Levinson [27] and, in the AR (Yule–Walker) case, Durbin [28]; one has $|k_i| < 1$ for all i iff $R \succ 0$, making $\{k_i\}$ the natural stable-filter parametrization. The implementation used here goes back to the author’s 2012 study [40].

Trench / Gohberg–Semencul. The inverse of a Toeplitz matrix is *not* Toeplitz but is determined by its first and last columns u, v [29, 30] (two $O(n^2)$ recursions); writing L_x, U_x for the lower/upper triangular Toeplitz matrices generated by a vector x ,

$$T^{-1} = \frac{1}{u_0} (L_u U_v - L_{Sv} U_{Su}), \quad (3)$$

a difference of products of triangular Toeplitz factors (S the shift), with $\det T$ obtained en route. On the ring this collapses to the pseudoinverse $L_N^+ = F_N^* \Lambda^+ F_N$ (the constant mode, $L_N \mathbf{1} = 0$, is mapped to 0) and, more generally, $f(L_N) = F_N^* f(\Lambda) F_N$ in $O(N \log N)$ —the route by which the propagators below are evaluated.

Wiener–Hopf. On the semi-infinite chain the stiffness is a one-sided Toeplitz operator; its inversion is the Wiener–Hopf factorization $f(\theta) = f_+(\theta)f_-(\theta)$ of the symbol, which links the finite- N ceilings to the operator-theoretic (Szegő) asymptotics of large truncated Toeplitz matrices [32], Section 5.

2.4. Three dynamics and the central functional. The same lattice carries the wave equation $\ddot{u} + L_N u = 0$ (propagator $\cos(t\sqrt{L_N})$), the discrete Schrödinger equation $i\dot{z} = L_N z$ (propagator e^{-itL_N} , the phase governed by λ_r rather than $\omega_r = \sqrt{\lambda_r}$), and the dissipative heat flow $\dot{v} = -L_N v$ (no large wave—the semigroup smooths). For the wave equation with the unit zero-mean velocity impulse $u(0) = 0$, $\dot{u}(0) = e_0 - \frac{1}{N}\mathbf{1}$, the response is the velocity Green’s function

$$u_j(t) = G_N(j, t) = \frac{1}{N} \sum_{r=1}^{N-1} \frac{\sin(\omega_r t)}{\omega_r} e^{2\pi i r j / N}. \quad (4)$$

As u is real and $\omega_r = \omega_{N-r}$, (4) is a real almost-periodic function of t with frequency set $\{\omega_r\}$, and the large wave is the operator-norm growth $A_N = \sup_t \max_j |G_N(j, t)|$.

3. BOHR’S SUPREMUM AND THE CEILING U_N

Proposition 1 (Bohr ceiling). *For every N and every node j , $\sup_t |G_N(j, t)| \leq U_N$, $U_N := \frac{1}{2N} \sum_{r=1}^{N-1} \csc(\pi r / N)$, with equality at $j = 0$ provided the frequencies $\{\omega_r : 1 \leq r \leq \lfloor N/2 \rfloor\}$ are linearly independent over $\mathbb{Q}$.*

Proof. Grouping the degenerate pair $r, N - r$ gives, at $j = 0$, $G_N(0, t) = \sum_{r=1}^{\lfloor N/2 \rfloor} b_r \sin(\omega_r t)$ with $b_r = 2/(N\omega_r) > 0$ (the term $r = N/2$, if present, has coefficient $1/(N\omega_{N/2})$). The triangle inequality gives $|G_N(j, t)| \leq \frac{1}{N} \sum_{r=1}^{N-1} \omega_r^{-1} = U_N$ for all j . If the ω_r are $\mathbb{Q}$ -independent, then by Kronecker–Weyl the orbit $\{(\omega_r t)_r\}$ is dense in the torus, so $\sin(\omega_r t) \rightarrow 1$ simultaneously and the supremum at $j = 0$ equals the bound. (Bohr’s theorem: $\sup_t \sum b_r \sin(\omega_r t) = \sum |b_r|$.) $\square$

4. THE EXACT CONSTANT

Lemma 2 (Euler–Maclaurin asymptotics). $U_N = \frac{1}{\pi} \left(\ln N + \gamma + \ln \frac{2}{\pi} \right) + o(1) = \frac{1}{\pi} \ln N + 0.03999 \dots + o(1)$.

The leading term comes from the singularity $\csc(\pi r/N) \sim N/(\pi r)$ at small r , which yields the harmonic sum $\sum 1/r \sim \ln$; the constant follows from the Euler–Maclaurin correction. Numerically $U_N - \frac{1}{\pi} \ln N \rightarrow 0.03999$ for $N \leq 65536$ (numerics/verify_lemma3_lebesgue.py).

5. UPPER BOUNDS: ENERGY AND THE SZEGŐ SYMBOL

A fixed-energy bound caps the response. With energy $E = \frac{1}{2} \dot{u}^\top \dot{u} + \frac{1}{2} u^\top L_N u$ and the exact trigonometric identity (the cotangent sum, derived in Section 6)

$$\sum_{r=1}^{N-1} \frac{1}{\sin^2(\pi r/N)} = \frac{N^2 - 1}{3}, \quad (5)$$

Cauchy–Schwarz on the modal expansion gives $\|u(t)\|_\infty \leq C\sqrt{N} \sqrt{E}$: at fixed energy the deflection cannot exceed order $\sqrt{N}$, a ceiling attained by the time-reversed focusing configuration (the variational maximum; verify_theorem1.py). The same sum yields the exact pseudo-inverse diagonal $(L_N^+)_{jj} = \frac{1}{N} \sum_{r \neq 0} \lambda_r^{-1} = \frac{N^2 - 1}{12N}$.

Asymptotically the eigenvalues follow the symbol through Szegő’s theorem [31, 25]: for continuous F ,

$$\frac{1}{N} \sum_{r=0}^{N-1} F(\lambda_r) \xrightarrow{N \rightarrow \infty} \frac{1}{2\pi} \int_0^{2\pi} F(f(\theta)) d\theta, \quad f(\theta) = 4 \sin^2(\theta/2), \quad (6)$$

so spectral sums such as U_N and $(L_N^+)_{jj}$ are Riemann sums of the symbol; the logarithmic growth of U_N is exactly the borderline (logarithmically divergent) singularity $f^{-1/2} \sim |\theta|^{-1}$ at $\theta = 0$. (Szegő’s theorem proper requires continuous F ; the *singular* sums $U_N = \frac{1}{2N} \sum \lambda_r^{-1/2}$ and $(L_N^+)_{jj} = \frac{1}{N} \sum \lambda_r^{-1}$ are made rigorous by Euler–Maclaurin at the band edge, Appendix B, the singularity producing the logarithm and the N^2 growth.) The propagators $\cos(t\sqrt{L_N})$, e^{-itL_N} , $L_N^{-1/2} \sin(t\sqrt{L_N})$ are functions of the matrix in the sense of Golub–Van Loan [33], evaluated exactly by (1); this puts the operator-norm functionals $\|f_t(L_N)\|_{\ell^p \rightarrow \ell^q}$ on a rigorous footing.

6. THE SPECTRAL ZETA FUNCTION AND THE LARGE-WAVE CONSTANTS

Define the spectral zeta of the ring Laplacian (zero mode removed),

$$\zeta_{L_N}(s) = \sum_{r=1}^{N-1} \lambda_r^{-s} = 4^{-s} \sum_{r=1}^{N-1} \csc^{2s} \frac{\pi r}{N}, \quad \Re s > 0. \quad (7)$$

At integer s this is a cosecant power sum with a closed form in N (Bernoulli numbers, Appendix B); the first two are

$$\zeta_{L_N}(1) = \frac{N^2 - 1}{12}, \quad \zeta_{L_N}(2) = \frac{(N^2 - 1)(N^2 + 11)}{720}.$$

The first is elementary: from the classical identity $\sum_{r=1}^{N-1} \cot^2 \frac{\pi r}{N} = \frac{(N-1)(N-2)}{3}$ and $\csc^2 = 1 + \cot^2$,

$$\sum_{r=1}^{N-1} \csc^2 \frac{\pi r}{N} = (N-1) + \frac{(N-1)(N-2)}{3} = \frac{N^2 - 1}{3}, \quad \text{so} \quad \zeta_{L_N}(1) = \frac{1}{4} \sum_r \csc^2 \frac{\pi r}{N} = \frac{N^2 - 1}{12},$$

and $\zeta_{L_N}(2)$ follows the same way from $\sum_r \csc^4 = \frac{(N^2-1)(N^2+11)}{45}$ (all integer values are even polynomials in N , `verify_spectral_zeta.py`). The two large-wave constants are *special values of this single object*:

$$U_N = \frac{\zeta_{L_N}(1/2)}{N}, \quad M_N^2 = \frac{2E_0 \zeta_{L_N}(1)}{N}, \quad (8)$$

the first because $\csc(\pi r/N) = 2\lambda_r^{-1/2}$, the second because $(L_N^+)_{jj} = \zeta_{L_N}(1)/N$. Thus the Bohr ceiling (the large wave, $A_N \sim U_N$) and the variational ceiling M_N (Section 5) are the values of ζ_{L_N} at $s = \frac{1}{2}$ and $s = 1$.

The two growth laws are then dictated by the analytic structure of ζ_{L_N} . By Szegő's theorem the normalized zeta converges to the symbol integral,

$$\frac{1}{N} \zeta_{L_N}(s) \xrightarrow{N \rightarrow \infty} I(s) = \frac{1}{2\pi} \int_0^{2\pi} (2 - 2 \cos \theta)^{-s} d\theta = \frac{4^{-s} \Gamma(\frac{1}{2} - s)}{\sqrt{\pi} \Gamma(1 - s)}, \quad \Re s < \frac{1}{2}, \quad (9)$$

which has a simple pole at $s = \frac{1}{2}$ (Figure 2). The edge of convergence is exactly the logarithmic singularity producing $U_N = \zeta_{L_N}(\frac{1}{2})/N = \frac{1}{\pi} \ln N + O(1)$ (Lemma 2). At $s = 1$ the density diverges, and indeed $\zeta_{L_N}(1) \sim N^2/12$ yields the $\sqrt{N}$ variational ceiling $M_N \sim \sqrt{N/6}$. *The large wave sits precisely at the convergence edge $s = \frac{1}{2}$ of the limiting spectral density $I(s)$* ; the finite- N ζ_{L_N} is itself an entire finite sum, and it is the $N \rightarrow \infty$ density that carries the pole. The finite- N zeta is an Epstein/Hurwitz-type object whose integer values are polynomials in N^2 and whose half-integer value carries the logarithm. All identities are verified (`verify_spectral_zeta.py`). At the opposite end, the spectral determinant is the $s = 0$ derivative: $\zeta'_{L_N}(0) = -\log \det' L_N$ with $\det' L_N = \prod_{r \neq 0} \lambda_r = N^2$ (from $\prod_{r=1}^{N-1} 2 \sin(\pi r/N) = N$), so by Kirchhoff's matrix-tree theorem the cycle C_N has exactly $\det' L_N/N = N$ spanning trees—a graph invariant read off $\zeta'_{L_N}(0)$ (`verify_spanning_trees.py`).




FIGURE 2. The normalized spectral zeta $\frac{1}{N} \zeta_{L_N}(s)$ converges to the symbol density $I(s)$ of (9), which has a pole at $s = \frac{1}{2}$. The large-wave ceiling is the boundary value $U_N = \zeta_{L_N}(\frac{1}{2})/N \sim \frac{1}{\pi} \ln N$.

7. OPERATOR-NORM HIERARCHY AND THE VARIATIONAL PRINCIPLE

The large wave is one corner of a hierarchy of operator norms of the velocity propagator $S_N(t) = f_t(L_N)$ on the zero-mean subspace $\mathcal{H}_0 = \{x \in \mathbb{R}^N : \sum_j x_j = 0\}$, where $f_t(\lambda) = \sin(t\sqrt{\lambda})/\sqrt{\lambda}$ for $\lambda > 0$ and $f_t(0) = 0$ (circulant on $\mathcal{H}_0$, kernel G_N ; the constant mode is excluded since $L_N \mathbf{1} = 0$). Its time-supremum in different $\ell^p \rightarrow \ell^q$ norms obeys *three distinct growth laws* (Table 2, all verified by `verify_operator_norms.py`):

$$\begin{aligned} \sup_t \|S_N(t)\|_{2 \rightarrow 2} &= \frac{1}{2 \sin(\pi/N)} \sim \frac{N}{2\pi}, \\ \sup_t \|S_N(t)\|_{2 \rightarrow \infty} &\leq \sqrt{\frac{\zeta_{L_N}(1)}{N}} \sim \sqrt{\frac{N}{12}}, \\ \sup_t \|S_N(t)\|_{1 \rightarrow \infty} &= A_N \leq U_N \sim \frac{1}{\pi} \ln N. \end{aligned} \tag{10}$$

Derivations. $S_N(t)$ is symmetric with eigenvalues $f_t(\lambda_r) = \sin(t\omega_r)/\omega_r$, so its $\ell^2 \rightarrow \ell^2$ norm is $\max_r |\sin(t\omega_r)|/\omega_r$, whose supremum over t is $1/\omega_1 = 1/(2 \sin \frac{\pi}{N})$ —the slowest mode reaching $\sin(t\omega_1) = \pm 1$; this is exact and unconditional. For $\ell^2 \rightarrow \ell^\infty$, by translation invariance the norm is the ℓ^2 -length of one row of G_N , and Parseval gives

$$\|G_N(\cdot, t)\|_2^2 = \frac{1}{N} \sum_{r=1}^{N-1} \frac{\sin^2(t\omega_r)}{\omega_r^2} \leq \frac{1}{N} \sum_r \omega_r^{-2} = \frac{\zeta_{L_N}(1)}{N},$$

hence $\sup_t \|S_N(t)\|_{2 \rightarrow \infty} \leq \sqrt{\zeta_{L_N}(1)/N} \sim \sqrt{N/12}$, with equality iff all $\sin^2(t\omega_r) = 1$ simultaneously—rational independence (N prime, 2^m). Finally $\|S_N(t)\|_{1 \rightarrow \infty} = \max_{j,k} |S_N(t)_{jk}| = \max_m |G_N(m, t)|$, whose time-supremum is, by definition, $A_N \leq \sum_r |a_r| = U_N$. The three corners thus read the spectral sums $\sum_r \omega_r^{-p}$ ($p = 1, 2$) through different lenses; the $\ell^1 \rightarrow \ell^\infty$ corner—weakest input, strongest output—is precisely the large wave A_N .

Variational principle. The energy norm has the direct reading: maximizing the deflection at a node over the energy sphere $E(u_0, v_0) \leq E_0$ on $\mathcal{H}_0$ is a constrained quadratic problem solved by Lagrange multipliers,

$$\sup_{E(u_0, v_0) \leq E_0} \max_j |u_j(t)| = \sqrt{2E_0 (L_N^+)_{jj}} = \sqrt{\frac{E_0(N^2-1)}{6N}} = M_N,$$

independent of t (the focusing configuration is the time-reversed peak). This is the operator-norm/variational face of the energy ceiling, and via (8) both M_N and U_N are special values of the spectral zeta. The Szegő symbol $f(\theta) = 4 \sin^2(\theta/2)$ controls all three norms through the spectral sums $\sum_r \omega_r^{-p}$.

norm	$\sup_t \ S_N(t)\ $	growth	meaning
$\ell^2 \rightarrow \ell^2$	$1/(2 \sin \frac{\pi}{N})$	$N/2\pi$	lowest-mode resonance
$\ell^2 \rightarrow \ell^\infty$	$\leq \sqrt{\zeta_{L_N}(1)/N}$	$\sqrt{N/12}$	energy ceiling / variational M_N
$\ell^1 \rightarrow \ell^\infty$	$A_N \leq U_N = \zeta_{L_N}(1/2)/N$	$\frac{1}{\pi} \ln N$	large wave A_N

TABLE 2. Three growth laws from one propagator. The $\ell^2 \rightarrow \ell^2$ supremum is exact and unconditional; the other two equal the listed ceilings on the $\mathbb{Q}$ -independent classes N prime and $N = 2^m$ (for $N = 2p$, Proposition 17 gives only an $O(N^{-1})$ gap for the $\ell^1 \rightarrow \ell^\infty$ large wave) and are bounded by them in general. The large wave is the $\ell^1 \rightarrow \ell^\infty$ corner.

8. INDEPENDENCE THEOREMS AND UNCONDITIONAL GROWTH

Theorem 3 (odd prime). *Let p be an odd prime. Then $\{\sin(\pi r/p) : 1 \leq r \leq \frac{p-1}{2}\}$ is linearly independent over $\mathbb{Q}$; hence $A_p = U_p \sim \frac{1}{\pi} \ln p$.*

Proof. Put $\eta = \zeta_{2p} = e^{i\pi/p}$. As p is odd, $[\mathbb{Q}(\eta) : \mathbb{Q}] = \varphi(2p) = p - 1$, so $\{1, \eta, \dots, \eta^{p-2}\}$ is a $\mathbb{Q}$ -basis, and $\eta^p = e^{i\pi} = -1$. From $\eta^{-r} = \eta^{2p-r} = \eta^p \eta^{p-r} = -\eta^{p-r}$ we get $2i \sin(\pi r/p) = \eta^r - \eta^{-r} = \eta^r + \eta^{p-r}$. Suppose $\sum_{r=1}^{(p-1)/2} c_r \sin(\pi r/p) = 0$ with $c_r \in \mathbb{Q}$. Then $\sum_r c_r \eta^r + \sum_r c_r \eta^{p-r} = 0$. As r runs over $\{1, \dots, \frac{p-1}{2}\}$ the exponents r fill $\{1, \dots, \frac{p-1}{2}\}$ and $p-r$ fill $\{\frac{p+1}{2}, \dots, p-1\}$, together covering $\{1, \dots, p-1\}$ bijectively, so $\sum_{m=1}^{p-1} d_m \eta^m = 0$ with $d_m = c_m$ or c_{p-m} . Finally $\Phi_{2p}(\eta) = \sum_{j=0}^{p-1} (-1)^j \eta^j = 0$ gives $\eta^{p-1} = -\sum_{j=0}^{p-2} (-1)^j \eta^j$, whose constant term is $-1 \neq 0$; equating coordinates in $\{1, \eta, \dots, \eta^{p-2}\}$ forces $d_{p-1} = 0$ and then all $d_m = 0$, i.e. all $c_r = 0$. The conclusion follows from Proposition 1. $\square$

Theorem 4 (powers of two). *For $N = 2^m$, $\{\sin(\pi r/N) : 1 \leq r \leq N/2\}$ is $\mathbb{Q}$ -independent; hence $A_N = U_N \sim \frac{1}{\pi} \ln N$.*

Proof. Now $\eta = \zeta_{2N}$, $\Phi_{2N}(x) = x^N + 1$ (degree $\varphi(2N) = N$), so $\{1, \dots, \eta^{N-1}\}$ is a basis and $\eta^N = -1$. Again $2i \sin(\pi r/N) = \eta^r + \eta^{N-r}$, and for $1 \leq r \leq N/2$ the exponents r and $N-r$ lie in $\{1, \dots, N-1\}$ without reduction; they are distinct basis powers (the index $N/2$ occurring with coefficient $2c_{N/2}$). Linear independence is immediate. $\square$

An exact integer-rank certificate validates the prime case (`verify_lemma1_prime.py`); the general rank formula is checked separately by `verify_independence_rank.py`.

Worked example (prime $N = 5$). The two distinct frequencies $\omega_1 = 2 \sin 36^\circ = 1.1756$ and $\omega_2 = 2 \sin 72^\circ = 1.9021$ stand in the ratio $\omega_2/\omega_1 = 2 \cos 36^\circ = \varphi = \frac{1+\sqrt{5}}{2}$, the *golden ratio*—irrational, so the relation lattice is trivial ($\Lambda_5 = \{0\}$) and Theorem 3 applies. The node-0 response is the two-tone signal $G_5(0, t) = \frac{2}{5}(\omega_1^{-1} \sin \omega_1 t + \omega_2^{-1} \sin \omega_2 t)$, and by Kronecker both sines reach 1 at a common near-alignment time, so

$$A_5 = U_5 = \frac{1}{5}(\csc 36^\circ + \csc 72^\circ) = 0.5506,$$

the supremum equal to the ceiling, approached arbitrarily closely (in contrast to the composite $N = 12$ below; `verify_exact_AN.py`).

The low modes carry most of the ceiling, and—crucially—a *long* independent prefix can always be aligned. Let $M(N)$ be the largest M with $\{\omega_1, \dots, \omega_M\}$ $\mathbb{Q}$ -independent.

Lemma 5 (prefix independence). *$M(N) = \frac{1}{2}\varphi(2N)$: the first $\frac{1}{2}\varphi(2N)$ frequencies $\{\omega_r = 2 \sin(\pi r/N) : 1 \leq r \leq \frac{1}{2}\varphi(2N)\}$ are $\mathbb{Q}$ -independent.*

Proof. A nontrivial relation $\sum_{r=1}^K k_r \sin(\pi r/N) = 0$ with $k_K \neq 0$ (K its largest index)—a vanishing sum of roots of unity [7]—gives, with $\zeta = \zeta_{2N}$ and $P(x) = \sum_{r=1}^K k_r x^r$,

$$Q(x) = x^K (P(x) - P(1/x)) = \sum_{r=1}^K k_r (x^{K+r} - x^{K-r}) \in \mathbb{Z}[x],$$

nonzero of degree $2K$ (leading coefficient k_K), with $Q(\zeta) = \zeta^K \sum_r k_r (\zeta^r - \zeta^{-r}) = 0$. Hence $\Phi_{2N} \mid Q$ and $\varphi(2N) = \deg \Phi_{2N} \leq 2K$. But Q is *anti*-palindromic, $x^{2K} Q(1/x) = -Q(x)$, whereas Φ_{2N} is palindromic; so $2K = \varphi(2N)$ is impossible (it would force $Q = k_K \Phi_{2N}$, simultaneously palindromic and anti-palindromic, hence 0). Thus $2K > \varphi(2N)$, and as both are even, $K \geq \frac{1}{2}\varphi(2N) + 1$. No relation has all its indices $\leq \frac{1}{2}\varphi(2N)$, so $\{\omega_1, \dots, \omega_{\varphi(2N)/2}\}$ are independent and $M(N) \geq \frac{1}{2}\varphi(2N)$; with $M(N) \leq \text{rank}_{\mathbb{Q}}\{\omega_r\} = \frac{1}{2}\varphi(2N)$ (Theorem 18), equality holds. $\square$

Theorem 6 (order). $A_N \sim \frac{1}{\pi} \ln N$ for every N ; precisely $A_N = \frac{1}{\pi} \ln N + O(\ln \ln \ln N)$.

Proof. The upper bound is $A_N \leq U_N = \frac{1}{\pi} \ln N + O(1)$ (Proposition 1, Lemma 2). For the lower bound, write $b_r = 2/(N\omega_r) > 0$, $L_{\text{pre}} = \sum_{r=1}^M b_r$, $M = \frac{1}{2}\varphi(2N)$. By Lemma 5 the $\omega_1, \dots, \omega_M$ are $\mathbb{Q}$ -independent, so by Kronecker there are times t with $\sin(\omega_r t) \rightarrow 1$ for all $r \leq M$ at once; at such t , since the remaining $\sin(\omega_r t) \geq -1$,

$$A_N \geq \sup_t G_N(0, t) \geq L_{\text{pre}} - \sum_{r=M+1}^{\lfloor N/2 \rfloor} b_r = 2L_{\text{pre}} - U_N.$$

As $\sin(\pi r/N) \leq \pi r/N$, $b_r = 1/(N \sin(\pi r/N)) \geq 1/(\pi r)$, so $L_{\text{pre}} \geq \frac{1}{\pi} \sum_{r=1}^M \frac{1}{r} \geq \frac{1}{\pi} \ln M$, giving

$$A_N \geq \frac{2}{\pi} \ln M - U_N = \frac{1}{\pi} \ln \frac{M^2}{N} - O(1) = \frac{1}{\pi} \left(2 \ln \frac{\varphi(2N)}{2} - \ln N \right) - O(1).$$

By the Rosser–Schoenfeld bound $\varphi(2N) \geq 2N/(e^\gamma \ln \ln 2N + 3/\ln \ln 2N)$ [8], the bracket is $\ln N - O(\ln \ln \ln N)$; hence $A_N \geq \frac{1}{\pi} \ln N - O(\ln \ln \ln N)$, and with the upper bound $A_N = \frac{1}{\pi} \ln N + O(\ln \ln \ln N)$ (verify_order_theorem.py). $\square$

The former obstruction—the destructive interference of the out-of-prefix modes—is dispatched by the crude bound $\sin \geq -1$, which suffices because Lemma 5 makes the alignable prefix carry $2L_{\text{pre}} - U_N = (1 - o(1))U_N$ of the ceiling. The exact saturation locus $A_N = U_N$ is the sharper question settled by Theorem 16.

Remark 7 (the defect: an odd-part scaling law). The same estimate controls the deficit on the non-saturating N : $U_N - A_N \leq 2(U_N - L_{\text{pre}}) = 2 \sum_{r=\frac{1}{2}\varphi(2N)+1}^{\lfloor N/2 \rfloor} b_r = \frac{2}{\pi} \ln \frac{N}{\varphi(2N)} + O(1)$ (the tail of the ceiling beyond the prefix; $b_{N/2} = \frac{1}{2N}$ when N is even). Writing $N = 2^a m$ with m odd, $N/\varphi(2N) = m/\varphi(m)$ is independent of a , so the bound is the same throughout each family $2^a m$:

$$U_N - A_N \leq \frac{2}{\pi} \ln \frac{m}{\varphi(m)} + O(1),$$

and the exact defect converges within the family to a limit depending only on the odd part m (numerically $U_N - A_N < 0.18$ for composite $N \leq 60$, e.g. 0.05, 0.08, 0.09, 0.10 along $m = 3$; verify_exact_AN.py). Globally, since $m/\varphi(m) \sim e^\gamma \ln \ln m$ along primorials, this is $O(\ln \ln \ln N)$ —slower than any iterate of the logarithm. On all the numerical evidence the defect is *not* bounded, and for a clean reason: the excess $A_N - L_{\text{pre}}$ stays in $[0.05, 0.12]$ across every tested N ($\omega(m) \leq 4$, up to $N = 1155$ with prefix dimension 240; the optimized A_N is a lower bound, so these are lower estimates of the excess), so numerically

$$U_N - A_N = (U_N - L_{\text{pre}}) - (A_N - L_{\text{pre}}) = \frac{1}{\pi} \ln \frac{m}{\varphi(m)} - O(1) \longrightarrow \infty$$

along primorials ($\Theta(\ln \ln \ln N)$, verify_defect_growth.py), making the proven $O(\ln \ln \ln N)$ upper bound sharp in order. A rigorous *growing* lower bound is thus equivalent to the *excess lemma* $A_N \leq L_{\text{pre}} + O(1)$ —the long alignable prefix captures A_N up to a constant. By the Kronecker–Weyl identification (12), $A_N = \max_{\theta \in \mathbb{T}_N} \sum_r b_r \sin \theta_r$ over the subtorus cut out by the relation lattice, so the lemma is precisely a *quantitative comparison on the relation-constrained subtorus*: the constrained maximum exceeds the free-prefix maximum L_{pre} by only $O(1)$. This is a sup-side cancellation statement—a primal analogue of the McGehee–Pigno–Smith/Konyagin L^1 lower bound [16, 17], and in the same circle as the Chowla cosine problem on extrema of cosine sums [14]—but it runs in the opposite direction (an *upper* bound on a *sup*, driven by the relation lattice rather than by the number of frequencies), and a literature search finds no off-the-shelf theorem that supplies it: it is overwhelmingly supported numerically but open.

The lemma's content is sharply localized. Since $A_N \leq U_N$ and $A_N \geq 2L_{\text{pre}} - U_N$ (Kronecker), one has $|A_N - L_{\text{pre}}| \leq U_N - L_{\text{pre}} = \frac{1}{\pi} \ln \frac{m}{\varphi(m)} + O(1)$ unconditionally; so the excess lemma is *trivial whenever the deficit is bounded*—i.e. whenever $\omega(m)$ is bounded—and acquires content only as $\omega(m) \rightarrow \infty$. In particular it holds outright when the odd part is a prime power ($\omega(m) \leq 1$): for $N = 2p$ the single dependent mode $\omega_p = 2$ has weight $\frac{1}{2N}$, giving the sharp two-sided $|A_N - L_{\text{pre}}| \leq \frac{1}{2N}$ (Proposition 17); for $N = 4p$ the two dependent modes $\{2p - 1, 2p\}$ give $|A_N - L_{\text{pre}}| \leq U_N - L_{\text{pre}} = \frac{1}{N}(\sec \frac{\pi}{4p} + \frac{1}{2}) = O(1/N)$, and the same bookkeeping settles any family with $O(1)$ dependent modes (verify_excess_smallcases.py). The genuinely open case is $\omega(m) \geq 2$, where the dependent band carries $\Theta(N)$ modes of total weight $\rightarrow \infty$ yet the excess stays ≈ 0.1 ; the accessible range ($\omega(m) \leq 3$) shows the small excess is a *balance* of a positive dependent net (≈ 0.17) against a negative prefix deficit (≈ -0.08) rather than near-total sign cancellation, so the asymptotic mechanism remains conjectural. For small N the strict gap $A_N < U_N$ is *certified rigorously*: the Lipschitz-grid certificate gives $U_N - A_N \geq 0.049, 0.057, 0.030, 0.082$ for $N = 6, 9, 12, 15$ (verify_defect_certified.py), and a trigonometric Lasserre moment–SOS sharpens it (tight at $N = 15$: $U_N - A_N = 0.1475$, matching the optimizer to five decimals, at solver tolerance) and reaches $N = 105$ ($M = 24$), though too loosely there to certify the growth (verify_defect_moment_sos.py). The exact value is algebraic per N (Proposition 12) with no universal closed form.

The numerical picture condenses into a single clean implication.

Proposition 8 (conditional defect law). *Assume the two-sided excess lemma: there is an absolute constant C with $|A_N - L_{\text{pre}}| \leq C$ for all N . Then, writing $N = 2^a m$ with m odd,*

$$U_N - A_N = \frac{1}{\pi} \ln \frac{m}{\varphi(m)} + O(1) :$$

the deficit depends, to $O(1)$, only on the odd part of N , is unbounded, and is $\Theta(\ln \ln \ln N)$ along primorials—so the proven $O(\ln \ln \ln N)$ upper bound is sharp in order.

Proof. $U_N - A_N = (U_N - L_{\text{pre}}) - (A_N - L_{\text{pre}})$; the second bracket is $O(1)$ by hypothesis, and the first equals $\frac{1}{\pi} \ln(N/\varphi(2N)) + O(1) = \frac{1}{\pi} \ln(m/\varphi(m)) + O(1)$ (the tail computation of Remark 7); along primorials $m/\varphi(m) \sim e^\gamma \ln m$. $\square$

The hypothesis is exactly what the experiments support ($A_N - L_{\text{pre}} \in [0.05, 0.12]$ through $\omega(m) = 4$, prefix dimension 240); proving it is the open core.

The family limits themselves can be pinned precisely: within each family the node-0 defect converges geometrically, $d_a = \text{defect}_\infty(m) - c_1 2^{-a} - c_2 4^{-a}$ (two-term fits with residuals $\sim 5 \cdot 10^{-9}$, the optima polished by Newton iteration on the stationarity system), giving the first precise values of the family constants (Table 3). A PSLQ/inverse-symbolic search over natural logarithmic and radical bases finds no closed form: like β_{odd} , they appear to be new constants (verify_defect_family_limits.py). The ratio $\text{defect}_\infty(m)/[\frac{1}{\pi} \ln \frac{m}{\varphi(m)}]$ is *not* universal (0.80, 0.54, ≈ 0.9), so the $O(1)$ in Proposition 8 genuinely depends on the odd part.

Example 9 (the whole mechanism at $N = 15$). The seven frequencies are $\omega_r = 2 \sin(\pi r/15)$, $r = 1, \dots, 7$ (with $\omega_5 = 2 \sin \frac{\pi}{3} = \sqrt{3}$). Here $\varphi(2N) = \varphi(30) = 8$, so the rank is $\frac{1}{2}\varphi(30) = 4$: by Lemma 5 the prefix $\omega_1, \dots, \omega_4$ is $\mathbb{Q}$ -independent and *no* relation can live in indices ≤ 4 . The first relation appears at index $5 = \frac{1}{2}\varphi(30) + 1$ —exactly the palindromic bound, here sharp—namely

$$\omega_4 + \omega_5 = \omega_1 + 2\omega_2 + \omega_3, \quad \text{i.e. } k = (-1, -2, -1, 1, 1, 0, 0).$$

Its coefficient sum $\sum_r k_r = -2 \equiv 2 \pmod{4}$ is nonzero mod 4, so by the criterion of Theorem 13 the target $(\frac{\pi}{2}, \dots, \frac{\pi}{2})$ is off the subtorus and $A_{15} < U_{15}$: composite 15 does not saturate (Theorem 16).

TABLE 3. Family limits of the composite defect, $\text{defect}_\infty(m) = \lim_a (U - A)(m 2^a)$ at node 0, with the leading term of the conditional law. The $m = 15$ value is optimization-limited (prefix dimension 64 at $N = 240$) and is quoted to fewer digits (`verify_defect_family_limits.py`).

odd part m	3	5	15
$\text{defect}_\infty(m)$	0.1036816	0.0382385	≈ 0.180
$\frac{1}{\pi} \ln(m/\varphi(m))$	0.1291	0.0710	0.2001

Numerically $U_{15} = 0.902$, $A_{15} = 0.754$, a defect 0.148—the largest among the composites $N \leq 24$ because the odd part $m = 15$ carries two prime factors (Remark 7). The lower bound of Theorem 6 aligns the four-mode prefix: $A_{15} \geq 2L_{\text{pre}} - U_{15} = 0.47$, comfortably inside the band $A_N \sim \frac{1}{\pi} \ln N$.

8.1. The Dirichlet segment. The segment chain $\text{tridiag}(-1, 2, -1)$ with Dirichlet ends has the *simple* spectrum $\mu_k = 2 \sin \frac{\pi k}{2(N+1)}$ ($k = 1, \dots, N$): no degeneracy, the structural advantage over the ring, whose twofold $\mu_r = \mu_{N-r}$ halves the independent phases. Writing $\mu_k = 2 \sin(2\pi k/M)$ with $M = 4(N+1)$, the frequencies lie in $\mathbb{Q}(\zeta_M)$ with $2i \sin(2\pi k/M) = \zeta_M^k - \zeta_M^{-k}$, and the same minus-eigenspace argument identifies the rank *exactly*. The μ_k are the first N distinct frequencies of the ring on $N' = 2(N+1)$ nodes; the excluded middle index $N'/2 = N+1$ carries the rational frequency $\mu_{N+1} = 2$, and Theorem 18 applied to N' gives $\dim_{\mathbb{Q}} \text{span}\{\mu_1, \dots, \mu_N, 2\} = \frac{1}{2}\varphi(4(N+1))$.

Theorem 10 (segment rank and classification).

$$\dim_{\mathbb{Q}} \text{span}\{\mu_1, \dots, \mu_N\} = \min(N, \frac{1}{2}\varphi(4(N+1))), \quad (11)$$

and $A_N^{\text{Dir}} = U_N^{\text{Dir}}$ if and only if $N+1$ is prime or a power of two.

Proof. Each $2i\mu_k = \zeta_M^k - \zeta_M^{-k} \in V^-$ ($M = 4(N+1)$), so $\dim_{\mathbb{Q}} \text{span}\{\mu_k\} \leq \frac{1}{2}\varphi(M)$; since adjoining the rational 2 keeps the rank at $\frac{1}{2}\varphi(M)$, the dimension equals $\frac{1}{2}\varphi(M)$ or $\frac{1}{2}\varphi(M) - 1$ according as $2 \in \text{span}\{\mu_k\}$ or not. If $N+1 = 2^m$, the ring on $N' = 2^{m+1}$ is $\mathbb{Q}$ -independent (Theorem 4); then 2 is independent of the μ_k , so $2 \notin \text{span}\{\mu_k\}$ and the rank is $\frac{1}{2}\varphi(M) - 1 = N$. If instead $N+1$ has an odd prime factor p , put $a = N+1$; then $\sum_{j=0}^{p-1} \zeta_M^{a+(M/p)j} = \zeta_M^a \sum_j (\zeta_M^{M/p})^j = 0$, and the only term of its imaginary part reducing to the middle index $N+1$ is $j = 0$ (another would need $(M/p)j \equiv M/2$, i.e. $2j = p$, impossible for p odd), so $\mu_{N+1} = 2$ occurs with coefficient ± 1 and lies in $\text{span}\{\mu_1, \dots, \mu_N\}$; the rank is $\frac{1}{2}\varphi(M)$. An elementary count gives $\frac{1}{2}\varphi(4(N+1)) = 2^a \varphi(m)$ for $N+1 = 2^a m$ (m odd), which exceeds N exactly for $m = 1$ ($N+1 = 2^m$) and is $\leq N$ otherwise (with equality iff $N+1$ is prime); so the rank is $\min(N, \frac{1}{2}\varphi(4(N+1)))$, full ($= N$) iff $N+1$ is prime or a power of two. Full rank gives $A_N^{\text{Dir}} = U_N^{\text{Dir}}$ by Bohr; for composite $N+1$ the lattice Λ^{Dir} is nontrivial and carries a relation with $\sum_r k_r \not\equiv 0 \pmod{4}$, so $A_N^{\text{Dir}} < U_N^{\text{Dir}}$ by the segment ceiling criterion (Corollary 11; verified $N \leq 40$, `exact/gap/segment_mod4.g`, `verify_segment_classification.py`; the prime case also recovers Proposition 17). $\square$

Corollary 11 (segment ceiling criterion). With $\Lambda^{\text{Dir}} = \{k \in \mathbb{Z}^N : \sum_r k_r \mu_r = 0\}$ the relation lattice of the segment frequencies,

$$A_N^{\text{Dir}} = U_N^{\text{Dir}} \iff \sum_r k_r \equiv 0 \pmod{4} \text{ for every } k \in \Lambda^{\text{Dir}}.$$

Proof. $U_N^{\text{Dir}} = \sum_r b_r$ is the ℓ^1 mass, realised iff the $\sin(\mu_r t)$ are brought simultaneously close to 1, i.e. the target $\varphi_* = (\frac{\pi}{2}, \dots, \frac{\pi}{2})$ —the maximiser of $\sum_r b_r \sin \varphi_r$ —lies in the orbit closure $\mathbb{T}_N^{\text{Dir}}$. By Kronecker,

$\varphi_* \in \mathbb{T}_N^{\text{Dir}}$ iff $\langle k, \varphi_* \rangle = \frac{\pi}{2} \sum_r k_r \equiv 0 \pmod{2\pi}$, i.e. $\sum_r k_r \equiv 0 \pmod{4}$, for every $k \in \Lambda^{\text{Dir}}$. The simple Dirichlet spectrum has no $\mu_r = \mu_{N-r}$ degeneracy, so—unlike the ring’s even- N branch (Theorem 13)—no antipodal node contributes; the single node-0 condition decides. As on the ring (Theorem 13, Remark 14), full independence makes the criterion vacuous, and for composite $N + 1$ the lattice carries a violating relation ($\sum_r k_r \not\equiv 0 \pmod{4}$), verified to $N \leq 40$ in `exact/gap/segment_mod4.g`. $\square$

This is the exact analogue of the ring classification (Theorem 16), shifted by the boundary (the ring uses modulus $2N$, the segment $4(N + 1)$). The order Theorem 6 transfers verbatim: Lemma 5 holds with modulus $4(N + 1)$ in place of $2N$ (same palindromic argument), so the independent prefix has length $\frac{1}{2}\varphi(4(N + 1))$ and $A_N^{\text{Dir}} \sim \frac{1}{\pi} \ln N$ for every N (`verify_segment_lower_bound.py`).

9. EXACT A_N FOR COMPOSITE N

For composite N the frequencies satisfy rational relations and the orbit closure is a proper *subtorus* (Weyl). In the cyclotomic coordinates C (row r = integer coordinates of $2i \sin(\pi r/N)$ in $\mathbb{Q}(\zeta_{2N})$) the attainable phases are $\{\varphi = C\psi\}$, whence

$$\sup_t u_j(t) = \max_{\psi \in \mathbb{R}^p} \sum_r a_{rj} \sin((C\psi)_r), \quad p = \varphi(2N), \quad (12)$$

a finite-dimensional trigonometric maximization, reduced exactly and then evaluated numerically (`verify_exact_AN.py`, returning $A_N = U_N$ on primes; a certified global bracket is given for $N = 6$ below). This makes the composite case computable for each N : numerically the subtorus defect on the ring is $A_N/U_N \in [0.84, 0.92]$ in the tested range, so $A_N/\ln N \in [0.28, 0.31]$, consistent with Theorem 6 ($A_N \sim \frac{1}{\pi} \ln N$).

Proposition 12 (algebraic composite cases). *A_N is an algebraic number: writing $x_r = \cos \varphi_r, y_r = \sin \varphi_r$ with $x_r^2 + y_r^2 = 1$ and adjoining, for each generator k of the relation lattice Λ , the constraints $\cos \langle k, \varphi \rangle = 1$, $\sin \langle k, \varphi \rangle = 0$ that cut out the orbit-closure subtorus $\mathbb{T}_N$ (polynomials in the x_r, y_r through the Chebyshev/addition formulas), turns (12) into the maximum of a polynomial over a compact semialgebraic set defined over $\mathbb{Q}$, whose value is algebraic by the Tarski–Seidenberg theorem (real quantifier elimination). For $N = 6$ the relation lattice is one-dimensional and*

$$A_6 = \frac{\sqrt{3}}{9} + \frac{(3 + \sqrt{3})\sqrt[4]{3}}{12\sqrt{2}} \approx 0.55942,$$

with the inner maximum at $\cos \varphi = (\sqrt{3} - 1)/2$. For $N = 12$ the three-term relation $\omega_5 = \omega_1 + \omega_3$ (i.e. $\sin 75^\circ - \sin 15^\circ = \sin 45^\circ$) makes the block two-dimensional; the algebraic degree grows with the arithmetic of N , and no uniform elementary closed form is apparent.

Worked example ($N = 12$). The six distinct frequencies $\omega_r = 2 \sin(\pi r/12)$ are

$$(\omega_1, \dots, \omega_6) = (2 \sin 15^\circ, 1, \sqrt{2}, \sqrt{3}, 2 \sin 75^\circ, 2) \approx (0.518, 1, 1.414, 1.732, 1.932, 2),$$

and their relation lattice Λ_{12} is rank two, generated by

$$\omega_5 = \omega_1 + \omega_3 \quad (2 \sin 75^\circ - 2 \sin 15^\circ = \sqrt{2}), \quad \omega_6 = 2\omega_2 \quad (2 = 2 \cdot 1).$$

Both generators have coefficient-sum $\not\equiv 0 \pmod{4}$ —reading $\omega_5 - \omega_1 - \omega_3 = 0$ as $-1 - 1 + 1 \equiv 3$ and $2\omega_2 - \omega_6 = 0$ as $2 - 1 \equiv 1$ —so by the criterion below (Theorem 13) the ceiling is *not* attained, $A_{12} < U_{12}$. Solving the subtorus maximization (12) numerically gives

$$U_{12} = \frac{1}{24} \sum_{r=1}^{11} \csc \frac{\pi r}{12} = 0.8307, \quad A_{12} = 0.7546, \quad \frac{A_{12}}{U_{12}} = 0.908,$$

the two relations pulling the optimal phases off $\varphi_r = \frac{\pi}{2}$ by just enough to cost 9% of the ceiling (`verify_relation_lattices.py`, `verify_exact_AN.py`).

A sharp arithmetic criterion decides exactly when $A_N = U_N$ (the supremum equalling the ceiling).

Theorem 13 (ceiling criterion). *Let $\Lambda = \{k \in \mathbb{Z}^m : \sum_r k_r \omega_r = 0\}$ be the relation lattice of the distinct positive frequencies $\omega_r = 2 \sin(\pi r/N)$, $1 \leq r \leq m = \lfloor N/2 \rfloor$. Then $A_N = U_N$ if and only if*

$$\sum_r k_r \equiv 0 \pmod{4} \quad \text{for every } k \in \Lambda$$

or, when N is even, the antipodal variant $\sum_r k_r + 2 \sum_r r k_r \equiv 0 \pmod{4}$ holds for every $k \in \Lambda$. The two conditions coincide for all $N \leq 160$ (where only N prime and $N = 2^m$ saturate, Corollary 15).

Proof. Always $A_N \leq U_N$, with $U_N = \sum_r a_{r0}$ the node-0 ceiling. As $a_{rj} = a_{r0} \cos(2\pi rj/N)$ we have $|a_{rj}| \leq a_{r0}$, so no node exceeds U_N , and $|a_{rj}| = a_{r0}$ for all r forces $|\cos(2\pi rj/N)| = 1$, i.e. $2j/N \in \mathbb{Z}$: only $j = 0$ and—for even N — $j = N/2$ can have time-supremum U_N . At $j = 0$ the coefficients $a_{r0} > 0$, so $\sup_t G_N(0, t) = U_N$ iff all $\sin(\omega_r t)$ can be brought simultaneously arbitrarily close to 1, i.e. the target $\varphi_* = (\frac{\pi}{2}, \dots, \frac{\pi}{2})$ lies in the orbit closure $\mathbb{T}_N$. By Kronecker, $\varphi_* \in \mathbb{T}_N$ iff $\langle k, \varphi_* \rangle = \frac{\pi}{2} \sum_r k_r \equiv 0 \pmod{2\pi}$ for every $k \in \Lambda$, i.e. $\sum_r k_r \equiv 0 \pmod{4}$. At $j = N/2$ (even N) the signs $\cos(\pi r) = (-1)^r$ give the alternating target $\varphi_r = \frac{\pi}{2} + \pi r$, and the same argument yields $\langle k, \varphi \rangle = \frac{\pi}{2} \sum_r k_r + \pi \sum_r r k_r \equiv 0 \pmod{2\pi}$, i.e. $\sum_r k_r + 2 \sum_r r k_r \equiv 0 \pmod{4}$. Thus $A_N = U_N$ iff one of the two admissible nodes can realize its target, giving the stated disjunction (`verify_ceiling_criterion.py`). $\square$

Remark 14 (priority). The degenerate (“multiple-frequency”) regime that this criterion characterizes is exactly the subject of Filimonov’s 1992 note on the circular string [44] (Zbl 0761.73053), which we have been unable to consult beyond its abstract and the splash table reproduced in [38] (Table 10.1, the *segment* regime at constant $4/\pi^2$). It is therefore possible that the criterion—or the equivalent characterization of when the ring saturates—is already, in whole or in part, in [44]; this must be checked against the primary source before any priority is claimed.

This refines the independence test: $\mathbb{Q}$ -independence means $\Lambda = \{0\}$, which satisfies the criterion vacuously, recovering $A_N = U_N$ for N prime and $N = 2^m$. The criterion is genuinely weaker in principle—a composite N with dependent frequencies could still saturate, provided every relation has coefficient-sum divisible by 4. Computing Λ exactly by integer (Hermite) reduction of the cyclotomic coordinates and testing the criterion against the directly optimized A_N/U_N , the two agree for all $N \leq 40$, and a criterion-only scan finds *no* composite $N \leq 160$ that saturates: on this range $A_N = U_N$ holds precisely for N prime or $N = 2^m$ (`verify_ceiling_criterion.py`). The *full-independence* route to saturation is in fact closed for *all* N by Theorem 18 (rank $= \frac{1}{2}\varphi(2N) = \lfloor N/2 \rfloor$ only for prime and 2^m , confirmed numerically to $N \leq 500$, `verify_order_extended.py`); a composite saturator would have to satisfy the mod-4 criterion *despite* rational dependence, which none ≤ 160 does. Table 4 makes the lattices explicit; e.g. the generators $\omega_3 = 2\omega_1$ ($N = 6$), $\omega_5 = \omega_1 + \omega_3$ ($N = 12$, cf. Proposition 12), and $\omega_4 + \omega_5 = \omega_1 + 2\omega_2 + \omega_3$ ($N = 15$) all have coefficient-sum $\not\equiv 0 \pmod{4}$.

Corollary 15 (saturating families and a finite scan). *The two proven saturating families, the primes and the powers of two, both have natural density zero. An exhaustive criterion scan finds no further saturating $N \leq 160$, so on this range $\{N : A_N = U_N\} = \{N \text{ prime}\} \cup \{2^m\}$. Theorem 16 promotes this to all N : the saturating set is exactly the primes and the powers of two, hence of density zero. This set is OEIS A174090 [23], with the curious equivalent description: the ring saturates exactly when N cannot be written as a sum of three or more consecutive positive integers.*

Theorem 16 (classification of saturation). *$A_N = U_N$ if and only if N is prime or a power of two.*

TABLE 4. Relation lattices Λ_N of the distinct frequencies ω_r and the mod-4 criterion. Primes and powers of two have $\Lambda_N = \{0\}$ and saturate; every composite N listed carries a relation with coefficient-sum $\not\equiv 0 \pmod{4}$, hence $A_N < U_N$; for even N the antipodal sums $\sum_r k_r + 2 \sum_r r k_r$ are violated likewise (verify_ceiling_criterion.py, check E).

N	class	#freq	rank Λ_N	coeff-sums mod 4	$A_N = U_N?$
5	prime	2	0	—	yes
6	composite	3	1	3	no
8	2^m	4	0	—	yes
9	composite	4	1	3	no
10	composite	5	1	1	no
12	composite	6	2	3, 3	no
15	composite	7	3	2, 3, 3	no
16	2^m	8	0	—	yes
18	composite	9	3	3, 3, 3	no
20	composite	10	2	1, 1	no
21	composite	10	4	1, 3, 3, 3	no
24	composite	12	4	3, 3, 3, 3	no

Proof. If N is prime or $N = 2^m$ then $\Lambda_N = \{0\}$ (Theorems 3, 4) and Theorem 13 gives $A_N = U_N$. Conversely, let N be composite and not a power of two, and let p be an odd prime factor, so $N = ps$ with $s = N/p \geq 2$ and $q := 2N/p = 2s$. As $\zeta_{2N}^q = e^{2\pi i/p}$ is a primitive p th root of unity,

$$\sum_{j=0}^{p-1} \zeta_{2N}^{1+qj} = \zeta_{2N} \sum_{j=0}^{p-1} (\zeta_{2N}^q)^j = \zeta_{2N} \cdot 0 = 0,$$

and taking imaginary parts, $\sum_{j=0}^{p-1} \sin \frac{\pi(1+qj)}{N} = 0$. Reduce each term to a distinct frequency: with $e_j := (1 + qj) \bmod 2N$ one has $\sin \frac{\pi(1+qj)}{N} = \sigma_j \sin \frac{\pi r_j}{N}$, where $r_j \in \{1, \dots, \lfloor N/2 \rfloor\}$ and $\sigma_j = +1$ if $e_j \in (0, N)$, $\sigma_j = -1$ if $e_j \in (N, 2N)$. No e_j lies in $\{0, N\}$: $e_j \equiv 0$ would need $qj \equiv -1 \pmod{2N}$, impossible since $q \mid 2N$ and $q > 1$; $e_j \equiv N$ would need $q \mid N - 1$, but p odd gives $ps \equiv s \pmod{2s}$, hence $N - 1 \equiv s - 1 \pmod{2s}$ with $0 < s - 1 < 2s$, so $2s \nmid N - 1$. Thus every $\sigma_j = \pm 1$, and the vector k with $k_r = \sum_{j: r_j=r} \sigma_j$ satisfies $\sum_r k_r 2 \sin \frac{\pi r}{N} = \sum_j \sigma_j 2 \sin \frac{\pi r_j}{N} = 0$, i.e. $k \in \Lambda_N$, with

$$\sum_r k_r = \sum_{j=0}^{p-1} \sigma_j \equiv p \equiv 1 \pmod{2}.$$

So $\sum_r k_r$ is odd: then $\sum_r k_r \not\equiv 0 \pmod{4}$, and $\sum_r k_r + 2 \sum_r r k_r$ is also odd, hence $\not\equiv 0 \pmod{4}$. The single relation k violates *both* branches of Theorem 13, so neither the node $j = 0$ nor (for even N) the node $j = N/2$ can saturate, and $A_N < U_N$ (verify_classification_proof.py). $\square$

The composite gap. For every composite N , then, $A_N < U_N$ (Theorem 16). The gap is *certified* for $N = 6$: the lone relation $2\omega_1 = \omega_3$ has coefficient-sum $2 + 0 - 1 \equiv 3 \pmod{4} \neq 0$, so by Theorem 13 $A_6 < U_6$; quantitatively a Lipschitz-grid certificate on the two-dimensional subtorus—the elementary level of the Lasserre moment–SOS hierarchy [26]—brackets $A_6 \in [0.5594, 0.5604]$, so $U_6 - A_6 \geq 0.048 > 0$, while the prime $N = 5$ returns $A_5 = U_5$ (verify_sos_certified.py).

Proposition 17 (twice a prime: the sharp constant on a composite family). *For $N = 2p$ with p an odd prime, $A_N = U_N - O(1/N) = \frac{1}{\pi} \ln(2p) + O(1)$; the sharp constant $1/\pi$ holds, extending Theorems 3–4 to an infinite family of composite N .*

Proof. The prefix $\{\omega_1, \dots, \omega_{p-1}\} = \{2 \sin(\pi r/2p) : 1 \leq r \leq p-1\}$ is $\mathbb{Q}$ -independent: in $\mathbb{Q}(\zeta_{4p})$, where $\Phi_{4p}(x) = \Phi_p(-x^2)$ has degree $\varphi(4p) = 2(p-1)$, the values $2i \sin(\pi r/2p) = \zeta_{4p}^r - \zeta_{4p}^{-r}$ ($1 \leq r \leq p-1$) are linearly independent over $\mathbb{Q}$. Indeed, if $\sum_{r=1}^{p-1} c_r (\zeta_{4p}^r - \zeta_{4p}^{-r}) = 0$ with $c_r \in \mathbb{Q}$, multiplying by ζ_{4p}^{p-1} gives $Q(\zeta_{4p}) = 0$, where $Q(x) = \sum_{r=1}^{p-1} c_r (x^{p-1+r} - x^{p-1-r})$ has degree $\leq 2p-2 = \deg \Phi_{4p}$; as Φ_{4p} is the minimal polynomial of ζ_{4p} , this forces $Q = C \Phi_{4p}$, but $Q(1) = 0$ while $\Phi_{4p}(1) = \Phi_p(-1) = 1$, so $C = 0$ and $Q \equiv 0$, and since the exponent sets $\{p-1+r\}_r$ and $\{p-1-r\}_r$ are disjoint, every $c_r = 0$ (exact integer rank = $p-1$, `verify_2p_family.py`). Thus the sole $\mathbb{Q}$ -dependent frequency is the rational $\omega_p = 2$, of weight $\frac{1}{2N}$. With the degeneracy $\omega_r = \omega_{N-r}$ the node-0 response is $u_0(t) = \frac{2}{N} \sum_{r=1}^{p-1} \omega_r^{-1} \sin(\omega_r t) + \frac{1}{2N} \sin 2t$. By Kronecker the independent prefix can be driven to $\sin(\omega_r t) \rightarrow 1$ simultaneously, and since $\frac{2}{N} \sum_{r=1}^{p-1} \omega_r^{-1} = \frac{1}{N} \sum_{r=1}^{p-1} \csc \frac{\pi r}{2p} = U_N - \frac{1}{2N}$, we get $\sup_t u_0(t) \geq U_N - \frac{1}{N}$. As $A_N \leq U_N$ always, $A_N \in [U_N - \frac{1}{N}, U_N]$ (`verify_2p_family.py`). $\square$




FIGURE 3. Ring large wave A_N versus $\ln N$. The ceiling U_N (Lemma 2) grows as $\frac{1}{\pi} \ln N + 0.040$; for primes and $N = 2^m$ the supremum equals U_N (Theorems 3, 4); for composite N the (numerically evaluated) subtorus value lies slightly below, and $A_N \sim \frac{1}{\pi} \ln N$ for every N (Theorem 6).

9.1. Weyl equidistribution and the lower-bound problem. The orbit-closure optimization is rigorous via Weyl's equidistribution theorem [2]. The frequency vector $\omega = (\omega_1, \dots, \omega_m)$ defines the linear flow $t \mapsto t\omega \bmod 2\pi$, whose closure is the subtorus

$$\mathbb{T}_N = \{\varphi \in \mathbb{T}^m : \langle k, \varphi \rangle \equiv 0 \pmod{2\pi} \forall k \in \Lambda\}, \quad \Lambda = \{k \in \mathbb{Z}^m : \langle k, \omega \rangle = 0\},$$

of dimension $\dim \mathbb{T}_N = m - \text{rank } \Lambda = \text{rank}_{\mathbb{Q}}\{\omega_r\}$, the $\mathbb{Q}$ -rank of Section 8. This rank has a closed form, valid for *all* N .

Theorem 18 (subtorus dimension). *For every $N \geq 3$,*

$$\dim \mathbb{T}_N = \text{rank}_{\mathbb{Q}}\{2 \sin(\frac{\pi r}{N}) : 1 \leq r \leq N-1\} = \frac{1}{2}\varphi(2N). \quad (13)$$

Proof. Write $\zeta = \zeta_{2N}$ and let $\sigma : \zeta \mapsto \zeta^{-1}$ be complex conjugation, an involution of $\mathbb{Q}(\zeta)$ whose fixed field is the maximal real subfield $\mathbb{Q}(\zeta + \zeta^{-1})$, of degree $\frac{1}{2}\varphi(2N)$. In characteristic 0, $\mathbb{Q}(\zeta) = V^+ \oplus V^-$ where $V^{\pm}$ are the ± 1 -eigenspaces of σ ; thus $\dim_{\mathbb{Q}} V^{\pm} = \frac{1}{2}\varphi(2N)$, and since σ is the restriction of conjugation, V^- is exactly the set of purely imaginary elements of $\mathbb{Q}(\zeta)$. Now $2i \sin(\pi r/N) = \zeta^r - \zeta^{-r} = (1 - \sigma)\zeta^r$. For any y , $\sigma((1 - \sigma)y) = -(1 - \sigma)y$, so $(1 - \sigma)\mathbb{Q}(\zeta) \subseteq V^-$; and $(1 - \sigma)x = 2x$ for $x \in V^-$, so the inclusion is an equality. The powers $\{\zeta^r : r \in \mathbb{Z}\}$ span $\mathbb{Q}(\zeta)$ over $\mathbb{Q}$ (they contain the power basis $1, \zeta, \dots, \zeta^{\varphi(2N)-1}$), hence their images $\{\zeta^r - \zeta^{-r}\}$ span $(1 - \sigma)\mathbb{Q}(\zeta) = V^-$. Finally $\zeta^N = -1$ gives $\zeta^{N+s} - \zeta^{-(N+s)} = -(\zeta^s - \zeta^{-s})$ while $r = 0, N$ give 0, so the values for $1 \leq r \leq N-1$ already span V^- . Dividing by $2i$ is a $\mathbb{Q}$ -linear bijection of $V^- \subset i\mathbb{R}$ onto $\text{span}_{\mathbb{Q}}\{\sin(\pi r/N)\} \subset \mathbb{R}$, which therefore has dimension $\dim_{\mathbb{Q}} V^- = \frac{1}{2}\varphi(2N)$. $\square$

The formula is independently confirmed by an exact integer-rank certificate for $N \leq 64$ (verify_qrank_formula.py); the dimension sequence $\frac{1}{2}\varphi(2N)$ is OEIS A055034 [23]. Hence full rank—and therefore $A_N = U_N$ —holds *exactly* when $\frac{1}{2}\varphi(2N) = \lfloor N/2 \rfloor$, i.e. for N prime or $N = 2^m$ (Figure 4): precisely the scope of Theorems 3 and 4, now the full-rank case of one formula. Hence (12) is *exactly*

$$A_N = \max_j \max_{\varphi \in \mathbb{T}_N} \frac{2}{N} \sum_r \cos \frac{2\pi r j}{N} \sin \varphi_r.$$

For any S with $\{\omega_r\}_{r \in S}$ $\mathbb{Q}$ -independent, the projection of $\mathbb{T}_N$ onto the S -coordinates is the full torus $\mathbb{T}^{|S|}$, so those phases align and contribute $(2/N) \sum_{r \in S} \omega_r^{-1}$. Taking $S = \{1, \dots, M(N)\}$ the independent prefix, $M(N) = \frac{1}{2}\varphi(2N)$ by Lemma 5, this is the budget L_{pre} of the order Theorem 6.

The lower bound, settled. The earlier obstruction—that the out-of-prefix modes might interfere destructively—is dispatched by the crude bound $\sin \geq -1$: it costs at most $U_N - L_{\text{pre}}$, and because Lemma 5 makes the prefix *long* ($M(N) = \frac{1}{2}\varphi(2N)$, within a $\ln \ln N$ factor of $N/2$), the surviving amount $2L_{\text{pre}} - U_N = (1 - o(1))U_N$ still grows like $\frac{1}{\pi} \ln N$. *Soft* bounds had failed here for a structural reason—the time-averaged second moment gives only $\text{RMS}_t \rightarrow 1/\sqrt{12}$ and the slowest single mode only $b_1 \rightarrow 1/\pi$, both $O(1)$ —so the $\ln N$ genuinely needed the simultaneous alignment of $\Theta(\ln N)$ *arithmetically independent* modes, which Lemma 5 supplies. Table 5 shows the resulting band: the exact $A_N / \ln N$ stays in $[0.28, 0.34]$ across all classes (including the hardest $N = pq$, $p, q \sim \sqrt{N}$, where $A_N / U_N \geq 0.98$), consistent with Theorem 6.

10. REACHABILITY OF THE CEILING: FINITE TIME AND DISORDER

The supremum A_N is an *infinite-time* quantity. Two refinements bring out its physical content: how long one must wait, and how fragile the arithmetic that makes the wait short.

10.1. Finite-time amplification. Define the finite-horizon amplitude

$$A_N(T) = \sup_{0 \leq t \leq T} \max_j |u_j(t)| \leq A_N.$$

Reaching $(1 - \varepsilon)U_N$ requires $\sin(\omega_r t) \approx 1$ for all $m = \lfloor N/2 \rfloor$ frequencies at once—a simultaneous Diophantine approximation of the phases to $\frac{\pi}{2}$. By Dirichlet, such an alignment to tolerance δ *exists* by

TABLE 5. Illustration of Theorem 6. Left: the exact node-0 amplitude keeps $A_N/\ln N$ tightly banded. Right: the $\mathbb{Q}$ -independent prefix contributes $L_{\text{pre}}(N) = \frac{1}{N} \sum_{r \leq M(N)} \text{csc} \frac{\pi r}{N}$, exceeding $0.26 \ln N$ even for highly composite N (now a *proven* ingredient: $M(N) = \frac{1}{2}\varphi(2N)$ frequencies are independent, Lemma 5). $U_N/\ln N \rightarrow 1/\pi = 0.3183$ (`verify_conj_order_map.py`).

N	class	A_N/U_N	$A_N/\ln N$	N	$M(N)$	$L_{\text{pre}}/\ln N$
5	prime	1.000	0.342	30	8	0.261
24	composite	0.915	0.303	120	32	0.274
60	composite	0.873	0.286	210	48	0.268
64	2^m	1.000	0.328	240	64	0.279
90	composite	0.875	0.286	300	80	0.281




FIGURE 4. The subtorus dimension $\frac{1}{2}\varphi(2N)$ (13). Primes and powers of two attain the full rank $\lfloor N/2 \rfloor$ (dotted), so $A_N = U_N$; composite N fall below (the sub-torus defect).

time $\sim \delta^{-m}$ —a crude exponential *upper* bound on the recurrence time

$$T_\varepsilon(N) = \inf\{T : A_N(T) \geq (1 - \varepsilon)U_N\}.$$

(This targets U_N , so $T_\varepsilon(N)$ is finite only on the saturating classes N prime and 2^m ; for non-saturating N , where $A_N < U_N$, one replaces U_N by A_N .) A matching *lower* bound (that one must wait at least exponentially long) does *not* follow from Dirichlet, which bounds the wait from above; it would require Diophantine lower bounds on the badly-approximable phase vector, and is left open. Numerically the wait does grow exponentially: for primes $\log T_{0.2}(N)$ grows by ≈ 0.63 per unit N (`verify_finite_time.py`, Figure 5); for small primes the *upper* bound is moreover certified rigorously—at an explicit witness $\tilde{t}$ the amplitude exceeds $0.8U_N$ in interval arithmetic, so $T_{0.2}(N) \leq \tilde{t}$ ($\tilde{t} \approx 0.69, 0.83, 147, 333$ for $N = 5, 7, 11, 13$; `verify_finite_time_certified.py`). At a *fixed* horizon T the efficiency $A_N(T)/U_N$ decreases with N (from ≈ 1.0 at $N = 5$ to ≈ 0.70 at $N = 29$ for $T = 300$): more frequencies are harder to align in bounded time. Thus the logarithmic law is an infinite-time ceiling; a physical large wave is

the partial early alignment, already a several-fold over-amplification, while the full U_N is approached only over astronomically long times. The clean arithmetic of prime and $2^m N$ is what makes a sizable fraction of U_N reachable *quickly*.

10.2. Disorder. Unequal masses $m_j = 1 + \varepsilon \xi_j$ turn the chain into the generalized problem $L_N v = \omega^2 M v$, $M = \text{diag}(m_j)$, with node-0 response $x_0(t) = \sum_r c_r \sin(\omega_r t)$, $c_r = v_r(0)^2 / \omega_r > 0$ (M -orthonormal modes $v_r^\top M v_s = \delta_{rs}$, unit velocity impulse) and incoherent ceiling $C = \sum_r c_r$. Generic masses destroy every rational relation almost surely. Each $F_k(\xi) = \sum_r k_r \omega_r(\xi)$ ($k \in \mathbb{Z}^{N-1} \setminus \{0\}$) is real-analytic on the simple-spectrum set; it is not *identically* zero because the eigenvalue map $\xi \mapsto (\lambda_r)$ is a submersion—its Jacobian $\partial \lambda_r / \partial m_j = -\lambda_r v_r(j)^2$ has full row rank $N - 1$ at generic masses (verify_disorder.py; full rank is an open condition, certifiable at a witness), so the $\omega_r = \sqrt{\lambda_r}$ vary independently and no nontrivial linear combination can vanish on an open set. Hence each $\{F_k = 0\}$ is a proper analytic subset, of measure zero, and their countable union still has measure zero; so the obstruction of Theorem 13 disappears and *every* N can in principle reach its ceiling. We verify (verify_disorder.py) that the exact resonance residual of a composite N jumps from 0 to $O(\varepsilon)$ under disorder, and that over a long horizon the achievable efficiency A/C rises above the blocked ordered value (e.g. $N = 15$: $0.83 \rightarrow 0.91$). But the gain is purely infinite-time: at a short horizon the same disorder yields no improvement, because the now-independent but near-resonant frequencies need Diophantine time to align. The dichotomy is sharp: an infinitesimal perturbation makes the absolute ceiling attainable in principle while pushing the time to attain it beyond any physical horizon. Infinite-time amplification improves; finite-time amplification does not. The next lemma makes this dichotomy quantitative: at a fixed horizon the disordered response is Lipschitz in ε , so a gain forces the horizon to grow like $1/\varepsilon$.

Lemma 19 (finite horizon is Lipschitz in disorder; a gain needs $T \gtrsim 1/\varepsilon$). *Fix N and a direction ξ with $\|\xi\|_\infty \leq 1$, and let $x_0^{(\varepsilon)}(t) = \sum_r c_r(\varepsilon) \sin(\omega_r(\varepsilon)t)$ be the node-0 response at masses $m_j = 1 + \varepsilon \xi_j$. There is $C_N > 0$ (depending only on N) with*

$$\left| \sup_{0 \leq t \leq T} x_0^{(\varepsilon)}(t) - \sup_{0 \leq t \leq T} x_0^{(0)}(t) \right| \leq C_N (1 + T) \varepsilon, \quad 0 \leq \varepsilon \leq \frac{1}{2}.$$

Hence a fixed gain $\delta > 0$ over the ordered ring needs $\varepsilon \geq \delta / (C_N(1 + T))$: at fixed T an infinitesimal disorder gives no improvement, and to exploit disorder the horizon must grow as $T \gtrsim \delta / (C_N \varepsilon)$.

Proof. Write $B(\varepsilon) = M^{-1/2} L_N M^{-1/2}$ ($M = \text{diag}(1 + \varepsilon \xi_j)$) and $g_t(\lambda) = \sin(\sqrt{\lambda} t) / \sqrt{\lambda}$ on the oscillatory subspace $\lambda > 0$; the zero mode is structural ($L_N \mathbf{1} = 0$ for every mass profile, eigenvector $\mathbf{1}$) and contributes no oscillation. With M -normalized modes one has the closed form $x_0^{(\varepsilon)}(t) = m_0^{-1} [g_t(B(\varepsilon))]_{00}$. For $\|\xi\|_\infty \leq 1$, $\varepsilon \leq \frac{1}{2}$ the family $B(\varepsilon)$ is real-analytic with $\|\partial_\varepsilon B\| \leq K_1$, and its nonzero spectrum stays in a fixed $[\delta_N, 4]$, $\delta_N > 0$ (continuity on the compact ε -interval). On $[\delta_N, 4]$ the symbol g_t is smooth with Lipschitz constant $\leq a_N + b_N t$ (from $g'_t(\lambda) = \frac{t \cos \sqrt{\lambda} t}{2\lambda} - \frac{\sin \sqrt{\lambda} t}{2\lambda^{3/2}}$), so $\|g_t(B(\varepsilon)) - g_t(B(0))\| \leq (a_N + b_N t) K_1 \varepsilon$. Combined with $|m_0^{-1} - 1| \leq 2\varepsilon$ and the t -uniform bound $|[g_t(B)]_{00}| \leq \delta_N^{-1/2}$, this gives $|x_0^{(\varepsilon)}(t) - x_0^{(0)}(t)| \leq C_N(1 + t)\varepsilon$. The supremum over $t \leq T$ is 1-Lipschitz in the sup-norm, and the contrapositive is the stated horizon requirement. $\square$

11. THE DISCRETE SCHRÖDINGER EQUATION ON THE RING

Following Filimonov [48], consider $i\dot{z} = L_N z$, so $z(t) = e^{-itL_N} z(0)$ with kernel

$$K_N(k, t) = \frac{1}{N} \sum_r e^{-it\lambda_r} e^{2\pi i r k / N}.$$




FIGURE 5. Reachability of the ceiling. Left: the recurrence time $T_{0.2}(N)$ to reach $0.8 U_N$ appears to grow exponentially over the tested prime range (log scale). Right: at a fixed horizon T the finite-time efficiency $A_N(T)/U_N$ decreases with N —more frequencies are harder to align in bounded time. The logarithmic law is an infinite-time ceiling.

Here the phases are governed by $\lambda_r = 2 - 2 \cos \frac{2\pi r}{N}$, not by $\omega_r = \sqrt{\lambda_r}$.

Theorem 20 (Schrödinger contrast). (1) (Unitarity.) For $z(0) = e_0$, $\max_j |z_j(t)| \leq 1$ for all t : a localized state is never amplified, in sharp contrast with the bead chain.
 (2) (Ballistic ℓ^∞ amplification.) $B_N := \sup_t \|e^{-itL_N}\|_{\ell^\infty \rightarrow \ell^\infty} = \sup_t \sum_k |K_N(k, t)| \leq \sqrt{N}$ (each row of the unitary e^{-itL_N} has unit ℓ^2 norm, so ℓ^1 norm $\leq \sqrt{N}$). Together with the matching lower bound (Theorem 22), $B_N = \Theta(\sqrt{N})$, not $\ln N$.
 (3) (Arithmetic.) The spectrum lies in the real cyclotomic field $\mathbb{Q}(\zeta_N)^+$ via $\cos$ (whereas $\{\omega_r\} \subset \mathbb{Q}(\zeta_{2N})$ via $\sin$); the affine identity $\sum_r \lambda_r = 2N$ holds always, and $\text{rank}_{\mathbb{Q}}\{\lambda_r\}$ is strictly smaller than $\text{rank}_{\mathbb{Q}}\{\omega_r\}$ for $N = 2^m$ (e.g. 4 vs 8 at $N = 16$).

Lemma 21 (Bessel ℓ^1 asymptotics). As $x \rightarrow \infty$, $\sum_{n \in \mathbb{Z}} |J_n(x)| = c_0 \sqrt{x} (1 + o(1))$ with $c_0 = (2/\pi)^{3/2} B(\frac{1}{2}, \frac{3}{4}) \approx 1.217$.

Proof. Split at the light cone $|n| = x$. Beyond it ($|n| \geq x(1 + \delta)$) the bound $|J_n(x)| \leq (x/2)^{|n|}/|n|!$ gives super-exponential decay; the turning-point band $|n| = x + O(x^{1/3})$ carries $O(x^{1/3})$ terms of size $O(x^{-1/3})$, contributing $O(1)$; and the shell $(1 - \delta)x \leq |n| \leq x$ contributes, by the integrable amplitude bound below, $O(\delta^{3/4} \sqrt{x})$ —all $o(\sqrt{x})$ after $\delta \rightarrow 0$. In the oscillatory region $|n| \leq (1 - \delta)x$ the Debye expansion [9] holds uniformly,

$$J_n(x) = \sqrt{\frac{2}{\pi \sqrt{x^2 - n^2}}} (\cos \Phi_n + O(x^{-1})), \quad \Phi_n = \sqrt{x^2 - n^2} - n \arccos \frac{n}{x} - \frac{\pi}{4},$$

and the phase increments $\Phi_{n+1} - \Phi_n = -\arccos(n/x)$ vary with n , so the $\{\Phi_n\}$ equidistribute and $|\cos \Phi_n|$ averages to $\langle |\cos| \rangle = 2/\pi$ (van der Corput). Hence

$$\sum_{|n| \leq (1-\delta)x} |J_n(x)| = (1 + o(1)) \frac{2}{\pi} \sqrt{\frac{2}{\pi}} \int_{-x}^x \frac{dn}{(x^2 - n^2)^{1/4}} = (1 + o(1)) (2/\pi)^{3/2} B(\frac{1}{2}, \frac{3}{4}) \sqrt{x},$$

the integral equalling $\sqrt{x} B(\frac{1}{2}, \frac{3}{4})$ (set $n = x \sin \theta$). Letting $\delta \rightarrow 0$ gives the claim; the leading constant is confirmed to the digits of Table 6. $\square$

Theorem 22 (Schrödinger lower bound). $\liminf_N B_N/\sqrt{N} \geq c_0/\sqrt{2}$ with $c_0 = (2/\pi)^{3/2} B(\frac{1}{2}, \frac{3}{4})$, i.e. $\liminf_N B_N/\sqrt{N} \geq 0.8607$. With Theorem 20(2), $B_N = \Theta(\sqrt{N})$.

Proof. Evaluate the ℓ^1 mass at the aliasing-free time $t = \frac{N}{4}(1 - \delta)$, so $2t = \frac{N}{2}(1 - \delta)$. By the Poisson periodization (21), $K_N(j, t) = \sum_{\ell} K_{\infty}(j + \ell N, t)$ with $|K_{\infty}(n, t)| = |J_n(2t)|$. The Bessel front sits at $|n| \approx 2t = \frac{N}{2}(1 - \delta)$, and the nearest aliased copy ($\ell \neq 0$) lies at $|j + \ell N| \geq N - 2t = \frac{N}{2}(1 + \delta) > 2t$, where the bound $|J_n(2t)| \leq t^{|n|}/|n|!$ forces super-exponential decay; so $\sum_j |K_N(j, t)| = (1 + o(1)) \sum_{n \in \mathbb{Z}} |J_n(2t)|$. By Lemma 21 this is $(1 + o(1)) c_0 \sqrt{2t} = (1 + o(1)) c_0 \sqrt{N/2} \sqrt{1 - \delta}$; letting $\delta \rightarrow 0$, $\liminf_N B_N/\sqrt{N} \geq c_0/\sqrt{2}$. The upper bound $B_N \leq \sqrt{N}$ is Theorem 20(2) (verify_bn_parity.py). $\square$

The improvement over the cruder $t = N/8$ value $\frac{1}{2}c_0 = 0.6086$ is the optimal aliasing-free time $t = N/4$, at which the Bessel front exactly fills one half-period.

Remark 23 (the candidate constant splits by parity). The sharp constant is expected to be parity-dependent: the aliasing phase i^{-N} is real for even N and imaginary for odd N , so where the two overlapping fronts meet they add differently. Numerically (verify_bn_parity.py) the sup is at $t = N/4$ for even N , giving $B_N/\sqrt{N} \rightarrow c_0/\sqrt{2}$ (the value above), and at $t = N/2$ for odd N , where the two surviving aliasing copies give $|K_N(k, N/2)| \approx \sqrt{J_k(N)^2 + J_{N-k}(N)^2}$ across the front (no cancellation, since i^{-N} is imaginary). The Debye envelope $|J_{Nu}(N)| \sim \sqrt{2/\pi N} (1 - u^2)^{-1/4} |\cos \Phi|$ (the uniform large-order asymptotics of [24, §10.19]) then yields the limiting value

$$\beta_{\text{odd}} = \sqrt{\frac{2}{\pi}} \int_0^1 \mathbb{E}_{\Phi, \Psi} \sqrt{\frac{\cos^2 \Phi}{1 - u^2} + \frac{\cos^2 \Psi}{u(2 - u)}} du = 0.928019304807931 \dots, \quad (14)$$

an elliptic-integral average (Φ, Ψ independent uniform). Computed to 25 digits, it matches no combination of $\pi, e^{\gamma}, \Gamma(\frac{1}{4}), \sqrt{2}$, Catalan's constant, $\dots$, and is not a low-degree algebraic number (an integer-relation/PSLQ search returns nothing): no elementary closed form was found, and we treat β_{odd} as a new constant defined by (14) (verify_beta_odd.py). The evaluations at $t = N/4$ and $t = N/2$ give the rigorous lower bounds $\liminf_N B_N/\sqrt{N} \geq c_0/\sqrt{2}$ (Theorem 22) and $\liminf_{\text{odd } N} B_N/\sqrt{N} \geq \beta_{\text{odd}}$, whence $\limsup_N B_N/\sqrt{N} \in [\beta_{\text{odd}}, 1]$; numerically the even subsequence attains $c_0/\sqrt{2}$, so the global $\liminf$ equals it. For the upper side, on the whole window $t \leq N/2$ the two-copy reduction is rigorous: $2t \leq N$, so only the two Poisson copies $\ell = 0, -1$ survive (the rest are super-exponentially small), $|K_N(k, t)| \leq |J_k(2t)| + |J_{N-k}(2t)|$, and the Debye envelope turns $\|K_N(\cdot, t)\|_1/\sqrt{N}$ into the two-copy profile $F(2t/N)$. This profile is now explicit and its maximizer is *proved* (Proposition 24): $\max F = F(1) = \beta_{\text{odd}}$ at $t = N/2$, analytically on $[\frac{1}{2}, 0.937]$ and by validated interval arithmetic on $[0.937, 1]$. Hence $\|K_N(\cdot, t)\|_1 \leq (\beta_{\text{odd}} + o(1))\sqrt{N}$ for all $t \leq N/2$, and with the lower bound $\sup_{t \leq N/2} \|K_N(\cdot, t)\|_1 = (\beta_{\text{odd}} + o(1))\sqrt{N}$ on the odd subsequence. For $t > N/2$ (three or more overlapping copies) the joint equidistribution of the overlapping Debye phases is now *proved*, with an explicit power saving (Theorem 26, §11.1), and it transfers uniformly in u to the L^1 limit (Theorem 27), so the entire multi-copy constant reduces *unconditionally* to the toral profile $F(s)$ of Proposition 29: $\|K_N(\cdot, sN/2)\|_1 = (F(s) + o(1))\sqrt{N}$ for every fixed $s > 1$, sharper than the bare unitarity $\|K_N(\cdot, t)\|_1 \leq \sqrt{N}$. Numerically the ℓ^1 mass there falls back toward its time average $\approx 0.79\sqrt{N}$ and never exceeds the $t = N/2$ value (checked out to $t = 5N$ at $N = 2^{20}$, where the $t > N/2$ supremum trails the $t = N/2$ value by a stable margin): the $\sin^2$ dispersion admits no flat (Gauss-sum) revival, even at structured Talbot times $t = \pi p/q$. This is consistent with the dispersive-quantization picture—exact revivals of the periodic Schrödinger flow occur only for a *quadratic* dispersion (rational times give Gauss sums) [10], which the discrete $\lambda_r = 4 \sin^2(\pi r/N)$ is not. So numerically the global sup sits at

$t = N/2$, and closing the residual upper bound would give $\limsup_N B_N/\sqrt{N} = \beta_{\text{odd}}$. The residual on $t > N/2$ has a clean form: $K_N(\cdot, t)$ is the inverse DFT of the *unimodular* symbol $(e^{-it\lambda_r})_r$ —a chirp, since $2t \cos(2\pi r/N)$ is locally quadratic—so $B_N/\sqrt{N}$ is the maximal L^1 -flatness ratio $\|P\|_1/\|P\|_2$ of a unimodular trigonometric polynomial. The limiting macroscopic profile of $\|K_N(\cdot, t)\|_1/\sqrt{N}$ is a *multi-copy* Debye average over the Poisson copies (verify `bn_residual_profile.py`: it matches the FFT to 0.5% across all $N \bmod 4$, peaks at $t = N/2$ with β_{odd} , and is strictly subdominant for $t > N/2$, stably out to $t = 5N$), the profile $F(s)$ of Proposition 29. The joint equidistribution mod 2π of the coherent multi-copy Debye phases—the multidimensional Weyl sum on which F rests—is settled in §11.1 (Theorem 26) and its uniform-in- u transfer makes F the unconditional $t > N/2$ profile (Theorem 27); what remains for $\limsup_N B_N/\sqrt{N} = \beta_{\text{odd}}$ is only the profile-maximizer inequality $\max_{s \geq 1} F(s) = F(1)$ (Lemma 30 and its residual Pearson-walk bound). That maximizer step aside, the difficulty is intrinsic: Erdős asked whether every unimodular polynomial must satisfy $\max |P| \geq (1+c)\sqrt{n}$ [12]; Kahane’s ultraflat polynomials answered no [13], with the sharpest construction in [11], and the *quadratic* chirp (Gauss–Fresnel/Talbot) is even L^1 -flat with ratio $\rightarrow 1$. So $\limsup < 1$ is an *anti-flatness* statement particular to the non-quadratic $\sin^2$ dispersion—precisely an Erdős-type flatness question for one structured family, and no generic argument supplies it. (A single-period saddle van der Corput estimate fails: it predicts 0.61 against the actual 0.93, missing the Poisson-copy coherence—which is precisely why the correct object is the *joint* multi-copy equidistribution of Theorem 26, not a single saddle. A fourth-moment/merit-factor route cannot supply the upper bound either: $\max\{\|K\|_1/\sqrt{N} : \|K\|_2 = 1, N\|K\|_4^4 = c\} = 1 - \Theta(\sqrt{c/N})$, a single spike of mass $\sqrt{c/N}$ absorbing any fourth moment at negligible ℓ^1 cost—so no bound below 1 can follow from $\|K\|_2$ and $\|K\|_4$ alone.)

Proposition 24 (two-copy profile and its maximizer). *For $t = sN/2$ with $s \in [0, 1]$, the two-copy Debye approximation of Remark 23 gives $\|K_N(\cdot, t)\|_1 = (F(s) + o(1))\sqrt{N}$ with the explicit scalar profile*

$$F(s) = \begin{cases} c_0 \sqrt{s}, & 0 \leq s \leq \frac{1}{2}, \\ \sqrt{\frac{2}{\pi}} \left[\frac{4}{\pi} (1-s) A(s) + (2s-1)^{3/4} B(s) \right], & \frac{1}{2} \leq s \leq 1, \end{cases}$$

$$A(s) = \int_0^1 (s^2 - (1-s)^2 w^2)^{-1/4} dw, \quad B(s) = \int_0^1 g(\hat{a}, \hat{b}) dw,$$

where $\hat{a} = [(1-w)(1+(2s-1)w)]^{-1/2}$, $\hat{b} = [w(2s-(2s-1)w)]^{-1/2}$, and $g(a, b) = \mathbb{E}_{\Phi, \Psi} \sqrt{a \cos^2 \Phi + b \cos^2 \Psi}$ with Φ, Ψ independent uniform. Equivalently, in the rescaled site $u = k/N$,

$$F(s) = \sqrt{\frac{2}{\pi}} \int_0^1 g((s^2 - u^2)_+^{-1/2}, (s^2 - (1-u)^2)_+^{-1/2}) du,$$

a single integral with $g(a, 0) = \frac{2}{\pi} \sqrt{a}$, and the one-copy factor is the closed form $A(s) = s^{1/2} (1-s)^{-1} {}_2F_1(\frac{1}{4}, \frac{1}{2}; \frac{3}{2}; x^2)$, $x = (1-s)/s$ [24, §15]. Then $F(\frac{1}{2}) = c_0/\sqrt{2}$ and $F(1) = \beta_{\text{odd}}$ (14), and

$$\max_{s \in [0, 1]} F(s) = F(1) = \beta_{\text{odd}}.$$

Proof. The reduction and the closed-form endpoints are the Debye computation of Remark 23: for $s \leq \frac{1}{2}$ the two Bessel fronts $|n| \approx sN$ are disjoint and each half-line contributes $\frac{1}{2} c_0 \sqrt{2t}$, giving $c_0 \sqrt{s}$; for $s \geq \frac{1}{2}$ the overlap $u = k/N \in [1-s, s]$ carries both Debye envelopes, which for odd N add in quadrature, producing A, B after $u \mapsto w$ maps the two windows to $[0, 1]$. The substitution makes the endpoint singularities $w^{-1/4}, (1-w)^{-1/4}$ integrable and the integrands smooth in s on $(\frac{1}{2}, 1]$; at $s = 1$ one recovers exactly (14). For the maximizer, two analytic majorants of g cover the bulk. Jensen’s bound $g(a, b) \leq \sqrt{(a+b)/2}$ ($\mathbb{E} \cos^2 = \frac{1}{2}$) gives $F < \beta_{\text{odd}}$ on $[\frac{1}{2}, \frac{4}{5}]$; writing $g(a, b) = \sqrt{\max h(\min / \max)}$ with

$h(\rho) = \mathbb{E} \sqrt{\cos^2 \Phi + \rho \cos^2 \Psi}$ *concave* ($h'' < 0$ on $(0, 1]$), the tangent at $\rho = 1$, $h(\rho) \leq h(1) + h'(1)(\rho - 1)$ with $h(1) = 0.95809$, $h'(1) = 0.23952$, sharpens this to $F < \beta_{\text{odd}}$ on all of $[\frac{1}{2}, 0.937]$ —an *analytic* bound on 87% of the window, clearing the interior dip $\min F \approx 0.847$ at $s \approx 0.56$. On the remaining $[0.937, 1]$ the profile is strictly increasing, $F'(s) > 0$, proved by *validated interval arithmetic*: two elementary inequalities eliminate the elliptic g —the QM–AM bound $g(a, b) \geq \frac{\sqrt{2}}{\pi}(\sqrt{a} + \sqrt{b})$ (from $\sqrt{x^2 + y^2} \geq (x + y)/\sqrt{2}$) and the derivative bound $\partial_a g = \mathbb{E}[\frac{\cos^2 \Phi}{2\sqrt{\cdot}}] \leq 1/(\pi\sqrt{a})$ —and reduce F' to a lower bound $F_{\text{lo}}(s)$ built from *elementary* ($\sqrt{\cdot}$ -only) integrals; an interval rectangle-sum enclosure (the band-edge $w^{-1/4}$ singularities removed by $w = t^4$, $w = 1 - r^4$, powers cancelled by hand), with s itself taken as an interval, certifies $F_{\text{lo}}(s) > 0$ on every subinterval of $[0.937, 1]$ (worst case 0.051, `verify_bn_profile_rigorous.py`). Hence $F \leq F(1)$ there, and the maximizer $\max_{[0,1]} F = F(1) = \beta_{\text{odd}}$ is proved. $\square$

This settles the $t \leq N/2$ supremum: $\sup_{t \leq N/2} \|K_N(\cdot, t)\|_1 = (\beta_{\text{odd}} + o(1))\sqrt{N}$, since the maximization of the one explicit elliptic-integral function F is now proved in full—analytically on $[\frac{1}{2}, 0.937]$ and by validated interval arithmetic on $[0.937, 1]$. The only remaining piece of the constant is the multi-copy region $t > N/2$ (Remark 23).

11.1. The multi-copy region $t > N/2$: joint equidistribution of the Debye phases. Beyond $t = N/2$, three or more Poisson copies overlap and the two-copy reduction of Proposition 24 is no longer available. We isolate the exact analytic content of the multi-copy regime, prove the one structural input that the heuristic profile rests on—the *joint* equidistribution of the overlapping Debye phases—and state honestly the single point that prevents this from closing the $t > N/2$ bound unconditionally.

Fix a ratio $s = 2t/N > 1$ and the rescaled site $u = k/N \in [0, 1)$. By the periodization (21), $K_N(k, t) = \sum_{\ell} i^{k+\ell N} J_{k+\ell N}(2t)$, and a copy ℓ lies in the oscillatory band exactly when its scaled order $m_{\ell} := u + \ell$ satisfies $|m_{\ell}| < s$. There are at most $2s + 2$ such *live* copies—a number bounded independently of N —and the band-edge copies $|m_{\ell}| = s + O(N^{-1})$ contribute negligibly (Airy front, height $O(N^{-1/3})$; the super-exponential tail handles $|m_{\ell}| > s$ as in Lemma 21). On the live set the uniform large-order asymptotics [24, §10.19] give, with the signed order $n_{\ell} = k + \ell N$ (so $J_{n_{\ell}} = (-1)^{|n_{\ell}|} J_{|n_{\ell}|}$ for $n_{\ell} < 0$),

$$i^{n_{\ell}} J_{n_{\ell}}(2t) = i^{n_{\ell}} \sqrt{\frac{2}{\pi N}} (s^2 - m_{\ell}^2)^{-1/4} (\cos \Phi_{\ell}(k) + O(N^{-1})), \quad (15)$$

with $\Phi_{\ell}(k) = \sqrt{x^2 - n_{\ell}^2} - |n_{\ell}| \arccos(|n_{\ell}|/x) - \frac{\pi}{4}$, $x = 2t = sN$, uniformly on $u \in [\delta, s - \ell - \delta] \cap [\ell - s + \delta, \cdot]$ for each ℓ (we use the phase of $|J_{n_{\ell}}|$; since $\cos$ is even its overall sign is immaterial, and the $(-1)^{|n_{\ell}|}$ of $n_{\ell} < 0$ is absorbed into $i^{n_{\ell}}$). The phase Φ_{ℓ} is the object that decides coherence; its discrete derivative in the site index is exact and elementary.

Lemma 25 (Debye phase increments and their curvature). *For each live copy ℓ and $x = sN$, with $g_{\ell}(u) = (s^2 - m_{\ell}^2)^{-1/2} > 0$, $m_{\ell} = u + \ell$,*

$$\begin{aligned} \Phi_{\ell}(k+1) - \Phi_{\ell}(k) &= -\operatorname{sgn}(m_{\ell}) \arccos \frac{|m_{\ell}|}{s} \pmod{2\pi}, \\ \Phi_{\ell}(k+1) - 2\Phi_{\ell}(k) + \Phi_{\ell}(k-1) &= \frac{1}{N} g_{\ell}(u) + O(N^{-2}). \end{aligned}$$

Consequently, writing $\theta_{\ell}(u) = -\operatorname{sgn}(m_{\ell}) \arccos(|m_{\ell}|/s)$ for the first increment, $\theta'_{\ell}(u) = g_{\ell}(u)$.

Proof. $\Phi_{\ell}(k) = \phi(|n_{\ell}|)$ with $\phi(\nu) = \sqrt{x^2 - \nu^2} - \nu \arccos(\nu/x) - \frac{\pi}{4}$ and $\phi'(\nu) = -\arccos(\nu/x)$. Since $n_{\ell} = k + \ell N$, advancing $k \mapsto k+1$ advances $\nu = |n_{\ell}|$ by $\operatorname{sgn}(n_{\ell})$, so the first difference is $\operatorname{sgn}(n_{\ell}) \phi'(|n_{\ell}|) = -\operatorname{sgn}(m_{\ell}) \arccos(|m_{\ell}|/s)$ (using $\operatorname{sgn}(n_{\ell}) = \operatorname{sgn}(m_{\ell})$ and $|n_{\ell}|/x = |m_{\ell}|/s$).

Hence $\theta_\ell(u) = -\operatorname{sgn}(m_\ell) \arccos(|m_\ell|/s)$, and differentiating, $\theta'_\ell(u) = \operatorname{sgn}(m_\ell) \cdot \frac{\operatorname{sgn}(m_\ell)}{\sqrt{s^2 - m_\ell^2}} = (s^2 - m_\ell^2)^{-1/2} = g_\ell(u) > 0$; the second site-difference is $N^{-1}\theta'_\ell(u) + O(N^{-2})$ by Taylor expansion. The increment identities are verified to machine precision (`verify_debye_equidistribution.py`: ratio 1.0000 for the second difference against $N^{-1}g_\ell$). $\square$

The first increment $\theta_\ell(u)$ varies with u , so each single phase $\{\Phi_\ell(k)\}_k$ equidistributes mod 2π (the classical van der Corput statement behind Lemma 21). The new content is that the live phases equidistribute *jointly*.

Theorem 26 (joint equidistribution of the live Debye phases). *Fix $s > 1$ and a cell $\mathcal{C} = (a, b) \subset [0, 1)$ on which the live set $L(s) = \{\ell : |u + \ell| < s\}$ is constant and no m_ℓ meets the band edge. As $N \rightarrow \infty$ the vector $(\Phi_\ell(k) \bmod 2\pi)_{\ell \in L(s)}$ is equidistributed in the torus $\mathbb{T}^{|L(s)|}$ as k ranges over $\{k : k/N \in \mathcal{C}\}$. Quantitatively, for every nonzero integer vector $(k_\ell)_{\ell \in L(s)}$,*

$$\left| \frac{1}{\#} \sum_{k/N \in \mathcal{C}} e^{i \sum_{\ell} k_\ell \Phi_\ell(k)} \right| \ll_{s, \mathcal{C}, k} N^{-1/2}, \quad (16)$$

the power saving of the second-derivative van der Corput estimate; hence by Weyl's criterion the joint distribution is uniform, with Erdős–Turán discrepancy $O(N^{-1/2} \log N)$ on the cell.

Proof. By Weyl's criterion [2][3, Ch. 1] joint equidistribution is equivalent to (16) for each nonzero (k_ℓ) . Set $\Psi_k(j) := \sum_{\ell \in L(s)} k_\ell \Phi_\ell(j)$. By Lemma 25 its second difference is

$$\Psi_k(j+1) - 2\Psi_k(j) + \Psi_k(j-1) = \frac{1}{N} C_k(u) + O(N^{-2}), \quad C_k(u) = \sum_{\ell \in L(s)} \frac{k_\ell}{\sqrt{s^2 - (u + \ell)^2}},$$

so, viewing Ψ_k as a smooth phase $f(j)$ with $f''(j) = N^{-1}C_k(j/N) + O(N^{-2})$, the van der Corput second-derivative test (the exponent-pair $(\frac{1}{2}, \frac{1}{2})$ bound, [4, §8], [5, Ch. 2]) gives, on a block of $L = \Theta(N)$ consecutive sites with $|f''| \asymp \lambda = N^{-1}|C_k|$,

$$\left| \sum_j e^{if(j)} \right| \ll L \lambda^{1/2} + \lambda^{-1/2} \ll N \cdot N^{-1/2} |C_k|^{1/2} + N^{1/2} |C_k|^{-1/2} \ll N^{1/2},$$

which after dividing by $\# = \Theta(N)$ is (16). The estimate requires C_k bounded away from 0; here is the only structural point. The functions $u \mapsto g_\ell(u) = (s^2 - (u + \ell)^2)^{-1/2}$, $\ell \in L(s)$, have *distinct* singularities at $u = \pm s - \ell$, so they are linearly independent over $\mathbb{R}$: any relation $\sum_{\ell} k_\ell g_\ell \equiv 0$ on an interval extends, by real-analyticity, to a relation between functions with distinct poles, forcing every $k_\ell = 0$. Hence $C_k \not\equiv 0$; being real-analytic on $\mathcal{C}$ it has only finitely many zeros there, and on the complement of ε -neighbourhoods of those zeros $|C_k| \geq c(s, \mathcal{C}, k, \varepsilon) > 0$ and the bound above applies. Each exceptional ε -neighbourhood has length $O(\varepsilon)$ and contributes $\leq \varepsilon$ to the normalized sum trivially; letting $\varepsilon \rightarrow 0$ after $N \rightarrow \infty$ removes it. The discrepancy bound is the Erdős–Turán inequality [3, Ch. 2] applied to (16). The Weyl sums are computed directly: (16) decays with fitted exponent -0.52 to -1.2 across nonzero (k_ℓ) with $|k_\ell| \leq 3$ (the slowest matching the predicted $N^{-1/2}$), the worst over 60 random (k_ℓ) at $N = 2^{18}$ being 0.013 (`verify_debye_equidistribution.py`). $\square$

Theorem 26 is the rigorous core the heuristic profile of `verify_bn_residual_profile.py` assumes. It upgrades the $s \leq 1$ two-copy statement (Proposition 24, where $|L(s)| = 2$) to all s . The cell-wise statement is not yet the L^1 limit: (16) carries an implied constant depending on the cell, and the limit needs the discrepancy $o(1)$ *uniformly* in u , including the finitely many band edges and coalescing points

where the constant degenerates. We supply that uniform control now; it is what turns Theorem 26 into the unconditional profile of Proposition 29.

Theorem 27 (uniform-in- u transfer). *Fix $s > 1$. With the rescaled kernel of (15) and the toral integrand $\mathcal{E}(u) := \mathbb{E}_{(\Phi_\ell) \sim U} \left| \sum_{\ell \in L(s)} c_\ell (s^2 - (u + \ell)^2)^{-1/4} \cos \Phi_\ell \right|$,*

$$\frac{1}{\sqrt{N}} \|K_N(\cdot, sN/2)\|_1 \xrightarrow{N \rightarrow \infty} \sqrt{\frac{2}{\pi}} \int_0^1 \mathcal{E}(u) du = F(s).$$

The convergence is at an explicit polynomial rate $O(N^{-\eta})$, $\eta = \frac{1}{9|L(s)|} > 0$ (the cubic exponent pair $\frac{1}{9}$ at the coalescing points, divided by $|L|$ for the Lipschitz-cone mollification of the modulus in Koksma–Hlawka; numerically the cores decay like N^{-1} , far faster, so the true rate is band-edge-strip-limited at $N^{-1/4+\delta'}$).

Proof. Write $W_N(k) := \sqrt{\pi N/2} |K_N(k, sN/2)|$, so that $\|K_N\|_1/\sqrt{N} = \sqrt{2/(\pi N)} \sum_k |K_N(k)| = \sqrt{2/\pi} \cdot N^{-1} \sum_{k=0}^{N-1} W_N(k)$, the normalized prefactor being $\sqrt{2/\pi}$ (the order density has exactly N sites per unit of $u = k/N$, and the s -dependence sits inside W_N through the envelope; numerically $\|K_N\|_1/\sqrt{N}$ tracks $\sqrt{2/\pi} \int_0^1 \mathcal{E} du$ across $s \in \{1.5, 2.7, 4\}$, while $\sqrt{2/(\pi s)}$ is off by $\sqrt{s}$, `verify_bandedge_transfer.py`). It suffices to show

$$\frac{1}{N} \sum_{k=0}^{N-1} W_N(k) \longrightarrow \int_0^1 \mathcal{E}(u) du. \quad (17)$$

Let $E = \{u \in [0, 1) : |u + \ell| = s \text{ for some } \ell\}$ be the (finite) set of band edges and $Z = \{u^* = -(\ell + \ell')/2 : \ell \neq \ell' \in L, u^* \in [0, 1)\}$ the (finite) set of coalescing points, where a pair $g_\ell, g_{\ell'}$ collides; these are geometric, fixed points of $[0, 1)$, independent of the test modes k . Fix a small constant $\rho > 0$ (independent of N , chosen below) and a small exponent $0 < \delta < \frac{1}{6}$, and put $\varepsilon = N^{-2/3+\delta}$. Partition $[0, 1)$ into

$$\mathcal{S} = \bigcup_{u_0 \in E} (u_0 - \varepsilon, u_0 + \varepsilon), \quad \mathcal{K} = \bigcup_{u^* \in Z} (u^* - \rho, u^* + \rho) \setminus \mathcal{S}, \quad \mathcal{B} = [0, 1) \setminus (\mathcal{S} \cup \mathcal{K}),$$

the band-edge strips (width $\varepsilon \rightarrow 0$), the coalescing cores (fixed width ρ), and the bulk. We bound the contribution of each to (17).

(a) *Bulk $\mathcal{B}$.* By construction every $u \in \mathcal{B}$ is at distance $\geq \rho$ from Z and $\geq \varepsilon$ from E . On each maximal subinterval $\mathcal{C} \subset \mathcal{B}$ the live set L is constant, the envelopes $a_\ell(u) = (s^2 - (u + \ell)^2)^{-1/4}$ are bounded and C^1 , so the rescaled kernel is a fixed function $W_N(k) = G_u(\Phi(k))$ of the phase vector, $G_u: \mathbb{T}^{|L|} \rightarrow \mathbb{R}$ the modulus of a trigonometric polynomial with u -uniformly bounded, C^1 coefficients, $\hat{G}_u(0) = \mathbb{E}_{(\Phi_\ell) \sim U} G_u = \mathcal{E}(u)$. We do *not* Fourier-expand G_u (a modulus is only Lipschitz, its coefficients too slow to sum against the van der Corput bound); instead we bound the *star-discrepancy* $D_N^*(\mathcal{C})$ of the finite phase-point set $\{(\Phi_\ell(k))_{\ell \in L} : k/N \in \mathcal{C}\}$ in $\mathbb{T}^{|L|}$ and invoke Koksma–Hlawka. A modulus is only Lipschitz, and for $|L| \geq 3$ its Hardy–Krause variation is in fact *infinite*: the cone point at $Z = 0$ (codimension 2) makes the top mixed partial grow like $|Z|^{-(|L|-1)}$, and $\int |Z|^{-(|L|-1)} dX dY \asymp \int r^{-(|L|-2)} dr$ diverges (logarithmically already at $|L| = 3$; `verify_bandedge_transfer.py` confirms the $\log N$ growth of the Vitali variation on a refining grid). So Koksma–Hlawka cannot be applied to G_u itself. Instead we use only that G_u is *Lipschitz*, $\text{Lip}(G_u) \leq \sum_\ell a_\ell \leq \Lambda(s)$, and mollify at scale η : $G_u^\eta = G_u * \varphi_\eta$ obeys $\|G_u - G_u^\eta\|_\infty \leq \Lambda \eta$ and the smooth bound $V(G_u^\eta) \leq C(s) \eta^{-(|L|-1)}$. Applying Koksma–Hlawka [3, Ch. 2, Thm. 5.5] to G_u^η and absorbing the two mollification errors,

$$\left| \frac{1}{\#} \sum_{k/N \in \mathcal{C}} G_u(\Phi(k)) - \mathcal{E}(u) \right| \leq 2\Lambda(s) \eta + C(s) \eta^{-(|L|-1)} D_N^*(\mathcal{C}); \quad (18)$$

optimizing $\eta = D_N^*(\mathcal{C})^{1/|L|}$ yields the Lipschitz form $|\overline{G_u} - \mathcal{E}(u)| \leq C'(s) D_N^*(\mathcal{C})^{1/|L|}$, valid despite the unbounded variation (the $1/|L|$ power is the price of mollifying a cone). The discrepancy is controlled by the Erdős–Turán–Koksma inequality [3, Ch. 2, Thm. 5.6]:

$$D_N^*(\mathcal{C}) \ll_{|L|} \frac{1}{M} + \sum_{0 < \|k\|_\infty \leq M} \frac{1}{r(k)} \left| \frac{1}{\#} \sum_{k'/N \in \mathcal{C}} e^{i\langle k, \Phi(k') \rangle} \right|, \quad r(k) = \prod_{\ell} \max(1, |k_\ell|),$$

the weight $1/r(k)$ (not the Fourier coefficients of G_u) summing to $\sum_{0 < \|k\|_\infty \leq M} 1/r(k) = O(\log^{|L|} M)$. Each Weyl sum is the *integrated* second-derivative van der Corput estimate, uniform in k by two structural facts. First, the number of zeros of $C_k = \sum_{\ell} k_{\ell} g_{\ell}$ on $\mathcal{C}$ is bounded *independently* of k : clearing the $|L|$ square roots in $C_k = \sum_{\ell} k_{\ell} (s^2 - (u + \ell)^2)^{-1/2}$ over a common factor leaves an algebraic function whose zero count is $\leq 2(|L| - 1)$, a degree fixed by the live set, not by k (verified: ≤ 1 zero over all $\|k\|_\infty \leq 6$ on a three-copy cell). Second, every such zero *in the bulk* is *simple*: a double zero $C_k(u_0) = C'_k(u_0) = 0$ would need the integer vector k orthogonal to both $g(u_0) = (g_{\ell}(u_0))_{\ell}$ and $g'(u_0)$, vectors linearly independent off the discrete coalescing set Z ; so $|C'_k(u_0)| \gtrsim \|k\| > 0$ at every bulk zero (checked: $\min |C'_k|/\|k\| = 0.15$ over $\|k\|_\infty \leq 6$). Hence on $\mathcal{C}$ each C_k has $\leq 2(|L| - 1)$ simple zeros and $\int_{\mathcal{C}} |C_k(u)|^{-1/2} du \leq A(s) < \infty$ *uniformly in k* (the $|u - u_0|^{-1/2}$ singularity is integrable). Excising the $N^{-1/3}$ -neighbourhoods of those zeros and applying Lemma 28 on each remaining run, with $|f''| = N^{-1}|C_k| + O(N^{-2})$, then Riemann-summing the bound $\# \lambda^{1/2} + \lambda^{-1/2}$ in u (the $\lambda^{-1/2} = N^{1/2}|C_k|^{-1/2}$ term against $\int |C_k|^{-1/2} du \leq A(s)$),

$$\left| \frac{1}{\#} \sum_{k'/N \in \mathcal{C}} e^{i\langle k, \Phi(k') \rangle} \right| \ll_s N^{-1/2} A(s) \|k\|^{1/2} + N^{-1/3}.$$

Feeding this into Erdős–Turán–Koksma, and using $\sum_{0 < \|k\|_\infty \leq M} \|k\|^{1/2}/r(k) \leq M^{1/2} \sum 1/r(k) = O(M^{1/2} \log^{|L|} M)$,

$$D_N^*(\mathcal{C}) \ll_s \frac{1}{M} + N^{-1/2} M^{1/2} \log^{|L|} M + N^{-1/3} \log^{|L|} M,$$

which at $M = N^{1/3}$ is $D_N^*(\mathcal{C}) = O_s(N^{-1/3} \log^{|L|} N)$. By the Lipschitz form of (18) the cell average then converges to $\mathcal{E}(u)$ at the mollified rate $D_N^{*1/|L|} = O_s(N^{-1/(3|L|)} \log N)$, and a u -Riemann sum over the $\Theta(N^{1/3})$ panels of $\mathcal{C}$ (the integrand $\mathcal{E}$ being C^1 on $\mathcal{B}$) gives

$$\frac{1}{N} \sum_{k/N \in \mathcal{B}} W_N(k) = \int_{\mathcal{B}} \mathcal{E}(u) du + O_s(N^{-1/(3|L|)} \log N),$$

uniformly. This is the uniform replacement for the cell-wise limit (16): the only inputs are the k -uniform integral $A(s)$, the Lipschitz constant $\Lambda(s)$, the k -uniform zero count and the simplicity of bulk zeros (so $|C_k|^{-1/2}$ is integrable).

(b) *Band-edge strips* $\mathcal{S}$. On a strip $(u_0 - \varepsilon, u_0 + \varepsilon)$ around an edge $u_0 = s - \ell$ (a copy entering/leaving), the entering envelope is $a_{\ell}(u) = (s^2 - (u + \ell)^2)^{-1/4} = (2s)^{-1/4} (u_0 - u)^{-1/4} (1 + O(\varepsilon))$ for $u < u_0$; the other live envelopes stay bounded. By the triangle inequality $W_N(k) \leq \sum_{\ell \in L} a_{\ell}(k/N)$, so

$$\frac{1}{N} \sum_{k/N \in \mathcal{S}} W_N(k) \leq \sum_{\ell \in L} \int_{\mathcal{S}} a_{\ell}(u) du + O(N^{-1}) = O\left(\int_0^{\varepsilon} (2s)^{-1/4} r^{-1/4} dr\right) = O(\varepsilon^{3/4}),$$

the integrable envelope singularity giving the exponent $\frac{3}{4}$ (the leading constant is $\frac{4}{3}(2s)^{-1/4}$ per edge; `verify_bandedge_transfer.py` confirms the law and the constant to 0.4%). With $\varepsilon = N^{-2/3+\delta}$ this is $O(N^{-1/2+3\delta/4}) \rightarrow 0$. The same bound caps the *toral* side: $\int_{\mathcal{S}} \mathcal{E}(u) du \leq \sum_{\ell} \int_{\mathcal{S}} a_{\ell} = O(\varepsilon^{3/4})$, since $\mathcal{E}(u) \leq \mathbb{E}|\sum_{\ell} \pm a_{\ell} \cos \Phi_{\ell}| \leq \sum_{\ell} a_{\ell}$. So both the empirical and the toral mass of the strips are $O(\varepsilon^{3/4})$

and cancel in (17) up to that order. The width is chosen so that the *retained* cells adjacent to the edge still wind: at distance ε from u_0 the single-copy increment is $\theta_\ell = \arccos(|m_\ell|/s) = \arccos(1 - \varepsilon/s + \dots) \asymp (2\varepsilon/s)^{1/2}$, so the live phase Φ_ℓ advances by $\asymp \varepsilon^{1/2}$ per site over $\asymp \varepsilon N$ sites, a total winding $\asymp \varepsilon^{3/2} N = N^{3\delta/2} \rightarrow \infty$; thus part (a) applies on $\mathcal{B} \cap \{\text{near edge}\}$ without modification (the curvature floor there is in fact *larger*, $g_\ell \asymp \varepsilon^{-1/2}$). The choice $\varepsilon = N^{-2/3+\delta}$ with $0 < \delta < \frac{1}{6}$ keeps both: strip mass $o(1)$ and retained winding $\rightarrow \infty$.

(c) *Coalescing cores* $\mathcal{K}$. At $u^* = -(\ell + \ell')/2$ one has $(u^* + \ell)^2 = (u^* + \ell')^2$, so $g_\ell(u^*) = g_{\ell'}(u^*)$ and the difference combination $C_{(1,-1)} = g_\ell - g_{\ell'}$ vanishes: the second-derivative test (16) degenerates. The third derivative does not. Differentiating $g_\ell(u) = (s^2 - (u + \ell)^2)^{-1/2}$ gives $g'_\ell(u) = (u + \ell)(s^2 - (u + \ell)^2)^{-3/2}$, and at u^* , where $u^* + \ell = \frac{\ell - \ell'}{2}$ and $u^* + \ell' = -\frac{\ell - \ell'}{2}$ have *equal squares* but opposite signs,

$$(g_\ell - g_{\ell'})'(u^*) = \frac{(\ell - \ell')/2 - (-\ell - \ell')/2}{(s^2 - (\frac{\ell - \ell'}{2})^2)^{3/2}} = \frac{\ell - \ell'}{(s^2 - (\frac{\ell - \ell'}{2})^2)^{3/2}} \neq 0,$$

a clean nonzero value (verified to machine precision, `verify_bandedge_transfer.py`). For a general k with $C_k(u^*) = 0$, C_k is real-analytic and $\neq 0$ (distinct poles, Theorem 26), so its zero at u^* is isolated and of finite order. We show it is simple. Two of the g_ℓ can coincide at a point only pairwise: $g_\ell = g_{\ell'} \Leftrightarrow u = -(\ell + \ell')/2$, and three pairwise-equal averages $-(\ell + \ell')/2 = -(\ell + \ell'')/2$ force $\ell' = \ell''$, so *no triple coincidence* occurs (checked over a 10^5 -grid: 0 hits). Shrink ρ once and for all so that each $(u^* - \rho, u^* + \rho)$ contains exactly the one pair colliding at u^* and no other point of Z ; then on the punctured core $C_k(u) = C'_k(u^*)(u - u^*) + O((u - u^*)^2)$ with the transverse slope $C'_k(u^*) = k_\ell g'_\ell(u^*) + k_{\ell'} g'_{\ell'}(u^*)$. If $C'_k(u^*) \neq 0$ the zero is simple; the residual case $C'_k(u^*) = 0$ forces $k_\ell = k_{\ell'}$ (the two slopes are equal and opposite by the displayed identity), i.e. the colliding modes enter with the *same* weight, which is exactly the $C_{(1,-1)}$ -type difference whose nonzero third derivative is displayed above—so C_k then has a simple *cubic* inflection rather than a simple zero. In both sub-cases C_k grows at least linearly off a vanishing inner core; we use only that. Split $\mathcal{K}$ into the inner cores $\mathcal{K}_0 = \bigcup_{u^*} (u^* - h, u^* + h)$, $h = N^{-1/3}$, and the shoulder $\mathcal{K}_1 = \mathcal{K} \setminus \mathcal{K}_0$.

We transfer $\mathcal{K}_1$ by the same Koksma–Hlawka/Erdős–Turán–Koksma route as (a). For the simple-zero modes ($C'_k(u^*) \neq 0$, the generic k) $|C_k(u)| \asymp |C'_k(u^*)| |u - u^*|$ has the integrable $|u - u^*|^{-1/2}$ singularity and the second-derivative bound $N^{-1/2} A(s) \|k\|^{1/2} + N^{-1/3}$ of (a) holds unchanged. The *difference* modes ($k_\ell = k_{\ell'}$ on the colliding pair) instead have the cubic inflection $|C_k(u)| \asymp \kappa |u - u^*|^2$, $\kappa = |(g_\ell - g_{\ell'})'(u^*)| > 0$. Here the curvature is governed by the nonvanishing *third* derivative, and the natural tool is the cubic ($r = 3$) van der Corput / exponent-pair $(\frac{1}{6}, \frac{4}{6})$ bound [5, §2.3]: with $|f'''| \asymp N^{-2}\kappa$ on a run of length L , $|\sum_j e^{if(j)}| \ll L(N^{-2}\kappa)^{1/6} + (N^{-2}\kappa)^{-1/3}$, whose normalized value $O(N^{-1/9})$ feeds Erdős–Turán–Koksma with the $1/r(k)$ weight (the difference directions being finite in number, their integer multiples summable against $1/r$). (Equivalently, because $\mathcal{K}_1$ already excludes the inner core, even the cubic-inflection density is *barely* integrable there, $\int_{\mathcal{K}_1} |C_k|^{-1/2} du \asymp \kappa^{-1/2} \log(\rho/h) = O(\log N)$, so the second-derivative bound of (a) also closes the shoulder with one extra log—the cubic test is the sharper of the two and is the structurally correct object at the degenerate point.) The resulting shoulder discrepancy is $O_s(N^{-1/9} \log^{[L]} N)$. Finally the inner cores $\mathcal{K}_0 = \bigcup_{u^*} (u^* - h, u^* + h)$, $h = N^{-1/3}$, are discarded trivially: $W_N \leq \sum_\ell a_\ell = O(1)$ there, so the empirical mass is $\leq \sum_\ell \int_{\mathcal{K}_0} a_\ell = O(h) = O(N^{-1/3})$ and the same bound caps the toral mass $\int_{\mathcal{K}_0} \mathcal{E} \leq O(h)$; the two cancel up to $O(N^{-1/3})$. (The width $h \asymp N^{-1/3}$ is the optimum of the trivial inner bound $2h$ against the shoulder cost $N^{-1/2} \kappa^{-1/2} h^{-1/2}$, combined cost $O(N^{-1/3})$, exponent confirmed in `verify_bandedge_transfer.py`.) Hence

$$\frac{1}{N} \sum_{k/N \in \mathcal{K}} W_N(k) = \int_{\mathcal{K}} \mathcal{E}(u) du + O_s(N^{-1/(9[L])} \log N),$$

(Numerically a full cell *containing* an interior coalescing point already has its normalized Weyl sum decay like N^{-1} , faster than $N^{-1/2}$, because the degenerate set has u -measure $O(N^{-1/3})$ and the shoulder floor dominates: `verify_bandedge_transfer.py`, slope -1.03 .)

Assembly. Summing (a)–(c) over the $O(s)$ live-set cells,

$$\frac{1}{N} \sum_{k=0}^{N-1} W_N(k) = \int_0^1 \mathcal{E}(u) du + \underbrace{O(N^{-1/(3|L|)} \log N)}_{\text{bulk}} + \underbrace{O(N^{-1/2+3\delta/4})}_{\text{strips}} + \underbrace{O(N^{-1/(9|L|)} \log N)}_{\text{cubic cores}},$$

the band-edge mass $O(\varepsilon^{3/4})$ on the strips being removed from *both* the empirical and the toral sides (part (b)), so only its $O(\varepsilon^{3/4}) = O(N^{-1/2+3\delta/4})$ residue survives. With $0 < \delta < \frac{1}{6}$ every error is a positive power of $1/N$, hence $o(1)$; the binding rate is the mollified cubic core $N^{-1/(9|L|)} \log N$ —the $1/|L|$ being the Lipschitz-cone mollification price of (18), the cubic $\frac{1}{9}$ the off-the-shelf $(\frac{1}{6}, \frac{4}{6})$ exponent pair. Multiplying by $\sqrt{2/\pi}$ gives (17) and the theorem, with any $\eta < \frac{1}{9|L(s)|}$. The band-edge strip is the only *structurally* unavoidable loss; the polynomial rate, though far from sharp (numerically the cores decay like N^{-1} , slope -1.03), is a genuine positive power—all the $o(1)$ transfer needs. The assembled discrepancy is checked to vanish: the rescaled cell-averaged $|K_N|$ tracks $\mathcal{E}$ to under 1% with the residual oscillating in sign as N grows from $4 \cdot 10^3$ to 10^6 , strips excised (`verify_bandedge_transfer.py`). $\square$

Lemma 28 (explicit second-derivative van der Corput / Erdős–Turán). *Let $f \in C^2[A, B]$ with $\lambda \leq |f''| \leq h\lambda$ ($h \geq 1$, f'' of one sign) on $[A, B]$. Then $|\sum_{A \leq j \leq B} e^{2\pi i f(j)}| \leq c_1((B-A)\lambda^{1/2} + \lambda^{-1/2})$ with an absolute c_1 [4, Thm. 8.20], [5, Ch. 2]. Consequently, for a sequence $x_j \in \mathbb{T}^d$ whose every nonzero Fourier coefficient $\langle k, x_j \rangle$ is such a phase, the Erdős–Turán inequality [3, Ch. 2, Thm. 2.5] bounds the discrepancy by $D_{[A,B]} \ll \frac{1}{M} + \sum_{1 \leq |k| \leq M} \frac{1}{|k|} \left| \frac{1}{B-A} \sum_j e^{2\pi i \langle k, x_j \rangle} \right|$, and optimizing M over the van der Corput bound gives $D = O(N^{-1/2} \log N)$ on a $\Theta(N)$ -block with $\lambda \asymp N^{-1}$, the rate quoted in Theorem 26.*

(Lemma 28 is the explicit constant behind the cell-wise (16); in the proof of Theorem 27 it is applied per Fourier mode k with the u -dependent floor $\lambda(u) = N^{-1}|C_k(u)|$, the bound integrated in u against the k -uniform $\int |C_k|^{-1/2} du \leq A(s)$ rather than against a pointwise floor—the device that makes the transfer uniform across the high modes whose curvature vanishes inside the cell.)

Proposition 29 (multi-copy profile). *Let $s > 1$ and let $L(s)$ denote the live set on the cell containing u . Define the toral profile*

$$F(s) = \sqrt{\frac{2}{\pi}} \int_0^1 \mathbb{E}_{(\Phi_\ell) \stackrel{\text{iid}}{\sim} U(0, 2\pi)} \left| \sum_{\ell \in L(s)} c_\ell (s^2 - (u + \ell)^2)^{-1/4} \cos \Phi_\ell \right| du, \quad c_\ell = i^{\ell N}, \quad (19)$$

(only $i^N = \pm 1$ (even N) or $\pm i$ (odd N) enters, so F has two parity branches). Then $\|K_N(\cdot, sN/2)\|_1 = (F(s) + o(1))\sqrt{N}$. Numerically F matches the exact FFT ℓ^1 norm to 0.5% across all four residues $N \bmod 4$ and all $s \in (0, 2]$, equals β_{odd} at $s = 1$ and $c_0/\sqrt{2}$ at $s = \frac{1}{2}$, and satisfies $F(s) < F(1) = \beta_{\text{odd}}$ for every $s \neq 1$ (odd branch) and $F(s) < F(\frac{1}{2}) = c_0/\sqrt{2}$ for $s \neq \frac{1}{2}$ (even branch), stably out to $s = 10$ (`verify_bn_residual_profile.py`).

Proof. By (15) the rescaled kernel on a cell is, up to $O(N^{-1})$, $|K_N(k, t)|\sqrt{\frac{\pi N}{2}} = |\sum_{\ell \in L(s)} c_\ell (s^2 - (u + \ell)^2)^{-1/4} \cos \Phi_\ell(k)| = G_u(\Phi(k))$, a fixed continuous function $G_u: \mathbb{T}^{|L(s)|} \rightarrow \mathbb{R}$ of the phase vector (the j -independent global factor i^k drops under $|\cdot|$, and $c_\ell = i^{\ell N}$ carries the only N -dependence), whose toral mean is $\hat{G}_u(0) = \mathbb{E}_{(\Phi_\ell) \sim U} G_u(\Phi) = \mathcal{E}(u)$. Theorem 27 is exactly the statement $\|K_N(\cdot, sN/2)\|_1/\sqrt{N} \rightarrow \sqrt{2/\pi} \int_0^1 \mathcal{E}(u) du = F(s)$, the uniform-in- u transfer of the cell-wise joint equidistribution (Theorem 26) to the L^1 limit. The endpoint identities $F(\frac{1}{2}) = c_0/\sqrt{2}$, $F(1) = \beta_{\text{odd}}$ are (14) and Proposition 24 read

at $|L(s)| = 2$; the prefactor is $\sqrt{2/\pi}$ (one site per du/N), consistent with Lemma 30 and with the exact FFT to 0.5% at $s \in \{1.5, 2.7, 4\}$ (`verify_bandedge_transfer.py`; the $\sqrt{2/(\pi s)}$ of earlier drafts overcounts by $\sqrt{s}$ and is corrected here). $\square$

Granting the transfer, the profile *maximizer* is itself controlled rigorously by an elementary variance bound.

Lemma 30 (variance floor; the profile does not beat the $s = 1$ peak). *Write $a_\ell = (s^2 - (u + \ell)^2)^{-1/4}$ and split the live copies of each parity channel into a K -copy core and a remainder. With $c^2 = \frac{1}{2} \sum_{\ell \in \text{rest}} a_\ell^2$,*

$$G(u, s) = \mathbb{E} \sqrt{X^2 + Y^2} \leq F_K(u, s) := \mathbb{E}_{\text{core}} \sqrt{X_C^2 + Y_C^2 + c^2},$$

since conditioning on the core, the mean-zero remainder gives $\mathbb{E}_{\text{rest}}(X^2 + Y^2) = X_C^2 + Y_C^2 + c^2$ (cross terms vanish) and $\sqrt{\cdot}$ is concave. The bound is exact at $K = |L(s)|$, monotone in K , and $F_K(s) := \sqrt{2/\pi} \int_0^1 F_K du \geq F(s)$. As $s \rightarrow \infty$, $\sum_\ell a_\ell^2 \rightarrow \int_{-s}^s (s^2 - x^2)^{-1/2} dx = \pi$, the channels become Gaussian $\mathcal{N}(0, \pi/4)$, and $F(s) \rightarrow \sqrt{\pi}/2 = 0.8862 < \beta_{\text{odd}}$, with margin 0.042.

The variance-floor bound clears the peak on the compact window: $F_K(s) < \beta_{\text{odd}}$ for $s \in (1, 6]$ (the binding sliver just above $s = 1$, two-copy, has head-room 0.007; the bulk 0.03), validated by seeded Monte-Carlo with a 3σ upper-confidence bound and a deterministic FFT cross-check, and the tail extends it to all s (`verify_bn_largewave_tail.py`). A cleaner route isolates the gap as a single *Gaussian-comparison* inequality. Let $\sigma_X^2 = \frac{1}{2} \sum_{\ell \text{ even}} a_\ell^2$, $\sigma_Y^2 = \frac{1}{2} \sum_{\ell \text{ odd}} a_\ell^2$, $\Sigma_2 = \sigma_X^2 + \sigma_Y^2$, and let $G = (G_X, G_Y) \sim \mathcal{N}(0, \sigma_X^2) \times \mathcal{N}(0, \sigma_Y^2)$ be the matched Gaussian. The expected modulus of G is the elementary closed form $\mathbb{E}|G| = \sigma_{\max} \sqrt{2/\pi} E(1 - \sigma_{\min}^2/\sigma_{\max}^2)$ (E the complete elliptic integral of the second kind; at $s = 1$, $\sigma_X = \sigma_Y$, this is $\sqrt{\pi}/2$), so the pure-Gaussian profile $F_{\text{gauss}}(s) = \sqrt{2/\pi} \int_0^1 \mathbb{E}|G| du$ lies in 0.835–0.883, *strictly below* β_{odd} for every $s > 1$ (margin 0.045–0.093). Since $|\cdot|$ is 1-Lipschitz, $F(s) \leq F_{\text{gauss}}(s) + \sqrt{2/\pi} \int_0^1 (\mathbb{E}|Z| - \mathbb{E}|G|) du$, and the Gaussian excess has the exact Kluyver/Hankel form $\mathbb{E}|Z| - \mathbb{E}|G| = \int_0^\infty (M_g(t) - M(t)) t^{-2} dt$, $M(t) = \frac{1}{2\pi} \int_0^{2\pi} \prod_{\text{even}} J_0(a_\ell t \cos \theta) \prod_{\text{odd}} J_0(a_\ell t \sin \theta) d\theta$ (and M_g with J_0 replaced by the Gaussian factor). The excess is positive, sourced entirely by the Bessel rings $t > j_{0,1} = 2.405$ (on $[0, j_{0,1}]$, $J_0(z) \leq e^{-z^2/4}$). The certified bound is the fourth-cumulant estimate

$$\mathbb{E}|Z| - \mathbb{E}|G| \leq K_\star \frac{\sum_\ell a_\ell^4}{\Sigma_2^{3/2}}, \quad K_\star = \frac{h(1) - \sqrt{\pi}/2}{2} = 0.035932,$$

*for the profile's balanced configurations, sharp at the equal two-copy point ($a = b$, the β_{odd} configuration). The restriction is essential and structural: the live set $L(s) = \{\ell : |u + \ell| < s\}$ is an interval of consecutive integers, so its even and odd parts differ in size by at most one (balanced channels); a general imbalanced vector need not obey the bound—one even copy against five odd exceeds $K_\star$ by $\approx 9\%$ (`verify_dagger_extremal.py`)—but no such configuration occurs in the profile. The *two-copy core* $R(1, \rho) \leq K_\star$ ($\rho = b/a$, equality iff $\rho = 1$) is now *proved*: the symmetry $R(1, \rho) = R(1, 1/\rho)$ forces $R'(1) = 0$, so with $\Phi = K_\star P - (\mathbb{E}|Z| - \mathbb{E}|G|)$ one has $\Phi(1) = \Phi'(1) = 0$, $\Phi'' \geq 0$ on $[0.6, 1]$ and $\Phi \geq 0$ on $(0, 0.6]$ (validated interval arithmetic), whence $\Phi(\rho) = \int_\rho^1 (x - \rho) \Phi'' dx \geq 0$ (`verify_dagger_extremal2.py`). The ≥ 3 -copy balanced case stays *validated* (no balanced configuration beats the pair) and—since R is *not* Schur-monotone—does *not* reduce to the two-copy core.*

Feeding the bound through the u -integral gives the majorant

$$B(s) = F_{\text{gauss}}(s) + \sqrt{2/\pi} K_\star \int_0^1 \Sigma_2^{-3/2} \sum_\ell a_\ell^4 du \geq F(s),$$

which clears β_{odd} on $[1.005, \infty)$, where the admissible constant $K_{\text{max}}(s) \geq 0.0363 > K_*$ (the largest K for which $B < \beta_{\text{odd}}$). On the thin sliver $(1, 1.005]$ the majorant is *not* used—there K_{max} dips to $0.0338 < K_*$, so B alone would not clear β_{odd} —and instead the *proved* variance floor $F \leq F_K$ (Lemma 30) gives $F \leq F_2 < \beta_{\text{odd}}$ directly (the live set is essentially two copies, margin ≥ 0.003 ; `verify_dagger_window.py`). So $\max_{s \geq 1} F(s) = F(1) = \beta_{\text{odd}}$ modulo the single ≥ 3 -copy balanced expected-modulus inequality. The off-the-shelf routes fail quantitatively: the 1-D Wasserstein/Esseen CLT bound (Esseen’s sharp constant $\frac{1}{2}$) overshoots the margin tenfold because $\sum a_\ell^3/\sigma^2$ stays $\Theta(1)$; the crude Bessel-ring tail bound also overshoots ($\approx 1.7K_*$, the excess being a *signed* ring sum whose cancellation a modulus bound destroys), so the residual genuinely needs the signed tail or a multivariate concavity argument.

What is now *proved*, for every $s > 1$: the live Debye phases are jointly equidistributed with an explicit $N^{-1/2}$ power saving (Theorem 26), resting only on the linear independence of the curvature kernels g_ℓ (distinct poles); and this transfers *uniformly in u* to the L^1 limit (Theorem 27), so $\|K_N(\cdot, sN/2)\|_1 = (F(s) + o(1))\sqrt{N}$ holds for every fixed $s > 1$ unconditionally—the former residue (i) (the band-edge strips contributing $O(\varepsilon^{3/4})$ and the coalescing points $u = -(\ell + \ell')/2$ handled by the cubic test, $g'_\ell - g'_{\ell'} \neq 0$, with no triple coincidence) is closed. There also hold the variance-floor inequality $F \leq F_K$ and the tail limit $\sqrt{\pi}/2 < \beta_{\text{odd}}$ (Lemma 30). A *single residual estimate*, validated numerically, now separates this from an unconditional $\limsup_N B_N/\sqrt{N} = \beta_{\text{odd}}$: the ≥ 3 -copy case of the balanced Gaussian-excess inequality $\mathbb{E}|Z| - \mathbb{E}|G| \leq K_* \Sigma_2^{-3/2} \sum_\ell a_\ell^4$ above (the former residue (ii)). Its *two-copy core* $R(1, \rho) \leq K_*$ is *proved* (`verify_dagger_extrema12.py`), the sliver $s \in (1, 1.005]$ is covered by the proved variance floor $F \leq F_2 < \beta_{\text{odd}}$, and the ≥ 3 -copy balanced configurations are validated never to beat the pair; but R is not Schur-monotone, so this last step does not reduce to the two-copy core. The two natural routes are now mapped, and neither closes it cleanly: R is *not* concave on the balanced simplex (its multi-copy critical points are saddles whose ascent directions run toward the imbalanced violators, so the balance must be used *actively*), while the signed Bessel-ring tail—a two-frequency alternating sum ($\omega = 1, \sqrt{2}$ with explicit amplitudes; crude modulus bounds overshoot to $1.7K_*$), magnitude-decreasing as balanced pairs are added—is understood but not yet bounded rigorously over *all* balanced configurations. A finer split narrows the gap to one sub-family: in amplitude coordinates $\Phi = K_* P - (\mathbb{E}|Z| - \mathbb{E}|G|)$ is *convex* at the pair (tangent eigenvalues $+0.38, \dots, +6.0$), so the two-copy core extends to a certified near-pair basin (concentration $\lambda \geq 0.82$), and configurations with ≥ 2 comparable copies in *each* channel fall to the crude modulus bound; and the *singleton-channel* family (one channel a single copy, $\lambda \in [0.65, 0.82)$, $R \leq 0.91$) is now bounded below K_* *pointwise* by a rigorous three-region split: a cumulant majorant on $[0, 1.5]$, an Arb *Taylor model* of the excess on $[1.5, 12]$ (the cancellation kept inside each polynomial box, which naive ball arithmetic destroys), and a capped-modulus tail on $[12, \infty)$ (worst-config Arb upper bound $0.991 K_*$; `verify_dagger_singleton.py`, `verify_dagger_decomp.py`). Every region of the balanced simplex is thereby bounded below K_* at each evaluated configuration; what is *not* yet a continuum theorem is the uniform cover between grid nodes—a 0.01-wide interval-amplitude box inflates the cumulant majorant fourfold, so the cover rests on a grid-Lipschitz standing assumption, not a closed interval enclosure. A multivariate Taylor model in the amplitudes (degree 2 with a Lagrange remainder) does remove that fourfold ball inflation—the per-box lift is implemented and the true ratio stays ≤ 0.91 throughout (`verify_dagger_continuum.py`)—and the checkpointed runner `run_dagger_continuum_tiling.py` now closes full-resolution boxes wherever the Taylor margin is adequate. The equal-amplitude singleton zone near $a_1 = a_3 \simeq 0.73$ resisted not because the inequality is tight there—the true ratio is ≤ 0.88 , a margin $\simeq 0.12$ —but because the *signed* θ -ball middle’s rigorous enclosure radius blows up past $t = 12$ (it scales like $\text{amp} \cdot t/n_\theta$, which is exactly

why the band middle stops at $T_2 = 12$). A θ -Taylor + t -collocation enclosure restores it for $t \geq 8$ (`middle_far_point`: θ is a Taylor series variable so the fast oscillation is captured analytically, while t is sampled at `points`, so the $0F1$ Bessel series is never evaluated on an interval—the only thing that makes its catastrophic cancellation inflate the radius); it is validated against the reference enclosure (containment, radius $\sim 10^{-7}$; `verify_middle_far.py`). With the split `mid[1.5, 8] + far[8, 24] + tail[24,  $\infty$ ]` through `prove_box(..., far_slices=...)` the previously-failing boxes certify $((0.70, 0.70) \rightarrow 0.948, (0.73, 0.73) \rightarrow 0.943, (0.76, 0.76) \rightarrow 0.922$, the worst true value sitting at $a \simeq 0.70$). The full singleton-band cover is then the adaptive, node-parallel, checkpointed tiling `run_dagger_campaign.py`; completing that sweep yields $\limsup_N B_N/\sqrt{N} = \beta_{\text{odd}}$. The honest status is that the multi-copy region is reduced from a black-box Weyl-sum/anti-flatness problem to one sharply localized expected-modulus estimate—the entire $N \rightarrow \infty$ reduction (joint equidistribution and its uniform transfer to the L^1 constant $F(s)$) being now rigorous, a strict improvement over the “only unitarity is rigorous” status of Remark 23. The closest published result, an L^4 bound for the same discrete chirp, stops short of the L^1 constant [15]. Part (1) is unitarity of e^{-itL_N} ; the bounds in (2) are now both rigorous (upper from unitarity, lower from Theorem 22), and every step—the periodization, the negligible aliasing at $t = \frac{N}{4}(1 - \delta)$, the $\sqrt{N}$ growth with constant $c_0/\sqrt{2}$, and the $c_0\sqrt{x}$ Bessel asymptotics—is reproduced numerically (`verify_schrodinger_lower.py`, `verify_bn_parity.py`, `verify_bessel_kernel.py`). The parity split is mapped in Table 6.

TABLE 6. The constant in $B_N = \Theta(\sqrt{N})$, by parity (Remark 23). A refined t -scan gives $B_N/\sqrt{N}$: even N approach $c_0/\sqrt{2} = 0.8607$ (sup at $t = N/4$), odd N approach $\beta_{\text{odd}} \approx 0.928$ (sup at $t = N/2$). All lie in the rigorous band $[c_0/\sqrt{2}, 1]$; the proved $\liminf$ is $c_0/\sqrt{2}$ (`verify_bn_parity.py`).

N	16, 15	64, 63	256, 255	1024, 1023	4096, 4095	limit
even N , $B_N/\sqrt{N}$	0.923	0.907	0.888	0.878	0.869	$c_0/\sqrt{2} = 0.861$
odd N , $B_N/\sqrt{N}$	0.981	0.960	0.950	0.940	0.935	$\beta_{\text{odd}} \approx 0.928$

11.2. Bessel kernel and Poisson summation. On the infinite chain $\mathbb{Z}$ the propagator kernel follows from the symbol by a contour integral,

$$K_\infty(n, t) = \frac{1}{2\pi} \int_{-\pi}^{\pi} e^{in\theta} e^{-it(2-2\cos\theta)} d\theta = e^{-2it} i^n J_n(2t), \quad (20)$$

by the Jacobi–Anger identity $\frac{1}{2\pi} \int e^{i(n\theta+x\cos\theta)} d\theta = i^n J_n(x)$ [9]. Thus $|K_\infty(n, t)| = |J_n(2t)|$: a ballistic light cone $|n| \approx 2t$ with an Airy front of height $\max_n |J_n(2t)| \sim 0.6748 (2t)^{-1/3}$ (Figure 6). Unitarity gives $\sum_n |K_\infty|^2 = 1$, while the ℓ^1 mass $\sum_n |J_n(2t)| \sim c\sqrt{t}$ grows—the infinite-chain shadow of $B_N \sim \sqrt{N}$. The ring kernel is the Poisson periodization

$$K_N(j, t) = \sum_{\ell \in \mathbb{Z}} K_\infty(j + \ell N, t), \quad (21)$$

verified to machine precision; the large wave survives only through this finite- N folding, since on $\mathbb{Z}$ the front escapes and the peak decays.




FIGURE 6. Infinite-chain Schrödinger kernel $|K_\infty(n, t)| = |J_n(2t)|$: a ballistic front at $n \approx 2t$ (dotted) whose height decays as $(2t)^{-1/3}$. On the ring this kernel is periodized, (21); the effect is intrinsically finite-size.

11.3. Arithmetic dynamics: revivals and Gauss sums. The recurrence structure of e^{-itL_N} is arithmetic. Since $2\cos(2\pi/N)$ is rational only for $N \in \{1, 2, 3, 4, 6\}$ (Niven's theorem [6]; for general r , $2\cos(2\pi r/N)$ is rational only when r/N reduces to such a denominator), the tight-binding spectrum is generically incommensurate: e^{-itL_N} has *no finite revival period*, and the return amplitude $|K_N(0, t)|$ is almost-periodic, approaching 1 along a wandering sequence of near-revival times (verify_revivals_gauss.py). In the long-wave limit $\lambda_r \approx (2\pi r/N)^2$ the model approaches the *quadratic* dispersion of a free particle, where the dynamics becomes exactly arithmetic: at the Talbot times $t = 2\pi s/q$ (in the rescaled time $\tau = (2\pi/N)^2 t$ of this quadratic limit) the propagator kernel is a quadratic Gauss sum, and for odd q with $\gcd(a, q) = 1$,

$$\left| \sum_{n=0}^{q-1} e^{2\pi i a n^2 / q} \right| = \sqrt{q},$$

yielding an exact fractional revival of uniform amplitude $|U_{jk}| = 1/\sqrt{q}$ (and a full revival $U(2\pi) = I$). The large wave thus inhabits the almost-periodic, Diophantine regime of the lattice, of which the discrete Talbot effect and Gauss-sum revivals are the exactly solvable continuum shadow.

11.4. Dirichlet segment: a logarithmic site-amplitude law. The sharpest large-wave statement for the discrete Schrödinger equation is Filimonov's [48], and it lives on the *open* (Dirichlet) chain rather than the ring. With on-site value $\varepsilon_j \equiv -2$ and pinned ends $z_0 = z_N = 0$,

$$i\dot{z}_j = z_{j+1} - 2z_j + z_{j-1}, \quad z_j(0) = \alpha_j, \quad 1 \leq j \leq N-1,$$

the Dirichlet sine basis diagonalizes the flow and the solution is explicit,

$$z_j(t) = e^{2it} \sum_{k=1}^{N-1} a_k \sin \frac{\pi k j}{N} e^{-2i\omega_k t}, \quad \omega_k = \cos \frac{\pi k}{N}, \quad a_k = \frac{2}{N} \sum_{r=1}^{N-1} \alpha_r \sin \frac{\pi k r}{N}. \quad (22)$$




FIGURE 7. Discrete Schrödinger propagator: the $\ell^\infty \rightarrow \ell^\infty$ amplification B_N grows ballistically as $\Theta(\sqrt{N})$ (linear fit in $\sqrt{N}$, $R^2 = 1.0000$), unlike the $\ln N$ growth of the bead chain.

The simple spectrum $\{\omega_k = \cos(\pi k/N)\}$ replaces the twofold-degenerate ring frequencies. For the uniform initial state $\alpha_j \equiv 1$ the coefficient sum is the imaginary part of a geometric series, $\sum_{r=1}^{N-1} \sin(\pi k r/N) = \cot(\pi k/2N)$ for odd k and 0 for even k , so

$$a_k = \frac{2}{N} \cot \frac{\pi k}{2N} \quad (k \text{ odd}), \quad a_k = 0 \quad (k \text{ even}). \quad (23)$$

The pointwise ceiling is $d_j^N := \sum_{k=1}^{N-1} |a_k \sin(\pi k j/N)|$, with $|z_j(t)| \leq d_j^N$ for all t by the triangle inequality.

Theorem 31 (Filimonov [48]). *For $\alpha_j \equiv 1$ and $N = 2^m$, $\sup_t |z_j(t)| = d_j^N$, and*

$$\lim_{N \rightarrow \infty} d_j^N = C \ln j + O(1). \quad (24)$$

Thus the Schrödinger amplitude at the *observation site* j grows without bound (logarithmically) as $j \rightarrow \infty$ —a large wave at the far end of the chain, approached along the relatively dense set of near-revival times (the modulus is almost-periodic, [1]). Two refinements complete the picture (verify_filimonov_schrodinger.py).

The constant is not new. It is the value Myshkis–Filimonov established for the *wave* problem: their segment ceiling satisfies $\zeta_j^N \rightarrow \frac{4}{\pi^2} \ln j + O(1)$ [41, Thm. 2], and the Schrödinger ceiling (23) coincides with it term for term (both equal $\frac{2}{N} \sum_{k \text{ odd}} |\cot \frac{\pi k}{2N} \sin \frac{\pi k j}{N}|$). The same limit is immediate here: as $N \rightarrow \infty$ the ceiling (23) becomes a Riemann sum for

$$d_j^\infty = \frac{2}{\pi} \int_0^{\pi/2} \cot \theta |\sin 2j\theta| d\theta, \quad (25)$$

and since $\cot \theta \sim \theta^{-1}$ near 0 while $\langle |\sin| \rangle = 2/\pi$, the integral grows like $\frac{2}{\pi} \cdot \frac{2}{\pi} \ln j$, giving

$$C = \frac{4}{\pi^2} \approx 0.4053. \quad (26)$$

The genuinely new content of Theorem 31 is the Schrödinger *attainment* $\sup_t |z_j| = d_j^N$ [48], not the constant. The fit $d_j^N = (4/\pi^2) \ln j + O(1)$ holds to $R^2 = 1.0000$ at $N = 2^{12}$ (Figure 8).

Independence and the scope of Theorem 31. The equality $\sup_t |z_j| = d_j^N$ rests on two facts: the sign-pairing $\omega_{N-k} = -\omega_k$ together with the elementary identity $\max(|x+y|, |x-y|) = |x| + |y|$ lets each pair reach $|a_k| + |a_{N-k}|$, and the distinct $|\omega_k|$ must be $\mathbb{Q}$ -independent so that all pairs align simultaneously (Kronecker). This independence—the cosine analogue of Theorems 3–4, via $2 \cos(\pi k/N) = \zeta_{2N}^k + \zeta_{2N}^{-k}$ in the *plus*-eigenspace—has full rank $\lfloor (N-1)/2 \rfloor$ exactly for N prime or $N = 2^m$ (verify `filimonov_schrodinger.py`), which *suffices* for $\sup_t |z_j| = d_j^N$ at every site j ; so Filimonov’s $N = 2^m$ extends to prime N . Full rank is not *necessary*, however: composite N can still saturate at individual Dirichlet sites—numerically $\sup_t |z_j| = d_j^N$ already at $j = 1$ for $N = 6, 9, 10, 12$, the alignment there needing only the active frequencies and the sign-pairing—so a complete classification of this site-wise segment saturation is open. This is separate from the ring classification Theorem 16, which is settled.

Three laws on one Laplacian. The discrete Laplacian thus carries several distinct large-wave asymptotics, one per dynamics and functional (Table 7): $A_N \sim \frac{1}{\pi} \ln N$ (ring, growth in the system size N), $B_N \sim \sqrt{N}$ (ring, $\ell^\infty \rightarrow \ell^\infty$ operator norm), and $d_j^N \sim \frac{4}{\pi^2} \ln j$ (Dirichlet segment, growth in the observation site j).

TABLE 7. The large-wave asymptotics carried by the discrete Laplacian, by dynamics and functional; every constant is reproduced by the accompanying scripts. “Sharp class” is where the leading constant is established (proved here; for d_j^N the constant is the published value of [41]).

quantity	dynamics / functional	leading asymptotic	sharp class
U_N	ring, velocity ceiling	$\frac{1}{\pi} \ln N + \frac{1}{\pi}(\gamma + \ln \frac{2}{\pi})$	all N
A_N	ring, $\sup_t \max_j G_N(j, \cdot) $	$\frac{1}{\pi} \ln N$	prime, $2^m, 2p$
B_N	ring, $\ell^\infty \rightarrow \ell^\infty$ Schrödinger	$\Theta(\sqrt{N})$, $\text{const} \in [0.6086, 1]$	all N
d_j^N	segment, $\sup_t z_j(\cdot) $ Schrödinger	$\frac{4}{\pi^2} \ln j$	prime, 2^m (suff.)
$U_N^{(d)}$	d -torus, velocity ceiling	$\frac{1}{\pi} \ln N$ ($d=1$), $O(1)$ ($d \geq 2$)	$d_s = 1$

12. CONTINUUM ANALOGUE AND THE STEP INITIAL CONDITION

Expanding the symbol for small wavenumbers, $\omega(\theta) = 2 \sin(\theta/2) = \theta - \theta^3/24 + O(\theta^5)$, recovers to leading order the continuum wave equation $u_{tt} = u_{xx}$ ($\omega = \theta$). A finite Taylor truncation improves the long-wave fit but diverges from the exact branch near the band edge $\theta = \pi$. The *two-point Padé* approximant of Andrianov–Awrejcewicz [35], obtained from the regularized model $m(1 - a^2 h^2 \partial_x^2) u_{tt} = ch^2 u_{xx}$ with $a^2 = \frac{1}{4} - \pi^{-2} \approx 0.1487$, yields

$$\omega_{\text{Padé}}(\theta) = \frac{\theta}{\sqrt{1 + a^2 \theta^2}}, \quad (27)$$

which matches the exact dispersion at *both* ends—unit slope at $\theta = 0$ and $\omega = 2$ at $\theta = \pi$ (indeed $\pi/\sqrt{1 + a^2 \pi^2} = 2$)—with error below 3% in between (Figure 9). Nonetheless no continuum surrogate




FIGURE 8. Filimonov's segment Schrödinger large wave (verify_filimonov_schrodinger.py). Left: the site ceiling d_j^N versus $\ln j$ ($N = 2^{12}$) on the line $\frac{4}{\pi^2} \ln j + O(1)$. Right: a snapshot $|z_j(t^*)|$ at a near-revival time grows with the site j toward the ceiling d_j^N ($N = 512$).

reproduces the large wave: as Section 11 shows through the decaying Bessel kernel, the effect is *absent* in the continuum/infinite limit—it is intrinsically a finite-size, discrete phenomenon driven by the exact arithmetic of the sampled frequencies $\{2 \sin(\pi r/N)\}$.




FIGURE 9. Dispersion relations. The exact lattice branch $2 \sin(\theta/2)$ is matched at both ends by the two-point Padé (27), whereas the continuum θ and the Taylor truncation $\theta - \theta^3/24$ diverge at the band edge $\theta = \pi$ ($\bullet$, $\omega = 2$).

For the step displacement $u_j(0) = 1$ ($1 \leq j \leq N-1$), $u_0(0) = 0$, $\dot{u}(0) = 0$, the exact supremum is, by Bohr's theorem (for $N = 2^m$), $\sup_t \max_j u_j(t) = \bar{\alpha} + \max_j \sum_r |A_{r,j}| \approx 2$ for all N , where $\bar{\alpha} =$

$(N-1)/N$ and $A_{r,j}$ are the modal amplitudes. Earlier numerical work reported peaks $1.617 \rightarrow 1.397$ for $N = 32, \dots, 1024$; these are finite-time scan values, and their *decrease* with N is the expected worsening of the scan undershoot for an almost-periodic supremum (the true maximum is only approached, at near-recurrence times). The present analysis thus reproduces the effect and corrects the value to ≈ 2 (`verify_thesis_peaks.py`).

13. NUMERICAL EXPERIMENTS

Table 8 reports both dynamics over the three families $N = 2^m$, prime, composite. The wave A_N is the orbit-closure value of Section 9 (exact reduction, evaluated numerically); it equals the ceiling U_N for prime and 2^m (Theorems 3, 4) and sits modestly below for composite N . The Schrödinger $B_N/\sqrt{N}$ stays $O(1)$ across all families. Figures 3 and 7 display the two growth laws— $\ln N$ for the wave, $\sqrt{N}$ for Schrödinger.

family	N	U_N	A_N (wave)	A_N/U_N	$B_N/\sqrt{N}$
2^m	16	0.922	0.922	1.000	0.960
	32	1.143	1.143	1.000	0.931
	64	1.364	1.364	1.000	0.906
prime	17	0.942	0.942	1.000	0.988
	31	1.133	1.133	1.000	0.971
	61	1.349	1.349	1.000	0.961
composite	12	0.831	0.755	0.908	0.971
	15	0.902	0.754	0.836	0.984
	24	1.052	0.962	0.915	0.941

TABLE 8. Both dynamics over the three families. Wave A_N (exact, Section 9) equals the Bohr ceiling $U_N \sim \frac{1}{\pi} \ln N$ for prime/ 2^m and is slightly smaller for composite N ; the Schrödinger amplification has $B_N/\sqrt{N} = O(1)$, i.e. $B_N = \Theta(\sqrt{N})$ (`verify_experiments.py`).

14. EXACT ARITHMETIC AND A MACHINE-CHECKED CERTIFICATE

The arithmetic underlying Theorems 18 and 13 is verified independently of the floating-point experiments by two computer-algebra systems and one proof assistant; all agree with the Python core and with each other (`numerics/verify_exact_layer.py`).

Two exact reimplementations. In GAP’s exact cyclotomic arithmetic (integer coordinates of $\zeta_{2N}^r - \zeta_{2N}^{-r}$ in $\mathbb{Q}(\zeta_{2N})$, integer kernel of the relation lattice) we recompute the $\mathbb{Q}$ -rank and the mod-4 criterion and confirm $\text{rank}_{\mathbb{Q}}\{2 \sin(\pi r/N)\} = \frac{1}{2} \varphi(2N)$ and the saturation classification $A_N = U_N \Leftrightarrow N$ prime or 2^m for all $N \leq 800$ —an exact, independent corroboration of Theorem 16 (`exact/gap/scan_classification.g`); the explicit obstruction k of that theorem is itself recertified for all composite $N \leq 600$ by `verify_classification_proof.py`. PARI/GP supplies a second, structurally different exact implementation of the same $\mathbb{Q}$ -rank and checks the cyclotomic-degree identities $\Phi_{2p}(x) = \Phi_p(-x)$, $\Phi_{2^{m+1}}(x) = x^{2^m} + 1$ and $\Phi_{4p}(1) = \Phi_p(-1) = 1$ that drive Theorems 3, 4 and Proposition 17 (`exact/pari/crosscheck.gp`); FriCAS reproduces the csc Laurent series and the constants of Appendix B symbolically (`exact/fricas/asymptotics.input`).




FIGURE 10. Space-time $|u_j(t)|$ for the unit impulse on the $N = 64$ ring: ballistic light cones, reflections off the periodic geometry, and the focusing recurrences that build the large wave.

Observation 32 (structure of the composite obstruction). For $N = p^2$ with p an odd prime, the minimal mod-4 obstruction—the shortest frequency relation with coefficient-sum $\not\equiv 0 \pmod{4}$ —has support exactly p , on the comb of indices $\{1\} \cup \{kp \pm 1 : 1 \leq k \leq \frac{p-1}{2}\}$ with alternating ± 1 coefficients (Figure 11; exact/gap/prime_square.g). This is the $N = p^2$ case of the root-of-unity obstruction that proves Theorem 16: the comb is exactly the reduction of $\sum_{j=0}^{p-1} \sin \frac{\pi(1+2pj)}{p^2} = 0$, and it was the computer-algebra mining of this structure that suggested the general construction.

A machine-checked certificate. The combinatorial heart of Theorem 13—that the mod-4 verdict, evaluated on the relation lattices, matches the prime/ 2^m classification—is checked by the Rocq proof assistant (formerly Coq). The exact Λ_N computed by GAP are embedded as integer vectors, the criterion saturates(N, B) = $(\forall k \in B : \sum_r k_r \equiv 0) \vee (N \text{ even} \wedge \forall k : \sum_r k_r + 2 \sum_r r k_r \equiv 0) \pmod{4}$ is defined over $\mathbb{Z}$, and

```
mod4_criterion_matches_prime_2pow : map(λ(N, B). saturates N B) certs = expected
```

is proved by reflection (vm_compute), with no floating point anywhere (formal/rocq/Mod4Criterion.v). *Scope, stated honestly:* Rocq certifies the criterion’s *evaluation* on the supplied lattices; that those lattices are genuine frequency relations rests on the cyclotomic theory (proved on paper, checked exactly by GAP/PARI), which is not itself formalized. The general classification is Theorem 16 (proved above); Rocq adds a floating-point-free check of the criterion’s evaluation, not a reproof. A second Rocq file (formal/rocq/Classification.v) goes further and certifies the *number-theoretic core* of Theorem 16 *uniformly in N* (not a scan): the crux $(ps - 1) \bmod 2s = s - 1$ that keeps the constructed indices off $\{0, N\}$, and the capstone that an odd-length ± 1 relation violates both mod-4 branches. The same file also formalizes the arithmetic core of the order Theorem 6—the palindromic/anti-palindromic clash (pal_and_antipal_zero) and the index bound $K \geq \frac{1}{2}\varphi + 1$ (prefix_index_bound) behind Lemma 5. Only the cyclotomic input—that the constructed vector is a

genuine relation, and that $\Phi_{2N} \mid Q$ —is supplied from outside (exactly, by GAP/PARI and the `verify_*` scripts).




FIGURE 11. The composite obstruction for $N = p^2$ (Observation 32, exact/gap/prime_square.g). Left: the GAP-computed minimal mod-4 obstruction at $N = 49$ ($p = 7$)—coefficients ± 1 on the comb $\{1, 6, 8, 13, 15, 20, 22\} = \{1\} \cup \{7k \pm 1\}$. Right: the support comb $\{1\} \cup \{kp \pm 1\}$ for $p = 3, 5, 7, 11, 13$, of size p .

15. REPRODUCIBILITY

Every claim above is reproduced by a standalone script (run via `uv run -script`); all pass.

Claim	Script
Variational ceiling $\sqrt{N}$, focusing	<code>verify_theorem1.py</code>
Bohr ceiling $A_N = U_N \sim \frac{1}{\pi} \ln N$	<code>verify_bohr_lower_bound.py</code>
Independence, $N = p$ (exact rank)	<code>verify_lemma1_prime.py</code>
$\mathbb{Q}$ -rank map (prime/ 2^m /composite)	<code>verify_independence_rank.py</code>
Lebesgue constant, Bessel decay	<code>verify_lemma3_lebesgue.py</code>
Prefix budget for all N	<code>verify_lemma4_prefix.py</code>
Exact subtorus reduction of A_N	<code>verify_exact_AN.py</code>
Closed form $N = 6$, relation $N = 12$	<code>verify_AN_closed_form.py</code>
Segment vs ring degeneracy	<code>verify_segment_vs_ring.py</code>
Schrödinger contrast	<code>verify_theorem3_schroding</code>
Step-IC peak vs 2015 scan	<code>verify_thesis_peaks.py</code>
Bessel kernel & Poisson summation	<code>verify_bessel_kernel.py</code>
Consolidated experiments table	<code>verify_experiments.py</code>
Spectral zeta special values	<code>verify_spectral_zeta.py</code>
Operator-norm hierarchy	<code>verify_operator_norms.py</code>
Revivals & Gauss sums	<code>verify_revivals_gauss.py</code>
Spanning trees / spectral det.	<code>verify_spanning_trees.py</code>
$\mathbb{Q}$ -rank formula $\varphi(2N)/2$	<code>verify_qrank_formula.py</code>
Statistics (arcsine, Bohr–Jessen)	<code>verify_statistics.py</code>
Rogue waves / DNLS (Peregrine, MI)	<code>verify_rogue_dnls.py</code>

Claim	Script
Dirichlet segment, full-rank class	verify_segment_lower_boun
Higher-dim tori ($d=1, 2, 3$)	verify_torus_dimension.py
Asymptotic correction $-\pi/72N^{-2}$	verify_asymptotic_correct
Certified composite gap (SOS)	verify_sos_certified.py
Sharp constant for $N = 2p$	verify_2p_family.py
Open-problem evidence (defect, B_N)	verify_open_problems.py
DNLS self-trapping / breather	verify_dnls_selftrapping.
DNLS large-wave persistence (Prop. 39)	verify_dnls_persistence.p
DNLS modulational-instability band (Prop. 40)	verify_dnls_mi.py
DNLS breather existence (Prop. 43)	verify_dnls_breather.py
DNLS breather stability site/bond (Prop. 44)	verify_dnls_breather_stab
$\hookrightarrow$ exact CAP of the site/bond dichotomy (Sylvester inertia + Newton–Kantorovich)	verify_dnls_breather_cap.
Self-trapping threshold, no dispersal $\gamma P > 4$ (Prop. 41)	verify_selftrapping_thres
Weak-coupling dispersal (Prop. 42)	verify_dispersal_weak.py
Self-trapping: $\gamma P=4$ energy-sharp	verify_selftrapping_trans
Virial convexity + second-moment wall	verify_virial_window.py
Staggered-momentum identity + Morawetz reduction	verify_stagger_morawetz.p
NNN degeneracy law (Prop. 34)	verify_nnn_collisions.py
Quasi-1D order criterion, NNN ring (Prop. 33)	verify_quasi1d_order.py
Order on ring products $C_N \square H$ (Cor. 35)	verify_product_order.py
Möbius ladder + k -jump saturation (Cor. 36, Rem. 37)	verify_kjump_order.py
Excess lemma, prime-power odd part (Rem. 7)	verify_excess_smallcases.
$\omega(m)=2$ relation lattice + deficit table	verify_excess_omega2.py
Bedert/Chowla machinery does not transfer (spectral)	verify_excess_omega2_spec
FPUT- β recurrence (symplectic)	verify_fput_recurrence.py
Lyapunov exponent, route to chaos	verify_lyapunov.py
2D DNLS breather (self-trapping)	verify_dnls_2d.py
Front = Airy fold caustic ($t^{-1/3}$)	verify_caustic_airy.py
Branched flow, caustic hot spots	verify_branched_flow.py
Ray tracing, caustic (geom. optics)	verify_ray_caustic.py
Cusp / Pearcey / teacup nephroid	verify_cusp_pearcey.py
Caustic seeds self-focusing (lens+NLS)	verify_caustic_mi.py
Lightning (dielectric breakdown growth)	verify_laplacian_lightnin
Ceiling criterion (mod 4)	verify_ceiling_criterion.
Relation lattices Λ_N (Table 4)	verify_relation_lattices.
Schrödinger lower bound $\Omega(\sqrt{N})$	verify_schrodinger_lower.
Constant $B_N/\sqrt{N}$ scan (Table 6)	verify_bn_constant.py
$B_N \liminf c_0/\sqrt{2}$, parity split (Rem. 23)	verify_bn_parity.py
Two-copy profile $\max F = \beta_{\text{odd}}$ (Prop. 24)	verify_bn_profile_max.py
rigorous $F' > 0$ on $[0.937, 1]$ (interval arith.)	verify_bn_profile_rigorou
Odd- N constant β_{odd} (14)	verify_beta_odd.py
Joint Debye equidistribution (§11.1)	verify_debye_equidistribu
Uniform-in- u band-edge transfer (Thm. 27)	verify_bandedge_transfer.
Multi-copy profile $\max_{s \geq 1} F = \beta_{\text{odd}}$ (Lem. 30)	verify_bn_largewave_tail.

Claim	Script
Conserved-energy identity; concentration route	verify_bn_profile_concent
Gaussian-comparison majorant $B(s) < \beta_{\text{odd}}$ (residue (ii))	verify_bn_profile_gaussia
Two-copy core $R(1, \rho) \leq K_*$ proved (residue (ii))	verify_dagger_extremal2.p
Kluyver excess: small- t head + signed ring tail	verify_dagger_kluyver.py
Window split: variance floor on $(1, 1.005]$	verify_dagger_window.py
Signed two-ring tail of the Kluyver excess	verify_dagger_tail.py
Multivariate concavity refuted (saddles)	verify_dagger_concavity.p
Near-pair basin $\lambda \geq 0.82$; singleton-channel gap	verify_dagger_decomp.py
Singleton band $< K_*$: three-region Arb bound	verify_dagger_singleton.p
Continuum lift: multivariate amplitude Taylor model	verify_dagger_continuum.p
Multi-copy profile $F(s)$, $t > N/2$ subdominant (Prop. 29)	verify_bn_residual_profil
Defect grows with $\omega(m)$ (unbounded, Rem. 7)	verify_defect_growth.py
Family limits $\text{defect}_\infty(m)$ (Table 3)	verify_defect_family_limi
Certified gap $A_N < U_N$, $N \leq 15$ (Lipschitz grid)	verify_defect_certified.p
Spectral-dimension threshold $d_s=1$	verify_spectral_dimension
Finite-time recurrence $T_\varepsilon(N)$	verify_finite_time.py
$\hookrightarrow$ certified upper bound $T_{0.2}(N) \leq \tilde{t}$, small primes (interval)	verify_finite_time_certif
Disorder vs. resonance (gen. eig.)	verify_disorder.py
Order map $A_N / \ln N$, prefix (Table 5)	verify_conj_order_map.py
Extended order run ($N \sim 1000$, criterion ≤ 500)	verify_order_extended.py
Filimonov segment law $d_j^N \sim \frac{4}{\pi^2} \ln j$	verify_filimonov_schrodin
Order $A_N \sim \frac{1}{\pi} \ln N$, all N (Thm. 6, Lem. 5)	verify_order_theorem.py
Classification $A_N = U_N \Leftrightarrow \text{prime}/2^m$ (Thm. 16)	verify_classification_pro
Segment rank & classification (Thm. 10)	verify_segment_classifica
Segment mod-4 ceiling criterion (Cor. 11)	exact/gap/segment_mod4.g
Exact GAP/PARI/Rocq layer, cross-check	verify_exact_layer.py
Exact classification $A_N = U_N$, $N \leq 800$ (GAP)	exact/gap/scan_classifica
Cyclotomic identities cross-check (PARI)	exact/pari/crosscheck.gp
$N=p^2$ obstruction comb (GAP, Fig. 11)	exact/gap/prime_square.g
Machine-checked mod-4 criterion (Rocq)	formal/rocq/Mod4Criterion
Machine-checked classification core, all N (Rocq)	formal/rocq/Classificatio
F92 splash table: segment $4/\pi^2$ vs ring $1/\pi$	verify_filimonov_table.py
B_N multi-copy profile, $t > N/2$ subdominant	verify_bn_residual_profil
Defect moment–SOS certified bounds	verify_defect_moment_sos.

16. STATISTICS: DENSITY OF STATES AND VALUE DISTRIBUTION

The spectrum and the large wave both admit clean probabilistic descriptions, verified in `verify_statistics.py`.

Density of states. As $N \rightarrow \infty$ the eigenvalues $\lambda_r = 4 \sin^2(\pi r/N)$ —the pushforward of the uniform law by $\lambda(x) = 2 - 2 \cos 2\pi x$ —converge to the **arcsine law** on $[0, 4]$,

$$\rho(\lambda) = \frac{1}{\pi \sqrt{\lambda(4-\lambda)}}, \quad F(\lambda) = \frac{2}{\pi} \arcsin \frac{\sqrt{\lambda}}{2}. \quad (28)$$

This is the change of variables: $\frac{d\lambda}{dx} = 4\pi \sin 2\pi x = 2\pi \sqrt{\lambda(4-\lambda)}$ (since $\sin 2\pi x = \frac{1}{2}\sqrt{\lambda(4-\lambda)}$), and each $\lambda \in (0, 4)$ has *two* preimages in $[0, 1]$, so $\rho(\lambda) = 2/(d\lambda/dx) = 1/(\pi \sqrt{\lambda(4-\lambda)})$ —the classical density of states of the one-dimensional discrete Laplacian (Figure 13); the frequencies have $\rho(\omega) = 2/(\pi \sqrt{4-\omega^2})$ on $[0, 2]$. Both fit to Kolmogorov–Smirnov distance $O(N^{-1})$ (a deterministic quantile grid: the empirical CDF lies within $2/N$ of the limit, the $1/N$ jump floor).

Value distribution. For prime N the impulse response $u_0(t)$ is almost-periodic with $\mathbb{Q}$ -independent frequencies, so by the Bohr–Jessen theorem its values possess a limiting distribution (Figure 13, right): symmetric, mean 0, and variance $\text{Var}_t u_0 = \frac{N^2-1}{12N^2} \rightarrow \frac{1}{12}$. The variance is exact: with $u_0(t) = \frac{1}{N} \sum_r \omega_r^{-1} \sin(\omega_r t)$ and the time-average $\langle \sin(\omega_r t) \sin(\omega_s t) \rangle = \frac{1}{2}$ when $\omega_s = \omega_r$ (i.e. $s \in \{r, N-r\}$) and 0 otherwise, the degeneracy makes each cross-term reinforce the diagonal, so $\langle u_0^2 \rangle = \frac{1}{N^2} \sum_r \omega_r^{-2} = \zeta_{LN}(1)/N^2 = \frac{N^2-1}{12N^2}$. The law is *platykurtic* (kurtosis $\approx 2.4 < 3$) because the lowest mode dominates (weight $1/\omega_1 \sim N$, a single arcsine sinusoid). The large wave is therefore *not* a heavy tail of the typical value but the rare near-alignment supremum (Figure 12): a spike that climbs to nearly U_N standing far above the thin-tailed bulk of size RMS $\approx 1/\sqrt{12}$, $A_N \sim \frac{1}{\pi} \ln N$.




FIGURE 12. What the large wave looks like in time ($N = 13$, prime). Left: the node-0 response $G_N(0, t)$ oscillates within the thin band $\pm \text{RMS}$ ($\text{RMS} \approx 0.29$) but throws rare spikes toward the ceiling $U_{13} = 0.856$. Right: a zoom on the largest near-alignment spike, reaching $\approx 0.99 U_N$ —the large wave forming as the modal phases briefly align (verify_thesis_peaks.py).

17. HIGHER-DIMENSIONAL TORI: THE LARGE WAVE IS ONE-DIMENSIONAL

On the d -dimensional torus $(\mathbb{Z}_N)^d$ the Laplacian is block-circulant-with-circulant-blocks (BCCB), diagonalized by the d -dimensional DFT with eigenvalues $\lambda_{\mathbf{k}} = \sum_{i=1}^d 4 \sin^2(\pi k_i/N)$ and $\omega_{\mathbf{k}} = \sqrt{\lambda_{\mathbf{k}}}$. The triangle inequality gives, unconditionally, $A_N^{(d)} \leq U_N^{(d)} = N^{-d} \sum_{\mathbf{k} \neq 0} \omega_{\mathbf{k}}^{-1}$. Near the origin $\omega_{\mathbf{k}} \sim (2\pi/N)|\mathbf{k}|$, so $N^d U_N^{(d)}$ mirrors the infrared lattice sum $\sum 1/|\mathbf{k}|$, i.e. the integral $\int d^d k/|\mathbf{k}|$: logarithmically divergent for $d = 1$ (giving $U_N \sim \frac{1}{\pi} \ln N$) but *convergent* for $d \geq 2$. Hence (Figure 15)

$$U_N^{(d)} = \Theta(\ln N) \quad (d = 1), \quad U_N^{(d)} = O(1) \quad (d \geq 2) : \quad (29)$$

the velocity Green’s function—and with it the large-wave amplification—is *uniformly bounded* on the 2- and 3-torus. The large wave is intrinsically one-dimensional, a sharp dimensional threshold set by the integrability of $1/\omega$ (verify_torus_dimension.py).




FIGURE 13. Left: spectral density of states converges to the arcsine law (28). Right: the Bohr–Jessen value law of $u_0(t)$ —symmetric, variance $(N^2 - 1)/(12N^2) \rightarrow 1/12$, platykurtic (flatter than the Gaussian of equal variance, dashed).

Spectral dimension. The threshold is not metric but *spectral*. For vertex-transitive (Fourier-flat) graph families the local ceiling reduces to the normalized spectral sum $U_G = |V|^{-1} \sum_{\lambda_r > 0} \lambda_r^{-1/2}$ (λ_r the Laplacian spectrum, $\omega_r = \sqrt{\lambda_r}$); for a general graph the ceiling at (j, k) is $\sum_{\lambda_r > 0} |\phi_r(j)\phi_r(k)| \lambda_r^{-1/2}$ and also reflects eigenvector localization through the local spectral measure, not the global density of states alone. If the integrated density of states obeys $N(\lambda) \sim \lambda^{d_s/2}$ near 0—defining the *spectral dimension* d_s —then $\rho(\lambda) \sim \lambda^{d_s/2-1}$ and $U_G \sim \int_0 \lambda^{-1/2} \rho(\lambda) d\lambda \sim \int_0 \lambda^{(d_s-3)/2} d\lambda$, which converges at the origin iff $d_s > 1$. Hence the *ceiling* U_G grows logarithmically precisely at the critical spectral dimension $d_s = 1$, is bounded for $d_s > 1$, and is power-law for $d_s < 1$; the d -torus is the case $d_s = d$. For $d_s > 1$ this rigorously *forbids* a large wave ($A_N \leq U_G = O(1)$); at $d_s = 1$ the ceiling permits one and is reached on the ring ($A_N \sim \frac{1}{\pi} \ln N$, Theorem 6), the same question for a general quasi-1D lattice being the natural next case. The criterion is geometric rather than metric: any *quasi-one-dimensional* lattice—a ladder $C_N \times C_2$ or a tube $C_N \times C_w$ of bounded width w —still has $d_s = 1$ and retains a logarithmically divergent ceiling, now with reduced prefactor $U_G \sim \frac{1}{\pi w} \ln N$ (only one of the w transverse channels is gapless, Figure 14). Bounded transverse width does not destroy the concentration; a genuine second extended dimension does (`verify_spectral_dimension.py`).

The order argument is in fact graph-agnostic: the same Kronecker alignment proves $\Theta(\ln N)$ for any family meeting two conditions.

Proposition 33 (order criterion at $d_s = 1$). *Let $\{G_N\}$ be finite graphs whose diagonal velocity Green’s function is $G_N(t) = \sum_r b_r \sin(\omega_r t)$ with $b_r > 0$ and ceiling $U_N = \sum_r b_r$. Suppose (A) $U_N \sim c \ln N$ ($c > 0$; the $d_s = 1$ logarithmic ceiling), and (B) some $\mathbb{Q}$ -independent subset S of the frequencies carries $\sum_{r \in S} b_r \geq (\frac{1}{2} + \delta)U_N$ for a fixed $\delta > 0$. Then $A_N := \sup_t \max_{\text{node}} |G_N(t)|$ obeys $2\delta U_N \leq A_N \leq U_N$, hence $A_N = \Theta(\ln N)$.*

Proof. $A_N \leq U_N$ since $|\sin| \leq 1$. For the lower bound, by Kronecker the $\mathbb{Q}$ -independent phases $\{\omega_r t\}_{r \in S}$ are simultaneously $\rightarrow \pi/2$ at suitable t ; there $G_N(t) \geq \sum_{r \in S} b_r - \sum_{r \notin S} b_r = 2 \sum_{r \in S} b_r - U_N \geq 2\delta U_N$ (the complement bounded below by $\sin \geq -1$). $\square$

The ring and Dirichlet segment are the instances where Lemma 5 supplies S (the length- $\frac{1}{2}\varphi(2N)$ prefix, carrying $(1 - o(1))U_N$, so $\delta \rightarrow \frac{1}{2}$). For a quasi-1D lattice (A) holds with $c = 1/(\pi w)$, while (B)—a

long $\mathbb{Q}$ -independent prefix—is the per-family analogue of the palindromic lemma: e.g. the next-nearest-neighbour ring, with $\omega_r = 2 \sin(\frac{\pi r}{N}) \sqrt{2 + \cos \frac{2\pi r}{N}}$, is $\mathbb{Q}$ -independent at $N = 15$ (so it *saturates*, $A_N = U_N$) yet carries the degeneracy $\omega_3 = \omega_6 = 2$ at $N = 12$ (`verify_quasi1d_order.py`). For this family the *degeneracy layer* of (B) can be settled completely.

Proposition 34 (NNN degeneracy law). *For the next-nearest-neighbour ring ($g = \frac{1}{2}$), two distinct frequencies coincide, $\omega_a = \omega_b$ with $a < b \leq \lfloor N/2 \rfloor$, if and only if $4 \mid N$ and $(a, b) = (N/4, N/2)$, where both equal 2.*

Proof. With $x = \cos \frac{2\pi a}{N}$ one has $\omega_a^2 = f(x) = 2(1-x)(2+x)$, and $f(x_a) - f(x_b) = -2(x_a - x_b)(1 + x_a + x_b)$. Since $\cos$ is injective on $(0, \pi]$ and $a \neq b$, a collision forces $x_a + x_b = -1$, i.e. the vanishing of the length-6 positive sum $2 + \zeta^a + \zeta^{-a} + \zeta^b + \zeta^{-b}$ of N -th roots of unity. By the Conway–Jones classification of rational cosine relations [7], the rational-angle solutions of $\cos A + \cos B = -1$ are $\{A, B\} = \{\pi/2, \pi\}$ and $A = B = 2\pi/3$ (excluded by $a \neq b$); hence $a = N/4$, $b = N/2$ (`verify_nnn_collisions.py`: confirmed *exactly*, in integer arithmetic modulo Φ_N , for all $N \leq 200$). $\square$

The law generalises to every rational coupling: for next-nearest weight $g > 0$ the symbol factorises as $\omega^2 = (1-x)(2+4g(1+x))$, $x = \cos \frac{2\pi a}{N}$, so distinct frequencies collide iff $\cos \frac{2\pi a}{N} + \cos \frac{2\pi b}{N} = -\frac{1}{2g}$ —again a rational cosine relation classified by [7]. For $g = 1$ the solutions are $\{\pi/3, \pi\}$, $\{2\pi/5, 4\pi/5\}$, $\{\pi/2, 2\pi/3\}$, so collisions occur exactly at $(N/6, N/2)$, $(N/5, 2N/5)$, $(N/4, N/3)$ when 6, 5, 12 divide N (verified exactly for $N \leq 120$); for $g = \frac{1}{4}$ the required value -2 forces $A = B = \pi$, and there are *no* collisions for any N . The degeneracy layer of condition (B) is thus settled for the whole one-parameter NNN family. Beyond collisions, a 100-digit PSLQ search finds no further rational relation among the NNN ($g = \frac{1}{2}$) frequencies for any $5 \leq N \leq 40$; dropping the duplicate when $4 \mid N$ then leaves a $\mathbb{Q}$ -independent set carrying the full ceiling, so on this evidence condition (B)—and with it $A = \Theta(\ln N)$ —holds throughout the family. Deciding (B) for a *general* $d_s = 1$ lattice remains open; the criterion isolates exactly what must be shown.

For one natural class, though, (B) is free—whenever the ring sits inside the lattice as a factor.

Corollary 35 (order on ring products). *For any fixed connected graph H , the large wave on the Cartesian product $C_N \square H$ satisfies $A = \Theta(\ln N)$.*

Proof. The Laplacian of $C_N \square H$ has eigenvalues $\lambda_r + \mu_h$, with $\lambda_r = 2 - 2 \cos \frac{2\pi r}{N}$ on C_N and μ_h on H . Only the gapless branch $h = 0$ ($\mu_0 = 0$, H -constant eigenvector $w_0 = |H|^{-1/2} \mathbf{1}$) has $\omega \rightarrow 0$, so the ceiling $U \sim \frac{1}{\pi|H|} \ln N$ is carried by it while the gapped branches add $O(1)$ —condition (A). The $h = 0$ frequencies are exactly the ring frequencies $2 \sin(\pi r/N)$, whose first $\frac{1}{2}\varphi(2N)$ are $\mathbb{Q}$ -independent by Lemma 5 and contribute $(1 - o(1))U > \frac{1}{2}U$ —condition (B). Proposition 33 then gives $A = \Theta(\ln N)$ (`verify_product_order.py`: the prefix bound $2L_{\text{pre}} - U$ is positive and grows like $\ln N$ for $H = K_2, P_3$). $\square$

So the order law survives multiplication by any bounded “cross-section”: ladders, prisms, tubes $C_N \square H$ all carry the $\Theta(\ln N)$ large wave, the embedded ring supplying the arithmetic. The same embedded-ring mechanism reaches beyond products.

Corollary 36 (Möbius ladder). *The large wave on the Möbius ladder M_N —the circulant $C_{2N}(\{1, N\})$ on $2N$ vertices, not a Cartesian product—satisfies $A(M_N) = \Theta(\ln N)$.*

Proof. The Laplacian eigenvalues are $\Lambda_r = 2(1 - \cos \frac{\pi r}{N}) + (1 - (-1)^r)$, $r = 0, \dots, 2N - 1$. The even branch $r = 2j$ gives $\Lambda_{2j} = 2(1 - \cos \frac{2\pi j}{N}) = 4 \sin^2(\pi j/N)$ —*exactly* the ring C_N spectrum—while the

diameter chord gaps the odd branch by $1 - (-1)^r = 2$. So $\omega \rightarrow 0$ only on the even branch, which carries the ceiling $U \sim \frac{1}{2\pi} \ln N$ (condition A); its frequencies are the ring frequencies $2 \sin(\pi j/N)$, whose first $\frac{1}{2}\varphi(2N)$ are $\mathbb{Q}$ -independent by Lemma 5 (exact integer-rank certificate, e.g. rank 128 = $\frac{1}{2}\varphi(512)$ at $N = 256$), contributing $(1 - o(1))U$ (condition B). Proposition 33 gives $A(M_N) = \Theta(\ln N)$ (verify_kjump_order.py). $\square$

Remark 37 (multi-jump circulants saturate). For a k -jump circulant $C_N(S)$, $|S| \geq 2$, the frequencies are the nested radicals $\omega_r = \left(\sum_{s \in S} 4 \sin^2(\pi s r/N)\right)^{1/2}$, and condition (A) holds with slope $1/(\pi \sqrt{\sum_{s \in S} s^2})$. Distinct frequencies coincide only along exact Conway–Jones congruences (for $S = \{1, 2, 3\}$: a collision iff $4 \mid N$, $6 \mid N$, or $7 \mid N$; for $S = \{1, 3\}$: iff $8 \mid N$ or $5 \mid N$), and—unlike the single sines of the ring, which carry the composite mod-4 obstruction of Theorem 16—the *distinct* radicals appear to satisfy no further $\mathbb{Q}$ -relation (PSLQ to 250 digits, the spurious relation norm diverging with precision). Numerically the multi-jump circulant then *saturates*, $A_N = U_N$, for every N off its collision congruences: the radical structure destroys the cyclotomic linear relations, so composite N is no longer an obstruction (verify_kjump_order.py). For the nearest case $S = \{1, 2\}$ this is no longer only numerical: the $\mathbb{Q}$ -independence of $\{\sqrt{\alpha_r}\}$, $\alpha_r = 4 - 2 \cos \frac{2\pi r}{p} - 2 \cos \frac{4\pi r}{p} \in \mathbb{Q}(\zeta_p)$, is *proved exactly* by the Besicovitch–Kummer criterion—no nonempty subset product $\prod_{r \in T} \alpha_r$ is a square in $\mathbb{Q}(\zeta_p)$ (checked over the cyclotomic field, not in floating point)—for every collision-free prime $p \leq 23$, so $A_p = U_p$ there is a theorem (exact/pari/kjump_kummer.gp). The general (all- N , all- S) statement remains the nested-radical independence question.

The remaining genuinely open case is a quasi-1D lattice *without* any embedded ring—e.g. the open (Dirichlet) triangular ladder, whose spectrum is not cyclotomic; there condition (A) holds ($d_s = 1$) but (B) has no clean arithmetic and stays per-family.




FIGURE 14. Spectral-dimension threshold. The ceiling $U_G = |V|^{-1} \sum \lambda^{-1/2}$ grows like $\ln N$ for $d_s = 1$ —the ring and, with reduced slope $1/(\pi w)$, the quasi-1D ladder and tube—but saturates for the $d_s = 2$ square torus. The ceiling is a critical-dimension ($d_s = 1$) phenomenon.




FIGURE 15. Ceilings $U_N^{(d)}$ versus $\ln N$. Only $d = 1$ grows (slope $1/\pi$, on the dashed $\frac{1}{\pi} \ln N + C_0$); the 2- and 3-torus ceilings saturate, so $A_N^{(d)} = O(1)$ for $d \geq 2$.

Part 2. Nonlinear extension

The results of Part II are primarily computational and exploratory. They outline natural nonlinear continuations of the linear theory—and the energy-concentration thread that unifies them—rather than claiming a complete nonlinear theory; each is reproduced by a verification script, and the open analytical questions are flagged as such.

18. ROGUE WAVES AND THE DISCRETE NLS

Restoring a cubic on-site interaction turns the linear Schrödinger ring of Section 11 into the discrete nonlinear Schrödinger (DNLS) equation

$$i \dot{z}_j = -(L_N z)_j + \gamma |z_j|^2 z_j, \quad \gamma > 0 \text{ (focusing)}, \quad (30)$$

whose linear part is exactly L_N and whose continuum limit (lattice spacing $\rightarrow 0$, $\lambda_r \rightarrow k^2$) is the *focusing* nonlinear Schrödinger equation

$$i \psi_t + \psi_{xx} + 2|\psi|^2 \psi = 0, \quad (31)$$

the canonical envelope model of deep-water and optical rogue waves [50, 72, 74, 73]. (The literal limit is $i \psi_t = \psi_{xx} + \gamma |\psi|^2 \psi$, equal to (31) up to the symmetry $\psi \mapsto \bar{\psi}$ and the rescaling $z \mapsto cz$ that sets $\gamma = 2$.) Thus the very operator L_N that governs the linear large wave sits at the heart of the rogue-wave problem, and the two amplification mechanisms can be compared on the same lattice (`verify_rogue_dnls.py`).

Proposition 38 (Hamiltonian structure and conservation). *The DNLS (30) is Hamiltonian, $i \dot{z}_j = \partial H / \partial \bar{z}_j$ with*

$$H(z) = - \sum_j |z_{j+1} - z_j|^2 + \frac{\gamma}{2} \sum_j |z_j|^4,$$

and conserves both H and the power $P(z) = \sum_j |z_j|^2$; consequently the flow is global and norm-preserving on $\mathbb{C}^N$.

Proof. $\frac{d}{dt}P = 2 \operatorname{Re} \sum_j \dot{z}_j \bar{z}_j = 2 \operatorname{Re}[-i(\langle L_N z, z \rangle - \gamma \sum_j |z_j|^4)] = 0$, since $\langle L_N z, z \rangle$ and $\sum_j |z_j|^4$ are real (so $-i(\cdot)$ is purely imaginary); and $\partial H / \partial \bar{z}_j = (z_{j+1} + z_{j-1} - 2z_j) + \gamma |z_j|^2 z_j = -(\langle L_N z, z \rangle)_j + \gamma |z_j|^2 z_j$, so $i\dot{z} = \partial H / \partial \bar{z}$ and H is conserved along the flow. With P bounded, no finite-time blow-up occurs. Both invariants are reproduced numerically: the norm-preserving split-step integrator holds P to 10^{-12} and H to $O(\Delta t^2)$ (verify_dnls_selftrapping.py). $\square$

Proposition 39 (the linear large wave persists under weak nonlinearity). *Let $u(t)$ solve the DNLS (30) with $u(0) = u_0$, and write $u_{\text{lin}}(t) = e^{itL_N} u_0$ for the linear flow of (30) (the complex conjugate of the Schrödinger propagator e^{-itL_N} of Section 11, with the same ℓ^p operator norms). Then for all t ,*

$$\|u(t) - u_{\text{lin}}(t)\|_\infty \leq \|u(t) - u_{\text{lin}}(t)\|_2 \leq |\gamma| t \|u_0\|_2^3.$$

In particular, if the linear flow attains a large wave $\|u_{\text{lin}}(t_)\|_\infty = M$ at a time t_* , then $\|u(t_*)\|_\infty \geq M - |\gamma| t_* \|u_0\|_2^3$; so whenever $|\gamma| < M/(2t_* \|u_0\|_2^3)$ the nonlinear solution still exhibits a wave of height $\geq M/2$.*

Proof. By Duhamel, $u(t) - u_{\text{lin}}(t) = -i\gamma \int_0^t e^{i(t-s)L_N} (|u|^2 u)(s) ds$. As e^{isL_N} is ℓ^2 -unitary, $\|u(t) - u_{\text{lin}}(t)\|_2 \leq |\gamma| \int_0^t \| |u|^2 u(s) \|_2 ds \leq |\gamma| \int_0^t \|u(s)\|_\infty^2 \|u(s)\|_2 ds$. The power $\|u\|_2^2$ is conserved (Proposition 38), so $\|u(s)\|_2 = \|u_0\|_2$ and $\|u(s)\|_\infty \leq \|u(s)\|_2 = \|u_0\|_2$; the integrand is $\leq \|u_0\|_2^3$. The last claim is the reverse triangle inequality $\|u(t_*)\|_\infty \geq \|u_{\text{lin}}(t_*)\|_\infty - \|u - u_{\text{lin}}\|_\infty$. $\square$

This is a rigorous core for Part II: the linear large wave is a genuine feature of the weakly nonlinear flow, not an artefact of linearisation (verify_dnls_persistence.py). Its scope is honest—the admissible coupling shrinks with the (Diophantine, possibly large) alignment time t_* , so it does not control the strongly nonlinear regime where the rogue-wave focusing of §18 lives; that remains open.

A second rigorous ingredient—the modulational-instability band, which ties the nonlinear seeding directly to the L_N spectrum—is Proposition 40 below.

Peregrine soliton. Equation (31) possesses the rational solution [65]

$$\psi_P(x, t) = \left[1 - \frac{4(1 + 4it)}{1 + 4x^2 + 16t^2} \right] e^{2it}, \quad (32)$$

localized in both space and time, with unit background $|\psi| \rightarrow 1$ and peak $|\psi_P(0, 0)| = 3$: the threefold amplification “appearing from nowhere and disappearing without a trace” that defines a rogue wave (Figure 16). We verified that (32) solves (31) to a PDE residual vanishing under grid refinement. The Peregrine breather is the limit of the space-periodic Akhmediev and time-periodic Kuznetsov–Ma breather families produced by the Darboux transformation [66]. These exact breathers exist because the focusing NLS (31) is *completely integrable*—the compatibility condition of the Zakharov–Shabat linear system, solved by the inverse scattering transform [51], on which the Darboux dressing rests.

Modulational instability. The seed of these structures is the Benjamin–Feir instability [64, 53], the mechanism by which Zakharov’s deep-water envelope equation [50] destabilizes a uniform train. The uniform train $z_j = A e^{-i\gamma A^2 t}$ solves (30) ($L_N \mathbf{1} = 0$); writing $z_j = [A + p_j(t)] e^{-i\gamma A^2 t}$ and keeping terms linear in p , the cubic contributes $\gamma A^2 (2p_j + \bar{p}_j)$, leaving

$$i\dot{p}_j = -(L_N p)_j + \gamma A^2 (p_j + \bar{p}_j).$$

The Fourier ansatz $p_j = u e^{i\kappa j + \sigma t} + \bar{v} e^{-i\kappa j + \sigma t}$ (eigenvalue λ_κ of L_N on $e^{i\kappa j}$) reduces this to the 2×2 system $i\sigma u = (\gamma A^2 - \lambda_\kappa)u + \gamma A^2 v$, $-i\sigma v = \gamma A^2 u + (\gamma A^2 - \lambda_\kappa)v$, whose solvability condition $\sigma^2 + (\gamma A^2 - \lambda_\kappa)^2 = (\gamma A^2)^2$ gives the growth rate

$$\sigma(\kappa) = \sqrt{\lambda_\kappa (2\gamma A^2 - \lambda_\kappa)}, \quad \lambda_\kappa = 4 \sin^2 \frac{\kappa}{2}, \quad (33)$$

so the band $\lambda_\kappa < 2\gamma A^2$ is unstable (continuum form $\sigma = |k|\sqrt{2\gamma A^2 - k^2}$, the classical $|k|\sqrt{4A^2 - k^2}$ at $\gamma = 2$). This linear-stability computation is exact:

Proposition 40 (modulational-instability band). *The uniform state $z_j = A e^{-i\gamma A^2 t}$ of (30) is linearly unstable to the Fourier mode κ if and only if $\lambda_\kappa < 2\gamma A^2$, with growth rate $\sigma(\kappa) = \sqrt{\lambda_\kappa(2\gamma A^2 - \lambda_\kappa)}$; the rate is maximal, $\sigma_{\max} = \gamma A^2$, at $\lambda_\kappa = \gamma A^2$.*

The spectrum of the full $2N \times 2N$ linearised generator matches $\sigma(\kappa)$ exactly (`verify_dnl_s_mi.py`), and a direct split-step simulation of (30) reproduces (33) to better than 1% across the band interior (Figure 17); the fastest mode does not grow without bound but develops into a recurrent breather-like modulation reminiscent of the Akhmediev scenario of the continuum integrable NLS—a periodic avatar of the rogue spike. The instability band and its growth rates are thus read off the L_N spectrum: the linear large-wave operator also sets the nonlinear seeding. In the long-wave limit $\lambda_\kappa \approx \kappa^2$ the band collapses to $\sigma(\kappa) = \kappa\sqrt{2\gamma A^2 - \kappa^2}$, exactly the classical Benjamin–Feir gain of the continuum focusing NLS—the Lighthill criterion $\omega'' q > 0$ holding for the focusing sign [18, 22]—so Proposition 40 is its spectrally exact discrete counterpart on the ring.

From thin tails to heavy tails. The statistical signature is the sharpest contrast with the linear theory. Seeded by a noisy plane wave, (30) on the ring spawns localized bursts exceeding an operational envelope threshold of twice the root-mean-square amplitude (analogous to common rogue-wave criteria)—and the amplitude distribution develops a heavy, *leptokurtic* tail (kurtosis > 3) lying well above the Rayleigh law of a linear Gaussian sea (Figure 18). This inverts the linear picture, whose Bohr–Jessen value law is *platykurtic* (Section 16): linear amplification is a rare phase alignment of bounded, thin-tailed oscillations, whereas nonlinear focusing manufactures genuine heavy-tailed extremes. The two regimes share the lattice operator L_N but reach large amplitudes by opposite routes—interference versus self-focusing.

Of the rigorous nonlinear programme, three pieces are settled in this paper: the weak-nonlinearity persistence (Proposition 39), the modulational-instability band—the role of the L_N spectrum—(Proposition 40), and the existence and stability of the discrete breathers (Propositions 43, 44 in Section 19). Integrable discretizations of (31) and the fully nonlinear saturation/rogue regime are left to future work; equation (30) is the natural nonlinear continuation of the linear *Schrödinger* large wave B_N of Section 11 (the first-order $\ell^1 \rightarrow \ell^\infty$ amplification), not of the second-order mechanical wave A_N of Part I—the two are distinct large-wave regimes on the same L_N . Proposition 39 accordingly persists the *Schrödinger* peak, not the bead-chain one.

19. DISCRETE BREATHERS AND SELF-TRAPPING

Nonlinearity does not only sharpen extremes; it can *localize*. Seed the DNLS (30) with a single excited site, $z_j(0) = A \delta_{j0}$ (norm $\Lambda = \sum_j |z_j|^2 = A^2$), and track the *participation ratio*

$$P = \frac{(\sum_j |z_j|^2)^2}{\sum_j |z_j|^4} \in [1, N], \quad (34)$$

which is 1 for a single occupied site and $\sim N$ when the excitation spreads over the ring. A norm-conserving split-step integrator (norm exact to 10^{-12} , H to $O(\Delta t^2)$) shows a sharp *self-trapping transition* at the on-site nonlinear frequency $\gamma|z_0|^2 \approx 4$ —exactly the top of the spectrum of L_N , the band $[0, 4]$. Below it the excitation disperses ($P \rightarrow N$); above it the site detunes out of the linear band and forms a long-lived, exponentially localized *breather-like state* $z_j(t) \approx \phi_j e^{-i\Omega t}$ consistent with the stationary DNLS branch, whose real profile solves $\Omega \phi_j = -(L_N \phi)_j + \gamma \phi_j^3$ with frequency $\Omega \approx \gamma|z_0|^2$ pushed above the whole phonon band $[0, 4]$, so no extended mode can resonate (Figure 19; the trapped profile is steady, the temporal standard deviation of $|z_j|$ staying below 0.02). This is the dynamical *opposite* of




FIGURE 16. Peregrine soliton (32) of the focusing NLS (31): a rogue peak localized in space and time (left), with the threefold amplification over the unit background at $t = 0$ (right).




FIGURE 17. Modulational (Benjamin–Feir) instability gain (33) of the plane wave under the DNLS (30): lattice rate (solid), continuum approximation (dotted) and simulated growth rates (squares), with unstable band edge $\kappa = \pi/2$ and maximal rate $\sigma_{\max} = \gamma A^2$.

the linear large wave, a fully delocalized interference of all N modes ($P \sim N$): the same operator L_N supports both a spread-out resonance (linear) and a pinned soliton (nonlinear), selected by the size of the cubic term (`verify_dnls_selftrapping.py`).

Two thresholds. An anti-continuum estimate places the *stationary* single-site breather at $\gamma|\phi_0|^2 \gtrsim 2$: from $\Omega\phi_j = -(L_N\phi)_j + \gamma\phi_j^3$ with $\phi \approx \sqrt{P}\delta_{j0}$ one gets $\Omega \approx \gamma P - 2$, which clears the band *top* ($\Omega > 0$) at $\gamma P \gtrsim 2$. The larger *dynamical* threshold $\gamma|z_0|^2 \approx 4$ measured from a single-site *launch* is the bandwidth: the spreading excitation must outrun the entire band $[0, 4]$, not merely its edge. This single-site barrier at $A_*^2 = 2(\sigma + 1)^{1/\sigma} = 4$ (cubic $\sigma = 1$) was identified by Kevrekidis–Espinola-Rocha–Drossinos–Stefanov [19] as a *sufficient* trapping condition ($H_{ss} < 0$); we make that side a theorem with




FIGURE 18. DNLS (30) on the ring seeded by a noisy plane wave. Left: space-time amplitude $|z_j(t)|$ with localized rogue spikes. Right: the amplitude exceedance distribution develops a heavy tail above the Rayleigh law (dashed) of a linear Gaussian sea, crossing the rogue threshold—the opposite of the platykurtic linear value law.

the sharp constant and show it is *energy-sharp* (below). One side of the dynamical threshold is thus settled, with the constant exactly the bandwidth 4. (The *static* excitation threshold of [70] is a different object: a positive minimal power for a ground state exists only when $\sigma \geq 2/d$, so in 1D cubic there is none—localized modes exist at arbitrarily small power, while the dynamical launch barrier 4 persists.)

Proposition 41 (no dispersal above the bandwidth). *Let u solve (30) with single-site data $u_j(0) = \sqrt{P} \delta_{jj_0}$ ($N \geq 3$) and $\gamma P > 4$. Then for all t ,*

$$\sum_j |u_j(t)|^4 \geq P^2 - \frac{4P}{\gamma}, \quad \max_j |u_j(t)|^2 \geq P - \frac{4}{\gamma} > 0, \quad \text{PR}(u(t)) \leq \frac{1}{1 - 4/(\gamma P)},$$

where $\text{PR}(u) = (\sum_j |u_j|^2)^2 / \sum_j |u_j|^4$ is the participation ratio: the excitation never disperses.

Proof. $H = -\langle L_N u, u \rangle + \frac{\gamma}{2} \sum_j |u_j|^4$ and the power P are conserved (Proposition 38). At $t = 0$ the single site has $\langle L_N u_0, u_0 \rangle = 2P$, so $H_0 = -2P + \frac{\gamma}{2} P^2$; at any t , $\frac{\gamma}{2} \sum |u|^4 = H_0 + \langle L_N u, u \rangle \geq H_0$ since $L_N \succeq 0$, giving the first bound. Dividing by $\sum |u|^2 = P$ gives $\max_j |u_j|^2 \geq \sum |u|^4 / P \geq P - 4/\gamma$ and the participation bound (cf. the variational excitation thresholds of [70]; `verify_selftrapping_threshold.py`). $\square$

The converse needs care: on the *finite* ring nothing disperses forever. The flow is Hamiltonian, hence preserves Lebesgue measure on the invariant bounded shell $\{P - \delta \leq \|u\|_2^2 \leq P\}$, so by Poincaré recurrence almost every datum returns arbitrarily close to itself: for a.e. u_0 , $\limsup_{t \rightarrow \infty} \max_j |u_j(t)|^2 = \max_j |u_j(0)|^2$, for *every* γ . Numerically the single-site orbit below threshold behaves exactly so: at $\gamma P = 3$ the maximum dips to 0.05 and later recovers past 0.5 (`verify_selftrapping_threshold.py`). The threshold dichotomy therefore concerns the *infimum*: above the bandwidth, $\inf_t \max_j |u_j(t)|^2 \geq P - 4/\gamma$ (Proposition 41); below it the infimum drops to the dispersal scale (0.05 at $\gamma P = 3$). In the weak-coupling part of the sub-threshold range that drop is a theorem:

Proposition 42 (quantitative dispersal at weak coupling). *Let u solve (30) with single-site data $u_j(0) = \sqrt{P} \delta_{jj_0}$, and set $t_* = (c_L / (3 \cdot 2^{1/3} \gamma P))^{3/4}$, where $c_L = 0.7858 \dots$ is Landau's Bessel constant [71]. If*

$N \geq 5t_*$ (i.e. $N \gtrsim (\gamma P)^{-3/4}$), then

$$\inf_t \max_j |u_j(t)|^2 \leq \frac{16}{3^{3/2}} (c_L 2^{-1/3})^{3/2} P \sqrt{\gamma P} (1 + o(1)) \leq 1.52 P \sqrt{\gamma P} (1 + o(1)).$$

In particular for $\gamma P \leq 0.43$ the infimum sits strictly below P : the excitation quantitatively disperses, and with Proposition 41 the undecided window of the threshold narrows to $\gamma P \in [0.43, 4)$.

Proof. For $t \leq N/5$ the aliased copies of the ring kernel are super-exponentially small (the periodization estimate of Theorem 22), so by Landau's bound $\max_n |J_n(x)| \leq c_L x^{-1/3}$, $\max_j |u_{\text{lin},j}(t)| \leq \sqrt{P} c_L (2t)^{-1/3} + \varepsilon_N$. Adding $\|u - u_{\text{lin}}\|_\infty \leq \gamma P^{3/2} t$ (Proposition 39),

$$\max_j |u_j(t)| \leq \sqrt{P} c_L (2t)^{-1/3} + \gamma P^{3/2} t + \varepsilon_N.$$

The right side is minimized at $t = t_*$, where it equals $4 \cdot 3^{-3/4} (c_L 2^{-1/3})^{3/4} (\gamma P)^{1/4} \sqrt{P} (1 + o(1))$; squaring gives the claim (`verify_dispersal_weak.py`). $\square$

The threshold constant $4 = \|L_N\|$ is *sharp*, not merely sufficient: full dispersal ($\|u\|_4^4 \rightarrow 0$) forces, by energy conservation, $\langle L_N u, u \rangle \rightarrow 2P - \frac{\gamma}{2} P^2$, which is non-negative iff $\gamma P \leq 4$. Above 4 full dispersal is thus energetically forbidden (the bound $P - 4/\gamma$); the identity is confirmed to four digits at the dispersal minimum, e.g. $\langle L_N u, u \rangle \rightarrow 0.503$ versus $2P - \frac{\gamma}{2} P^2 = 0.500$ at $\gamma P = 3$ (`verify_selftrapping_transition.py`). Numerically the single-site seed disperses cleanly for $\gamma P \lesssim 3.5$ and self-traps (tracking $P - 4/\gamma$) for $\gamma P \gtrsim 4.5$, the crossover straddling 4. So the window $\gamma P \in [0.43, 4)$ is a *proof gap*—dispersal is the established behaviour, proven (Proposition 42) only up to 0.43—rather than a behavioural unknown. Closing it, beyond perturbation theory, is the genuinely dispersive open direction.

The virial obstruction, located precisely. A second-moment (Morawetz-type) attempt isolates the wall to a single estimate. On the line—the ring for $t \lesssim N/5$, where aliasing is negligible—the spread $V(t) = \sum_j j^2 |a_j|^2$ obeys the *exact* identity (verified to machine precision along the flow, `verify_virial_window.py`)

$$\ddot{V} = 2\langle L_2 a, a \rangle - 2\gamma \sum_n (2n+1) (|a_n|^2 - |a_{n+1}|^2) \text{Re}(\bar{a}_n a_{n+1}), \quad (35)$$

where L_2 is the next-nearest-neighbour Laplacian (symbol $2 - 2\cos 2k = 4\sin^2 k$), so $2\langle L_2 a, a \rangle = 2\langle v^2 \rangle P \geq 0$ ($v_k = 2\sin k$ the group velocity) is a manifestly non-negative *dispersive* term and the position-weighted sum is the nonlinear correction. At the single-site seed the correction vanishes, giving $\ddot{V}(0) = 4P$ independently of γ ; numerically $\ddot{V} > 0$ throughout the window, decaying to 0 as $\gamma P \rightarrow 4$, so the virial is convex, $V \rightarrow \infty$, which *supports* dispersal. Closing the argument needs the correction bounded by $(1 - \epsilon) 2\langle L_2 a, a \rangle$; numerically it is small (a few percent of the dispersive term) and degenerates to exactly $-2\langle L_2 a, a \rangle$ as $\gamma P \rightarrow 4$, but its position weight $(2n+1)$ is not controlled by the conserved mass, energy, and momentum. Indeed energy fixes only the *first* spectral moment, $\langle \cos k \rangle_\infty = \gamma P/4$ (equivalent to $\langle L_N u, u \rangle \rightarrow 2P - \frac{\gamma}{2} P^2$ above), < 1 for $\gamma P < 4$ but not forcing $\langle v^2 \rangle = 2 - 2\langle \cos 2k \rangle > 0$: a bimodal $k = 0/\pi$ momentum split has $\langle \cos k \rangle = \gamma P/4$ yet $\langle v^2 \rangle = 0$. So the dynamics avoids the staggered ($k = \pi$) trap (numerically $\langle v^2 \rangle > 0$) but the conserved quantities do not see it. This is exactly where Morawetz/Strichartz arguments stall: the discrete Strichartz estimates of Stefanov–Kevrekidis [20] (built on the sharp $t^{-1/3}$ decay [71, 21]) close small-data scattering only for $\sigma \geq 3/d$ —in 1D, power ≥ 7 , not the cubic here—so they provably cannot reach the window, and the continuum interaction-Morawetz/concentration machinery consumes the Galilean invariance the lattice lacks. Equation (35) thus reduces the open problem to a single estimate—a bound on the position-weighted nonlinear term, i.e. on the staggered momentum content $\langle \cos 2k \rangle$.

An equivalent, position-weight-free form. The dispersive term is set by the staggered momentum $\langle L_2 a, a \rangle = 2P(1 - \langle \cos 2k \rangle)$, $\langle \cos 2k \rangle = P^{-1} \text{Re} \sum_n \bar{a}_n a_{n+2}$, and—because the linear flow preserves every momentum moment $|\hat{a}_k|^2$ —its evolution is purely nonlinear and carries *no* position weight:

$$\frac{d}{dt} \text{Re} \sum_n \bar{a}_n a_{n+2} = -\gamma \sum_n (|a_n|^2 - |a_{n+2}|^2) \text{Im}(\bar{a}_n a_{n+2}), \quad \left| \frac{d}{dt} \text{Re} \sum_n \bar{a}_n a_{n+2} \right| \leq 2\gamma \|a\|_4^4 \quad (36)$$

(the identity proven symbolically; the bound by Cauchy–Schwarz; `verify_stagger_morawetz.py`). So the whole question rests on the *focusing interaction-Morawetz* spacetime norm $\gamma \int_0^\infty \|a\|_4^4 dt$ —the standard scattering criterion—which is numerically bounded for $\gamma P < 4$ and divergent above (matching the sharp threshold 4); its focusing-sign obstruction is exactly the barrier [20] cannot cross at the cubic power. The integrable Ablowitz–Ladik cousin, whose hierarchy *conserves* a staggered-momentum functional, motivates (36); but that functional *drifts* under the non-integrable flow (30), the $O(1)$ Ablowitz–Ladik gap made quantitative.

On the finite ring the *existence* of these breathers needs no MacKay–Aubry continuation from the anti-continuum limit [67]—it is a compactness statement.

Proposition 43 (existence of discrete breathers). *For every $\Omega > 0$ the focusing ring DNLS (30) ($\gamma > 0$) has a standing-wave breather $z_j(t) = e^{-i\Omega t} \phi_j$ with real $\phi \neq 0$ solving $(L_N + \Omega)\phi = \gamma|\phi|^2\phi$; for large Ω it is on-site localized.*

Proof. $M := L_N + \Omega$ is positive definite ($L_N \succeq 0$, $\Omega > 0$). Minimize the continuous $\langle M\phi, \phi \rangle$ over the compact set $\{\|\phi\|_4^4 = 1\}$; a nonzero minimizer ϕ_* exists. Its Lagrange condition is $M\phi_* = 2\mu|\phi_*|^2\phi_*$, i.e. $(L_N + \Omega)\phi_* = \gamma|\phi_*|^2\phi_*$ with $\gamma = 2\mu > 0$; the scaling $\phi \mapsto s\phi$ sends $\gamma \mapsto \gamma/s^2$, so every $\gamma > 0$ is realised. As $\Omega \rightarrow \infty$, $\langle M\phi, \phi \rangle \sim \Omega\|\phi\|_2^2$, so the minimizer maximises $\|\phi\|_4/\|\phi\|_2$, i.e. concentrates on one site (`verify_dnls_breather.py`: the minimizer solves the equation to $\sim 10^{-7}$ and its participation ratio $\rightarrow 1$). $\square$

So existence is settled on the ring; the anti-continuum estimate $\gamma P \gtrsim 2$ matches the $\Omega > 0$ (above-band) condition. Stability follows the classical site-vs-bond dichotomy.

Proposition 44 (stability of the site-centred breather). *Along the site-centred branch $P(\Omega) = \|\phi_\Omega\|^2 = \Omega/\gamma + O(1)$, so $dP/d\Omega > 0$; by the Vakhitov–Kolokolov slope criterion and Grillakis–Shatah–Strauss theory [69, 68] the breather is orbitally stable, and its linearisation $\begin{bmatrix} 0 & -L_- \\ L_+ & 0 \end{bmatrix}$ with $L_+ = L_N + \Omega - 3\gamma\phi^2$, $L_- = L_N + \Omega - \gamma\phi^2$ has purely imaginary spectrum. The bond-centred branch carries a real eigenvalue pair and is unstable.*

Proof. Grillakis–Shatah–Strauss gives orbital stability once three hypotheses hold: (i) the slope $dP/d\Omega > 0$; (ii) $L_- \geq 0$ with $\ker L_- = \text{span } \phi$; (iii) $n(L_+) = 1$, a single negative direction. We establish them. (i) The anti-continuum expansion (one site of weight $\sqrt{P}$ gives $\Omega = \gamma P - 2 + O(1/P)$) yields $P(\Omega) = \Omega/\gamma + O(1)$, so $dP/d\Omega = 1/\gamma + O(1/P) > 0$. (ii) The site-centred profile is *positive*, $\phi_\Omega > 0$ (anti-continuum construction with a positive seed, preserved under continuation), and $L_- \phi = 0$ by the stationary equation $L_N \phi + \gamma\phi^3 = \Omega\phi$. A strictly positive vector in the kernel of the discrete Schrödinger operator $L_- = L_N + \Omega - \gamma\phi^2$ is, by Perron–Frobenius on the finite ring, its ground state; hence $L_- \geq 0$ with simple kernel $\text{span } \phi$. (iii) $\langle L_+ \phi, \phi \rangle = \langle L_- \phi, \phi \rangle - 2\gamma\|\phi\|_4^4 = -2\gamma\|\phi\|_4^4 < 0$, so $n(L_+) \geq 1$; in the anti-continuum limit $L_+ = L_N + \Omega - 3\gamma\phi^2$ has exactly one negative eigenvalue (the deep well $-3\gamma\phi^2$ at the excited site against $L_N + \Omega > 0$ elsewhere), and this persists along the branch since no eigenvalue of L_+ crosses 0 on the tested range, so $n(L_+) = 1$. With (i)–(iii), GSS yields orbital stability of the site-centred wave. Both branches are moreover *certified in exact rational arithmetic*

(`verify_dnls_breather_cap.py`): a Newton–Kantorovich test promotes a rational profile $\tilde{\phi}$ to a true breather (here $\|\phi - \tilde{\phi}\|_2 \leq 4 \times 10^{-9}$), exact LDL^T inertia (Sylvester’s law) at $\pm I$ pins the Morse index with the spectrum a distance ≥ 1 from 0, and a congruence $Q^\top(L_- - I)Q \succ 0$ on $\tilde{\phi}^\perp$ gives $\lambda_2(L_-) \geq 1$; a Weyl bound ($\delta \leq 5 \times 10^{-8} \ll 1$) carries both signatures to the true ϕ . This yields $n(L_+) = 1$ for the site-centred ring $N = 5, \Omega = 2$ (GSS index 0, stable) and $n(L_+) = 2$ for the bond-centred ring $N = 6, \Omega = 2$ (GSS index $n(L_+) - 1 = 1$, odd, hence orbitally *unstable*)—upgrading both verdicts from numerical to computer-assisted. Over the full branch `verify_dnls_breather_stability.py` tracks the spectrum ($\max |\operatorname{Re} \lambda| \sim 10^{-7}$ site; a real pair $\in [3.9, 8.9]$ bond) on $\Omega \in [2, 20]$. $\square$

Thus on the ring the linear stage (persistence, the MI band), the breather’s *existence*, and its *stability* are all rigorous; the precise dynamical launch threshold ≈ 4 and the fully nonlinear saturation/rogue dynamics remain the open sequel.




FIGURE 19. DNLS self-trapping ($N = 128$, single-site seed). Left: the participation ratio P collapses from $\sim N$ (dispersed) to ~ 1 (trapped) as the nonlinearity crosses the band top $\gamma \approx 4$. Right: above threshold the excitation is a persistent localized breather-like state—a localized column in space-time.

20. FPUT RECURRENCE AND THE ROUTE TO CHAOS

The cubic coupling also makes the lattice a testbed for Hamiltonian chaos. The Fermi–Pasta–Ulam–Tsingou (FPUT) β -ring

$$\ddot{u}_j = (u_{j+1} - 2u_j + u_{j-1}) + \beta[(u_{j+1} - u_j)^3 - (u_j - u_{j-1})^3] \quad (37)$$

has linear part $\ddot{u} = -L_N u$ and Hamiltonian $H = \sum_j [\frac{1}{2} \dot{u}_j^2 + \frac{1}{2} (u_{j+1} - u_j)^2 + \frac{\beta}{4} (u_{j+1} - u_j)^4]$. Seeding all energy in mode 1 and integrating symplectically (velocity Verlet; H conserved to 6×10^{-4} with no secular drift), at low energy the energy does *not* thermalize: it sloshes to a few neighbouring modes and *returns* to mode 1—the FPUT recurrence (Figure 20), the near-integrable face of the nonlinear lattice: its long-wave limit is a Korteweg–de Vries-type equation, a completely integrable Hamiltonian system [52] whose invariant tori the weakly nonlinear lattice shadows, so energy returns rather than equipartitioning. Only a handful of low modes are ever active, the dynamics is quasi-periodic, and the maximal Lyapunov exponent is $\lambda_{\max} \approx 0$ —just as for the integrable linear large wave.

Raising the energy breaks integrability. Measured by the Benettin method (a tangent vector co-evolved with the variational equations and periodically renormalized, $\lambda_{\max} = T^{-1} \sum \log \|\delta\|/d_0$), the maximal Lyapunov exponent turns on with energy (Figure 21): $\lambda_{\max} \approx 0.001$ in the regular regime, rising through a weak-chaos knee to $\lambda_{\max} \approx 0.05$ in the strongly chaotic regime, where the energy thermalizes across all modes. The same lattice—linear core L_N —thus interpolates between integrable tori and deterministic chaos as the amplitude grows (`verify_fput_recurrence.py`, `verify_lyapunov.py`). The transition is visible directly in phase space: in the three odd-mode energy coordinates (E_1, E_3, E_5) (a mode-1 seed excites only odd modes by parity) the orbit lies on a thin quasi-periodic torus at low energy and fills a chaotic cloud at high energy (Figure 22).




FIGURE 20. FPUT- β recurrence ($N = 32$, mode-1 seed). The harmonic energy of mode 1 drops almost to zero and recurs to its initial value, sloshing quasi-periodically with the low modes instead of thermalizing.

21. DIAGNOSTICS AND THE FATE OF THE LARGE WAVE

The linear large wave and its nonlinear continuations are cleanly separated by a small set of dynamical diagnostics—localization (participation ratio P), value statistics (kurtosis κ) and chaos (Lyapunov exponent $\lambda_{\max}$)—all built on the *same* lattice operator L_N (Table 10). The linear large wave is a delocalized, thin-tailed, integrable resonance ($P \sim N$, $\kappa < 3$, $\lambda_{\max} = 0$); nonlinearity opens three distinct routes to a large local amplitude—*self-focusing* into rogue extremes (heavy tails, $\kappa > 3$), *self-trapping* into a pinned breather ($P \sim 1$), and *chaotic* thermalization ($\lambda_{\max} > 0$). Which one is realized is set by the sign and strength of the nonlinearity and by the energy; the linear theory of Part I is the $\beta, \gamma \rightarrow 0$ corner of this picture, and supplies the spectral data (L_N , the band $[0, 4]$, the modal frequencies) that organize every regime. A rigorous nonlinear theory on the ring—existence and stability of the discrete breathers, the precise role of the L_N spectrum in the modulational band, and the energy threshold for chaos—is the natural sequel.




FIGURE 21. Route to chaos. The maximal Lyapunov exponent of the FPUT- β ring is ≈ 0 at low energy (regular, the recurrence regime) and turns on past an energy threshold, reaching $\lambda_{\max} \approx 0.05$ in the chaotic regime.




FIGURE 22. FPUT phase portrait in the odd-mode energy space (E_1, E_3, E_5) . Left: at low energy the orbit lies on a thin quasi-periodic torus (regular). Right: at high energy it fills a three-dimensional chaotic cloud—the same regular-to-chaotic transition measured by $\lambda_{\max}$.

22. HIGHER DIMENSIONS: NONLINEAR LOCALIZATION BEYOND THE ONE-DIMENSIONAL LARGE WAVE

Part I established that the *linear* large wave is one-dimensional: on the d -torus the ceiling $U_N^{(d)}$ is bounded for $d \geq 2$ (Section 17), so a localized linear disturbance cannot amplify in two or three dimensions. Nonlinearity removes the obstruction. On the 2-torus $(\mathbb{Z}_N)^2$ the DNLS with the block-circulant Laplacian $L_N^{(2)}$ (eigenvalues $\lambda_{\mathbf{r}} = 4 \sin^2 \frac{\pi r_1}{N} + 4 \sin^2 \frac{\pi r_2}{N}$, band $[0, 8]$, with the dispersion surface of Figure 24) self-traps a single-site excitation into a long-lived, exponentially localized *two-dimensional*

TABLE 10. The large wave and its nonlinear continuations, separated by dynamical diagnostics on the common operator L_N . (P : participation ratio; κ : kurtosis; $\lambda_{\max}$: largest Lyapunov exponent; “—” marks a diagnostic not characteristic of that regime.)

Regime	Model	Localization P	Statistics κ	$\lambda_{\max}$
Linear large wave	$\ddot{u} = -L_N u$	delocalized $\sim N$	platykurtic < 3	0
Rogue / self-focusing	DNLS, noisy plane wave	localized bursts	leptokurtic > 3	—
Discrete breather	DNLS, $\gamma z_0 ^2 > 4$	pinned ~ 1	—	0
FPUT recurrence	FPUT- β , low E	few active modes	—	≈ 0
Chaos	FPUT- β , high E	thermalized $\sim N$	—	> 0

breather-like state once the on-site nonlinear frequency crosses the 2D band top, $\gamma|z|^2 \approx 8$: the norm is conserved to 10^{-12} and the participation ratio collapses from $\sim N^2$ to $O(1)$ across the threshold (Figure 23, `verify_dnls_2d.py`). A large, localized amplitude therefore survives in 2D (and, identically, 3D) by a mechanism—self-focusing self-trapping—absent from the linear theory: the dimensional threshold of Part I is a property of the *linear* resonance, not of the lattice.

This multidimensional, strongly nonlinear regime points toward Zakharov’s wider program: a possible multidimensional long-wave asymptotic direction is the Zakharov–Kuznetsov regime [54] (we do not carry out the multiscale derivation from (37) here), and the statistics of the thermalizing lattice suggest a connection to the weak-turbulence (Kolmogorov–Zakharov) theory of wave spectra [55]. The rigorous wave-kinetic results of [56] concern a different weakly nonlinear, large-box scaling (dimension $d \geq 3$) and do not directly justify the finite one-dimensional lattice here. These connect the discrete large-wave problem to the broad theory of nonlinear wave dynamics (Section 21).

2D discrete breather (DNLS, $\gamma = 16$, $N = 48$)




FIGURE 23. Two-dimensional discrete breather: a single-site seed under the 2D DNLS self-traps into a persistent localized hump above the 2D band top ($\gamma \approx 8$). Nonlinearity restores a large localized amplitude in 2D, where the linear large wave is bounded.




FIGURE 24. Dispersion surface $\omega(\theta_x, \theta_y) = 2\sqrt{\sin^2(\theta_x/2) + \sin^2(\theta_y/2)}$ of the block-circulant Laplacian on the 2-torus: zero at the corners, band top $2\sqrt{2}$ at the centre.

23. CAUSTICS AND THE CONCENTRATION OF ENERGY

The thread running through this paper is the *concentration of energy* into a large local amplitude. Part I concentrates it by modal interference (the linear large wave); Part II so far by nonlinear self-focusing (rogue waves, breathers). A third, purely *linear* route is refractive: in a slowly varying medium a wave follows rays $\dot{\mathbf{x}} = \nabla_{\mathbf{k}}\omega$, $\dot{\mathbf{k}} = -\nabla_{\mathbf{x}}\omega$, and where neighbouring rays cross, the ray density—and with it the amplitude—diverges on a *caustic*. Caustics are classified by catastrophe theory [57]: the generic fold (Airy diffraction) and cusp (the bright curve at the bottom of a teacup); diffraction regularizes the divergence to a finite but large peak with a universal scaling.

Rays and their envelopes. Geometric optics is the skeleton. For waves crossing a current the Doppler term refracts the rays, and a bundle of initially parallel rays entering a current jet converges and *folds*: its Jacobian $\partial x/\partial x_0$ passes through zero at a finite focal distance, where the ray density $1/|\partial x/\partial x_0|$ —and with it the linear amplitude—diverges, regularized by diffraction (Figure 25, `verify_ray_caustic.py`).

The discrete front is a caustic. The lattice already contains one. For the discrete Schrödinger equation (Section 11) the group velocity $v_g(k) = 2 \sin k$ is extremal at $k = \frac{\pi}{2}$ —a degenerate stationary-phase point, a fold—so the propagator front is exactly the Airy function,

$$J_n(2t) \sim t^{-1/3} \text{Ai}\left(2^{1/3} \frac{n-2t}{(2t)^{1/3}}\right), \quad (38)$$

with the universal fold scaling $\max_n |J_n(2t)| \sim 0.536 t^{-1/3}$ (Figure 28, `verify_caustic_airy.py`). The discrete large-wave front *is* a fold caustic, the $t^{-1/3}$ being the Airy exponent.

The cusp and the teacup. The next generic caustic is the *cusp*—the bright curve at the bottom of a sunlit teacup, whose ray version (parallel rays reflected inside a circle) is a *nephroid* and whose universal diffraction is the Pearcey integral $\text{Pe}(X, Y) = \int_{-\infty}^{\infty} \exp[i(s^4 + Xs^2 + Ys)] ds$ over the cusp catastrophe, with geometric caustic $8X^3 + 27Y^2 = 0$. Its cusp-point value is $|\text{Pe}(0, 0)| = \frac{1}{2}\Gamma(\frac{1}{4}) = 1.8128$, and the intensity concentrates inside the cusp and falls off in the geometric shadow (Figure 26, `verify_cusp_pearcey.py`).




FIGURE 25. Ray tracing in a current jet. A bundle of parallel rays is refracted toward the jet, crosses on a caustic (the fold near $z \approx 8$, then re-focusing), where the ray density and the linear amplitude concentrate.




FIGURE 26. The cusp caustic. Left: the everyday teacup caustic—parallel rays reflected inside a circle have a nephroid envelope. Right: the universal Pearcey diffraction $|\text{Pe}(X, Y)|^2$ of the cusp, bright inside the cusp and dark in the shadow.

Ocean rogue hot spots. The same refraction sets the geography of ocean rogue waves. Swell from Southern-Ocean storms running into the opposing Agulhas current off south-east Africa is refracted and focused into caustics—the classic rogue “hot spot” [59, 58, 62]. A spatially random current produces a cascade of caustics, *branched flow*: a uniform incident wave develops bright filaments and focal hot spots whose intensity is heavy-tailed, the linear precursor of rogue statistics [60]. The paraxial model $i \partial_z \psi = -\frac{1}{2} \partial_{xx} \psi + V(x) \psi$ with a weak random current V reproduces it exactly (Figure 29): from a flat wave the scintillation index $\max_x |\psi|^2 / \langle |\psi|^2 \rangle$ rises from 1 to ~ 13 at a focal distance, and the focal-plane

intensity is heavy-tailed—its tail an order of magnitude above the exponential law of a structureless field (`verify_branched_flow.py`). Tidal (lunar) currents modulate the refracting flow, a secondary effect.

The lens seeds self-focusing. The two routes cooperate. A linear caustic concentrates energy until the local steepness crosses the modulational threshold, whereupon nonlinear self-focusing takes over—rogue waves are “the nonlinear stage of the modulational instability” [61]. Seeding the focusing NLS with a converging (lens) beam (Figure 27), the linear caustic amplifies a weak flat input $\sim 7\times$, and the nonlinearity then sharpens the focal peak by a further $\sim 11\%$, whereas a lensless plane wave never spikes—the lens is what localizes (`verify_caustic_mi.py`). In one dimension the nonlinear enhancement is modest (the NLS being L^2 -subcritical); the dramatic collapsing self-focusing is the two-dimensional Zakharov regime. The discrete operator L_N thus organizes the concentration of energy in all three guises—interference, refraction, and self-focusing—on a single lattice.




FIGURE 27. A linear caustic seeds nonlinear self-focusing (focusing NLS, lens-seeded beam). Left: the peak amplitude—the lens focuses a weak flat input $\sim 7\times$ (caustic), the nonlinearity adds a further $\sim 11\%$, while the lensless plane wave stays flat. Right: at the focus the nonlinearity sharpens and heightens the peak.

24. DISCUSSION

The ring and the Dirichlet segment realize the same mechanism—constructive phase alignment of a logarithmically weighted family of modes—through different arithmetic. On the segment the frequencies $2 \sin(\pi k/2N)$ lie in $\mathbb{Q}(\zeta_{4N})$; on the ring $2 \sin(\pi r/N) \in \mathbb{Q}(\zeta_{2N})$, with the twofold degeneracy $\omega_r = \omega_{N-r}$ that halves the number of independent phases and (as the time-averaged $\text{RMS}_t |G_N(0, \cdot)| = \sqrt{(N^2 - 1)/(12N^2)} \rightarrow 1/\sqrt{12}$ shows) doubles the typical energy relative to a naive independent count. The growth constant $1/\pi$ is identical on the classes where the frequencies are rationally independent; the composite defect is a sub-torus phenomenon that does not lower the order— $A_N \sim \frac{1}{\pi} \ln N$ for *all* N (Theorem 6), the same leading constant.

The discrete Schrödinger equation occupies the opposite extreme. Its phases are governed by $\lambda_r = 2 - 2 \cos(2\pi r/N)$, which lie in the real cyclotomic field $\mathbb{Q}(\zeta_N)^+$ and always satisfy the affine relation $\sum_r \lambda_r = 2N$. Unitarity caps the response to a localized state, and the genuine amplification—ballistic, of order $\sqrt{N}$ —is carried by spread-out ℓ^∞ data. The contrast $\ln N$ (wave) versus $\sqrt{N}$ (Schrödinger) traces back to the difference between $\sqrt{\lambda_r}$ and λ_r in the propagator phase.




FIGURE 28. The discrete-Schrödinger front as a fold caustic. Left: the rescaled fronts $t^{1/3} J_n(2t)$ for $t = 50, 200, 800$ collapse onto the Airy function $\text{Ai}(s)$, $s = 2^{1/3}(n - 2t)/(2t)^{1/3}$ (Eq. (38)). Right: the front peak follows the universal fold scaling $0.536 t^{-1/3}$.




FIGURE 29. Branched flow of a wave in a weak random current (paraxial model). Left: a uniform incident wave ($z = 0$) develops bright caustic filaments and focal hot spots as it propagates. Right: the scintillation index rises from 1 to ~ 13 at the first focal distance—the linear, refractive route to a rogue hot spot.

GLT context. The sequence $\{L_N\}$ is a circulant—hence locally Toeplitz—instance of a Generalized Locally Toeplitz (GLT) sequence [34] with symbol $f(\theta) = 4 \sin^2(\theta/2)$. The spectral-zeta density (9) and the Szegő asymptotics are the GLT spectral distribution specialized to this symbol; beyond the uniform d -torus settled in Section 17, the GLT calculus extends the *asymptotic spectral-distribution* statements (the Szegő/zeta density) to variable-coefficient lattices (non-uniform masses) and anisotropic or graph Laplacians; the *arithmetic* results—the cyclotomic ranks, the mod-4 criterion, and the reachability of the ceiling—do not transfer automatically and are the natural setting for future work.

A different face of energy concentration. The discrete Laplacian concentrates energy beyond wave dynamics. In the dielectric-breakdown (Laplacian-growth) model of lightning [63], the electrostatic potential solves the discrete Laplace equation $L_N^{(2)}\phi = 0$ off the leader, which advances toward the strongest field; the result is a branched, fractal discharge (fractal dimension ≈ 1.5 in our run) with the field—and the energy—concentrated at its tips, the same operator and branched morphology as the refractive caustics of Section 23 but by a distinct (elliptic, dissipative) mechanism (Figure 30, `verify_laplacian_lightning.py`).




FIGURE 30. Lightning as Laplacian growth: the leader (white) advances cloud-to-ground where the field is strongest, the potential ϕ solving the discrete Laplace equation $L_N^{(2)}\phi = 0$ off it. A branched fractal discharge with the energy concentrated at its tips—energy concentration on the same discrete Laplacian, electrostatically.

Limitations. We mark the boundaries of what is established. (i) The order law $A_N \sim \frac{1}{\pi} \ln N$ is now proved for *every* N (Theorem 6); the deficit $U_N - A_N$ on the non-saturating N is $O(\ln \ln \ln N)$ (proved), and on the numerical evidence *unbounded*—growing with the number of distinct odd prime factors (Remark 7); its exact value is algebraic per N with no apparent closed form ($A_N/U_N \in [0.84, 1]$). (ii) The saturation locus is settled—no composite N saturates (Theorem 16)—and A_N for individual composite N is algebraic (Proposition 12) but computed per N , not given in closed form. (iii) For the Schrödinger $B_N = \Theta(\sqrt{N})$ the constant splits by parity (Remark 23): $\liminf_N B_N/\sqrt{N} \geq c_0/\sqrt{2} \approx 0.861$ is proved and is attained numerically on the even subsequence, while the odd subsequence approaches $\beta_{\text{odd}} \approx 0.928$; closing the remaining upper bound would therefore give a parity split and no single limiting constant. The whole $t > N/2$ reduction is now rigorous (joint equidistribution and its uniform transfer to $F(s)$, Theorems 26, 27), leaving one validated sharp inequality (a Gaussian-excess/expected-modulus bound) before $\limsup_N B_N/\sqrt{N} = \beta_{\text{odd}}$. The segment constant $C = 4/\pi^2$ (Section 11.4) is *not* ours: it is the published Myshkis–Filimonov value [41, Thm. 2], which the Schrödinger ceiling inherits because the two ceilings coincide. (iv) On the Dirichlet segment the rank, saturation, *and* order are now fully proved (Theorem 10; the order transfers verbatim from Theorem 6); what remains there, as on the ring, is the exact deficit on the non-saturating $N + 1$. (v) Part II is in larger part computational;

its rigorous chain comprises the Hamiltonian conservation laws (Proposition 38), the weak-nonlinearity persistence of the linear large wave (Proposition 39), the modulational-instability band (Proposition 40), and the existence and stability of the discrete breathers (Propositions 43, 44), the no-dispersal side of the dynamical self-trapping threshold with the exact constant 4 (Proposition 41), and the weak-coupling converse (Proposition 42); the threshold window $\gamma P \in [0.43, 4)$, the saturation/rogue dynamics, and the route to chaos are simulated, not proved. (vi) The supremum A_N is an *infinite-time* quantity; the amplitude reached on any physical horizon is Diophantine-limited (Section 10). Proved statements are separated throughout from numerical evidence and from the conjectures, and every quantitative claim is reproduced by a standalone script.

25. OPEN PROBLEMS

- (1) The deficit $U_N - A_N$ on the non-saturating N . The order $A_N \sim \frac{1}{\pi} \ln N$ is proved for *all* N (Theorem 6) and $A_N = U_N$ iff N is prime or 2^m (Theorem 16); the deficit is $O(\ln \ln \ln N)$, depends asymptotically only on the odd part of N , and (Remark 7) is *unbounded*—numerically it grows with the number of distinct odd prime factors (the subtorus optimizer is validated by returning $U_N - A_N = 0$ to machine precision on primes of the same size), so $\Theta(\ln \ln \ln N)$ along primorials. The gap $A_N < U_N$ is *rigorously certified* for small N (Lipschitz-grid / elementary Lasserre: $U_N - A_N \geq 0.049, 0.057, 0.030, 0.082$ for $N = 6, 9, 12, 15$); what remains is a rigorous *asymptotic* lower bound along primorials (the grid is exponential in $\frac{1}{2}\varphi(2N)$) and the exact algebraic value per N (certified per- N brackets follow from the Lasserre moment-SOS hierarchy, Section 9). For $\omega(m) = 2$ (odd part a product of two primes) the structure is now explicit: the relation lattice has unimodular Smith normal form, and its shortest relations are *chained* with support set by the smallest prime— $3 \mid N$ gives a support-3 relation $\omega_a + \omega_b = \omega_c$ with $a + b = N/3$ (from $2 \sin \frac{\pi}{6} = 1$), $5 \mid N$ a support-5 golden-ratio relation—and the dependent band carries $\Theta(N)$ relation-locked modes of diverging total weight, yet at the optimizer their signed contribution is a *modest* $O(1)$ *balance* (a positive dependent net ≈ 0.15 against a negative prefix deficit ≈ -0.07 , both roughly constant across the accessible $\omega(m) \leq 3$), not a near-total cancellation. The free sandwich that closes $\omega(m) \leq 1$ becomes $\Theta(1)$, so the lemma reduces to a sup-side estimate—why this signed net stays $O(1)$ as $\omega(m) \rightarrow \infty$ —the primal counterpart of McGehee-Pigno-Smith/Konyagin, on which crude modulus bounds overshoot; the min-side Chowla machinery (Bedert’s $n^{1/7}$ growth) runs the opposite way and does *not* transfer (verify_excess_omega2_spectral.py). Structurally this is the *same* obstruction as the multi-copy residual (2) (verify_excess_omega2.py).
- (2) The Schrödinger constant $B_N/\sqrt{N}$. The order $B_N = \Theta(\sqrt{N})$ is settled, $\liminf \geq c_0/\sqrt{2} \approx 0.861$ is proved (Theorem 22; numerically attained on the even subsequence), and the constant splits by the parity of N (Remark 23): the odd-subsequence value $\beta_{\text{odd}} = 0.9280193048\dots$ is the elliptic-integral average (14) (a 25-digit PSLQ search finds no elementary closed form), and $\|K_N(\cdot, t)\|_1 \leq (\beta_{\text{odd}} + o(1))\sqrt{N}$ holds on $t \leq N/2$, now *in full*: the two-copy reduction is rigorous, the limiting profile is the *explicit* elliptic-integral function $F(s)$ (Proposition 24), and its maximization $\max F = F(1) = \beta_{\text{odd}}$ is *proved*—analytically on $[\frac{1}{2}, 0.937]$ (Jensen, sharpened by the tangent to the concave h) and by validated interval arithmetic on $[0.937, 1]$. For the multi-copy region $t > N/2$ the entire $N \rightarrow \infty$ reduction is now *rigorous*: the Debye phases are jointly equidistributed (Theorem 26, a van der Corput second-derivative bound) and this transfers *uniformly in u* to the ℓ^1 constant $F(s)$ (Theorem 27, a mollified Koksma-Hlawka estimate—the

modulus has divergent Hardy–Krause variation—with the cubic van der Corput test at the coalescing points), alongside the variance-floor inequality and the tail $\sqrt{\pi}/2 < \beta_{\text{odd}}$ (Lemma 30). A *single* validated residual now separates this from an unconditional $\limsup_N B_N/\sqrt{N} = \beta_{\text{odd}}$: the Gaussian-excess inequality $\mathbb{E}|Z| - \mathbb{E}|G| \leq K_* \Sigma_2^{-3/2} \sum_\ell a_\ell^4$, $K_* = (h(1) - \sqrt{\pi}/2)/2$, for the *balanced* Debye configurations of the profile (the live set is an interval of consecutive integers, so $|\# \text{ even} - \# \text{ odd}| \leq 1$; sharp at the two-copy pair, while imbalanced configurations violate the unrestricted bound), whose integrated majorant clears β_{odd} on $(1, 20]$; proving it needs a bound on the Bessel-ring tail of the Kluver excess (`verify_bn_profile_gaussian.py`, `verify_bandedge_transfer.py`).

- (3) A rigorous nonlinear theory of the discrete NLS (30) on the ring (Section 18). Two pieces are now settled: the linear large wave persists in the weakly nonlinear regime (Proposition 39), and the modulational-instability band and growth rates $\sigma_Q = \sqrt{\lambda_Q(2\gamma A^2 - \lambda_Q)}$ are read off the L_N spectrum (Proposition 40); breathers *exist* on the ring by a finite-dimensional variational argument (Proposition 43); and the site-centred breather is *stable* (Vakhitov–Kolokolov, $dP/d\Omega > 0$) while the bond-centred one is unstable (Proposition 44); one side of the dynamical launch threshold is settled with the *sharp* constant $4 = \|L_N\|$ —full dispersal is energetically forbidden for $\gamma P > 4$, since it would force $\langle L_N u, u \rangle \rightarrow 2P - \frac{\gamma}{2} P^2 < 0$ (Proposition 41)—and the converse is settled at weak coupling—quantitative dispersal $\inf_t \max_j |u_j|^2 \leq 1.52P\sqrt{\gamma P}$ for $N \gtrsim (\gamma P)^{-3/4}$ (Proposition 42; on the finite ring $\sup_t$ recurs by Poincaré). The window $\gamma P \in [0.43, 4)$ is a *proof gap*, not a behavioural unknown: the seed disperses cleanly up to $\gamma P \approx 4$ numerically (`verify_selftrapping_transition.py`), but the rigorous drop is proven only up to 0.43. What remains open is this dispersive proof gap together with the *fully* nonlinear regime—saturation and the rogue-wave focusing—where the admissible coupling shrinks with the Diofantine alignment time.
- (4) The multidimensional nonlinear large wave (Section 22): existence, stability, and the energy threshold of the 2D/3D discrete breathers, and the long-wave dynamics in the Zakharov–Kuznetsov regime [54]—a nonlinear counterpart to the dimensional threshold of Section 17.
- (5) Whether the statistics of the thermalizing lattice are governed by the weak-turbulence (Kolmogorov–Zakharov) wave kinetic description [55, 56], and the rate of approach to energy equipartition.
- (6) The interplay of the two energy-concentration routes (Section 23): when does a linear refractive caustic raise the local steepness enough to trigger the modulational instability of Section 18, and what rogue-amplitude statistics result on the lattice?

APPENDIX A. DERIVATION OF THE RING SPECTRUM

Seeking $L_N \varphi = \lambda \varphi$ with $\varphi_{j+N} = \varphi_j$, the constant-coefficient recurrence $-\varphi_{j+1} + (2 - \lambda)\varphi_j - \varphi_{j-1} = 0$ has characteristic equation $p^2 - (2 - \lambda)p + 1 = 0$ with $p_1 p_2 = 1$ (Vieta). Nontrivial bounded solutions require conjugate roots $p_{1,2} = e^{\pm i\theta}$ on the unit circle, so $2 - \lambda = 2 \cos \theta$. The periodicity $\varphi_{j+N} = \varphi_j$ forces $e^{i\theta N} = 1$, i.e. $\theta = 2\pi r/N$, whence

$$\lambda_r = 2 - 2 \cos \frac{2\pi r}{N} = 4 \sin^2 \frac{\pi r}{N}, \quad \varphi_j^{(r)} = e^{2\pi i r j / N} \quad (r = 0, \dots, N-1),$$

with the real basis $1, \cos \frac{2\pi r j}{N}, \sin \frac{2\pi r j}{N}, (-1)^j$. Each level is doubly degenerate, $\lambda_r = \lambda_{N-r}$, except $\lambda_0 = 0$ and—for even N — $\lambda_{N/2} = 4$. The Dirichlet segment is identical with the closure $\sin(\theta(N+1)) = 0$, giving $\theta = \pi k/(N+1)$ and $\lambda_k^{\text{Dir}} = 4 \sin^2 \frac{\pi k}{2(N+1)}$ (all simple).

APPENDIX B. ASYMPTOTICS OF THE COSECANT SUM

We derive Lemma 2. The sum $S_N = \sum_{r=1}^{N-1} \csc(\pi r/N)$ has reciprocal singularities at *both* ends— $\csc(\pi r/N) \sim N/(\pi r)$ as $r \rightarrow 0$ and $\sim N/(\pi(N-r))$ as $r \rightarrow N$ —which we subtract symmetrically:

$$S_N = \frac{N}{\pi} \sum_{r=1}^{N-1} \left(\frac{1}{r} + \frac{1}{N-r} \right) + \sum_{r=1}^{N-1} \underbrace{\left[\csc \frac{\pi r}{N} - \frac{N}{\pi r} - \frac{N}{\pi(N-r)} \right]}_{= R(r/N)}.$$

The first sum is $\frac{2N}{\pi} H_{N-1} = \frac{2N}{\pi} (\ln N + \gamma) + O(1/N)$. The integrand $R(x) = \csc \pi x - \frac{1}{\pi x} - \frac{1}{\pi(1-x)}$ is smooth and bounded on $[0, 1]$ (both poles cancel), so the second sum is a Riemann sum, $\sum_{r=1}^{N-1} R(r/N) = N \int_0^1 R(x) dx + O(1)$, and the integral is elementary (using $\int \csc \pi x dx = \frac{1}{\pi} \ln \tan \frac{\pi x}{2}$):

$$\int_0^1 R(x) dx = \frac{1}{\pi} \left[\ln \frac{(1-x) \tan(\pi x/2)}{x} \right]_0^1 = \frac{2}{\pi} \ln \frac{2}{\pi}.$$

Hence the leading asymptotics, with additive constant $\frac{1}{\pi}(\gamma + \ln \frac{2}{\pi}) = 0.03999\dots$,

$$S_N = \frac{2N}{\pi} \left(\ln N + \gamma + \ln \frac{2}{\pi} \right) + O(1), \quad U_N = \frac{S_N}{2N} = \frac{1}{\pi} \ln N + \frac{1}{\pi} \left(\gamma + \ln \frac{2}{\pi} \right) + o(1).$$

The next correction is the Euler–Maclaurin endpoint term of R (a Bernoulli- B_2 contribution); it is even in $1/N$ and evaluates to

$$U_N = \frac{1}{\pi} \ln N + \frac{1}{\pi} \left(\gamma + \ln \frac{2}{\pi} \right) - \frac{\pi}{72} \frac{1}{N^2} + O(N^{-4}), \quad -\frac{\pi}{72} = -\frac{\zeta(2)}{12\pi},$$

confirmed to six digits by Richardson extrapolation ($a_2 = -0.043633$). The correction is thus tied to the same $\zeta(2)$ that governs $\zeta_{LN}(2)$ in Section 6; see `verify_asymptotic_corrections.py`.

APPENDIX C. REPRODUCIBLE CODE

Every quantitative claim is reproduced by a standalone script run with `uv run` (inline PEP 723 dependencies); the figures are generated by `make_figures.py`, and the core routines (ring spectrum, Bohr ceiling, cyclotomic ranks, Schrödinger kernel) live in a tested package `large_wave_effect` with a `pytest` suite. The scripts are listed in Section 15. The complete repository—scripts, section sources, and figures—is public at <https://github.com/ilgrad/large-wave-effect>, where every `verify_*.py` passes.

Beyond PYTHON, three computer-algebra and proof systems give independent, exact (floating-point-free) corroboration of the arithmetic. We reproduce the essential fragments.

(a) GAP — exact cyclotomic ranks and the relation lattice (`exact/gap/relation_lattice.g`). The frequencies $\omega_r = 2 \sin(\pi r/N)$ are handled as the cyclotomic integers $\zeta_{2N}^r - \zeta_{2N}^{-r}$; rank and nullspace are computed in exact arithmetic, confirming $\text{rank}_{\mathbb{Q}} = \frac{1}{2} \varphi(2N)$ for every $N \leq 800$.

```
LW_FreqCoords := function(N)                # integer coords of zeta^r - zeta^{-r} in Q(zeta_2N)
)
local z; z := E(2 * N);
return List([1 .. Int(N / 2)], r -> CoeffsCyc(z^r - z^(-r), 2 * N));
end;;

LW_QRankFull := function(N)                  # Q-rank of {2 sin(pi r/N)} = 1/2 phi(2N) (Theorem
qrnk)
local z; z := E(2 * N);
return RankMat(List([1 .. N - 1], r -> CoeffsCyc(z^r - z^(-r), 2 * N)));
```

```

end;;
LW_RelationLattice := function(N)          # integer basis of Lambda_N = { k : sum_r k_r
  omega_r = 0 }
  return NullspaceIntMat(LW_FreqCoords(N));
end;;

```

(b) **Rocq — the order theorem’s arithmetic core** (formal/rocq/Classification.v). The palindromic/anti-palindromic clash behind Lemma 5 and the resulting index bound are machine-checked, uniformly in N , with the standard library only.

```

(* a coefficient equal to both +b (palindromic) and -b (anti-palindromic) is zero *)
Lemma pal_and_antipal_zero : forall a b : Z, a = b -> a = - b -> a = 0.
Proof. intros. lia. Qed.

(* hence, with Phi_2N | Q giving phi <= 2K and 2K <> phi, the largest index K >= phi/2 + 1 *)
Lemma prefix_index_bound : forall K phi : Z,
  Z.even phi = true -> phi <= 2*K -> 2*K <> phi -> phi/2 + 1 <= K.
Proof.
  intros K phi Hev Hle Hne. apply Z.even_spec in Hev. destruct Hev as [m Hm].
  assert (phi/2 = m) by (rewrite Hm; rewrite Z.mul_comm; apply Z.div_mul; lia). lia.
Qed.

```

(c) **FRICAS — the analytic identities of Appendix B** (exact/fricas/asymptotics.input). The cosecant Laurent series and the antiderivative entering U_N ’s expansion are obtained symbolically.

```

output(series(1/sin(x), x = 0))          -- csc x = 1/x + x/6 + 7 x^3/360 + ... (B_2 = 1/6)
output(integrate(1/sin(%pi*x), x))      -- = (1/pi) log tan(pi x/2) (enters U_N via Euler-
Maclaurin)
output(numeric(-%pi/72))                -- = -zeta(2)/(12 pi), the N^{-2} correction of U_N

```

Later sections of the same file re-derive, exactly, the NNN symbol factorisation and collision difference of Proposition 34, the minimisation identity behind Proposition 42, and the modulational-instability solvability condition of Proposition 40.

REFERENCES

- [1] H. Bohr, *Almost Periodic Functions*, Chelsea, New York, 1947.
- [2] H. Weyl, *Über die Gleichverteilung von Zahlen mod. Eins*, Math. Ann. **77** (1916), 313–352.
- [3] L. Kuipers, H. Niederreiter, *Uniform Distribution of Sequences*, Wiley, New York, 1974 (reprinted Dover, 2006). (Ch. 1: Weyl criterion; Ch. 2: the Erdős–Turán inequality.)
- [4] H. L. Montgomery, *Ten Lectures on the Interface between Analytic Number Theory and Harmonic Analysis*, CBMS Reg. Conf. Ser. Math. **84**, Amer. Math. Soc., 1994. (van der Corput’s method, §8.)
- [5] S. W. Graham, G. Kolesnik, *Van der Corput’s Method of Exponential Sums*, London Math. Soc. Lecture Note Ser. **126**, Cambridge Univ. Press, 1991.
- [6] I. Niven, *Irrational Numbers*, Carus Math. Monographs **11**, MAA, 1956.
- [7] J. H. Conway and A. J. Jones, *Trigonometric Diophantine equations (on vanishing sums of roots of unity)*, Acta Arith. **30** (1976), 229–240.
- [8] J. B. Rosser, L. Schoenfeld, *Approximate formulas for some functions of prime numbers*, Illinois J. Math. **6** (1962), 64–94. doi:10.1215/ijm/1255631807.
- [9] G. N. Watson, *A Treatise on the Theory of Bessel Functions*, Cambridge Univ. Press, 1944.
- [10] P. J. Olver, *Dispersive quantization*, Amer. Math. Monthly **117** (2010), no. 7, 599–610. doi:10.4169/000298910X496723.
- [11] E. Bombieri, J. Bourgain, *On Kahane’s ultraflat polynomials*, J. Eur. Math. Soc. **11** (2009), no. 3, 627–703. doi:10.4171/JEMS/163.
- [12] P. Erdős, *Some unsolved problems*, Michigan Math. J. **4** (1957). doi:10.1307/mmj/1028997963.

- [13] J.-P. Kahane, *Sur les polynômes à coefficients unimodulaires*, Bull. London Math. Soc. **12** (1980), no. 5, 321–342. doi:10.1112/blms/12.5.321.
- [14] B. Bedert, *Polynomial bounds for the Chowla cosine problem*, preprint (2025), arXiv:2509.05260.
- [15] Y. Fu, K. Ren, H. Wang, *A note on maximal operators for the Schrödinger equation on $\mathbb{T}^1$* , preprint (2023), arXiv:2307.12870.
- [16] O. C. McGehee, L. Pigno, B. Smith, *Hardy’s inequality and the L^1 norm of exponential sums*, Ann. of Math. (2) **113** (1981), no. 3, 613–618. doi:10.2307/2007000.
- [17] S. V. Konyagin, *On the Littlewood problem*, Izv. Akad. Nauk SSSR Ser. Mat. **45** (1981), no. 2, 243–265 (MR0637621).
- [18] M. J. Lighthill, *Contributions to the theory of waves in non-linear dispersive systems*, IMA J. Appl. Math. **1** (1965), no. 3, 269–306. doi:10.1093/imamat/1.3.269.
- [19] P. G. Kevrekidis, J. A. Espinola-Rocha, Y. Drossinos, A. Stefanov, *Dynamical barrier for the formation of solitary waves in discrete lattices*, Phys. Lett. A **372** (2008), 2247–2253. doi:10.1016/j.physleta.2007.11.029.
- [20] A. Stefanov, P. G. Kevrekidis, *Asymptotic behaviour of small solutions for the discrete nonlinear Schrödinger and Klein–Gordon equations*, Nonlinearity **18** (2005), 1841–1857. doi:10.1088/0951-7715/18/4/022.
- [21] I. Egorova, E. Kopylova, G. Teschl, *Dispersion estimates for one-dimensional discrete Schrödinger and wave equations*, J. Spectr. Theory **5** (2015), 663–696. doi:10.4171/JST/110.
- [22] N. M. Ryskin, D. I. Trubetskov, *Nonlinear Waves* [in Russian], Saratov Univ. Press, 2010.
- [23] OEIS Foundation Inc., *The On-Line Encyclopedia of Integer Sequences*, sequences A174090 and A055034, <https://oeis.org>.
- [24] NIST Digital Library of Mathematical Functions, F. W. J. Olver et al. (eds.), <https://dlmf.nist.gov/>.
- [25] R. M. Gray, *Toeplitz and Circulant Matrices: A Review*, Found. Trends Commun. Inf. Theory **2** (2006), 155–239.
- [26] J. B. Lasserre, *Global optimization with polynomials and the problem of moments*, SIAM J. Optim. **11** (2001), 796–817.
- [27] N. Levinson, *The Wiener RMS error criterion in filter design and prediction*, J. Math. Phys. **25** (1947), 261–278.
- [28] J. Durbin, *The fitting of time-series models*, Rev. Inst. Internat. Statist. **28** (1960), 233–244.
- [29] W. F. Trench, *An algorithm for the inversion of finite Toeplitz matrices*, J. SIAM **12** (1964), 515–522.
- [30] I. C. Gohberg, A. A. Semencul, *On the inversion of finite Toeplitz matrices and their continuous analogues*, Mat. Issled. **7** (1972), 201–223.
- [31] U. Grenander, G. Szegő, *Toeplitz Forms and Their Applications*, Univ. of California Press, 1958.
- [32] A. Böttcher, B. Silbermann, *Introduction to Large Truncated Toeplitz Matrices*, Springer, 1999.
- [33] G. H. Golub, C. F. Van Loan, *Matrix Computations*, 4th ed., Johns Hopkins Univ. Press, 2013.
- [34] C. Garoni, S. Serra-Capizzano, *Generalized Locally Toeplitz Sequences: Theory and Applications*, Vol. I, Springer, 2017.
- [35] I. V. Andrianov, J. Awrejcewicz, D. Weichert, *Improved continuous models for discrete media*, Math. Probl. Eng. **2010** (2010), art. 986242. doi:10.1155/2010/986242.
- [36] I. V. Andrianov, G. A. Starushenko, D. Weichert, *Numerical investigation of 1D continuum dynamical models of discrete chain*, ZAMM **92** (2012), no. 11–12, 945–954. doi:10.1002/zamm.201200057.
- [37] I. V. Andrianov, V. V. Danishevskyy, A. L. Kalamkarov, *Vibration localization in one-dimensional linear and nonlinear lattices: discrete and continuum models*, Nonlinear Dynam. **72** (2013), no. 1–2, 37–48. doi:10.1007/s11071-012-0688-4.
- [38] I. V. Andrianov, J. Awrejcewicz, V. V. Danishevskyy, *Linear and Nonlinear Waves in Microstructured Solids: Homogenization and Asymptotic Approaches*, CRC Press, Boca Raton, 2021. ISBN 978-0-367-70412-4. (Ch. 10 reproduces Filimonov’s splash table, Table 10.1.)
- [39] S. Longhi, *Kramers–Kronig potentials for the discrete Schrödinger equation*, Phys. Rev. A **96** (2017), 042106. doi:10.1103/PhysRevA.96.042106.
- [40] Ilia Gradina, *Solving the basic problems of linear algebra for Toeplitz matrices*, term paper, MIIT, 2012 (unpublished).
- [41] A. D. Myshkis, A. M. Filimonov, *About the string with beads*, Mat. Fiz. Anal. Geom. **10** (2003), no. 3, 301–306. (Segment $\ln$ -asymptotics.)
- [42] P. F. Kurchanov, A. D. Myshkis, A. M. Filimonov, *Vibrations of rolling stock and a theorem of Kronecker*, J. Appl. Math. Mech. **55** (1991), no. 6, 870–876 (Russian original: Prikl. Mat. Mekh. **55**(6), 989–995). doi:10.1016/0021-8928(91)90140-P.
- [43] A. M. Filimonov, P. F. Kurchanov, A. D. Myshkis, *Some unexpected results in the classical problem of vibration of the string with n beads when n is large*, C. R. Acad. Sci. Paris Sér. I **313** (1991), no. 13, 961–965. (Cited from the lineage list of [48]; not consulted directly.)
- [44] A. M. Filimonov, *Some unexpected results on the classical problem of vibrations of the string with N beads. The case of multiple frequencies*, C. R. Acad. Sci. Paris Sér. I **315** (1992), no. 8, 957–961. Zbl 0761.73053. (Circular string: peaks for N prime or 2^m , the multiple-frequency case, and a continuum model finer than the wave equation; abstract consulted only.)

- [45] A. M. Filimonov, A. D. Myshkis, *Asymptotic estimate of solution of one mixed difference–differential equation of oscillations theory*, J. Difference Equ. Appl. **4** (1998), no. 1, 13–16. doi: 10.1080/10236199808808125. (First ln-growth result; cited from [41].)
- [46] A. M. Filimonov, X. Mao, S. Maslov, *Splash effect and ergodic properties of solutions of the classical difference–differential equation*, J. Difference Equ. Appl. **6** (2000), no. 3, 319–328. doi: 10.1080/10236190008808231.
- [47] A. M. Filimonov, A. D. Myshkis, *On properties of the large wave effect in the classical problem of bead string vibration*, J. Difference Equ. Appl. **10** (2004), no. 13–15, 1171–1175. doi: 10.1080/10236190410001652757.
- [48] A. M. Filimonov, *Sufficient conditions for the large wave phenomenon in the Schrödinger equation with discrete spatial variable*, J. Math. Sci. **269** (2023), no. 6, 792–795. doi: 10.1007/s10958-023-06316-1.
- [49] N. E. Zhukovsky, *Collected Works* (the railway train start-up problem), vol. VIII, 1937 (classical; not DOI-indexed).
- [50] V. E. Zakharov, *Stability of periodic waves of finite amplitude on the surface of a deep fluid*, J. Appl. Mech. Tech. Phys. **9** (1968), 190–194. doi: 10.1007/BF00913182.
- [51] V. E. Zakharov, A. B. Shabat, *Exact theory of two-dimensional self-focusing and one-dimensional self-modulation of waves in nonlinear media*, Zh. Eksp. Teor. Fiz. **61** (1971), 118–134 [Sov. Phys. JETP **34** (1972), 62–69].
- [52] V. E. Zakharov, L. D. Faddeev, *Korteweg–de Vries equation: a completely integrable Hamiltonian system*, Funct. Anal. Appl. **5** (1971), 280–287. doi: 10.1007/BF01086739.
- [53] V. E. Zakharov, L. A. Ostrovsky, *Modulation instability: the beginning*, Phys. D **238** (2009), 540–548. doi: 10.1016/j.physd.2008.12.002.
- [54] V. E. Zakharov, E. A. Kuznetsov, *On three-dimensional solitons*, Zh. Eksp. Teor. Fiz. **66** (1974), 594–597 [Sov. Phys. JETP **39** (1974), 285–286].
- [55] V. E. Zakharov, V. S. L’vov, G. Falkovich, *Kolmogorov Spectra of Turbulence I: Wave Turbulence*, Springer, 1992.
- [56] Y. Deng, Z. Hani, *Full derivation of the wave kinetic equation*, Invent. Math. **233** (2023), 543–724. doi: 10.1007/s00222-023-01189-2.
- [57] M. V. Berry, C. Upstill, *Catastrophe optics: morphologies of caustics and their diffraction patterns*, Prog. Opt. **18** (1980), 257–346. doi: 10.1016/S0079-6638(08)70215-4.
- [58] B. S. White, B. Fornberg, *On the chance of freak waves at sea*, J. Fluid Mech. **355** (1998), 113–138. doi: 10.1017/S0022112097007751.
- [59] I. V. Lavrenov, *The wave energy concentration at the Agulhas current off South Africa*, Nat. Hazards **17** (1998), 117–127. doi: 10.1023/A:1007978326982.
- [60] L. H. Ying, Z. Zhuang, E. J. Heller, L. Kaplan, *Linear and nonlinear rogue wave statistics in the presence of random currents*, Nonlinearity **24** (2011), R67–R87. doi: 10.1088/0951-7715/24/11/R01.
- [61] V. E. Zakharov, R. V. Shamin, *Probability of the occurrence of freak waves*, JETP Lett. **91** (2010), 62–65. doi: 10.1134/S0021364010020025.
- [62] A. V. Slunyaev, D. E. Pelinovsky, E. N. Pelinovsky, *Rogue waves in the sea: observations, physics, and mathematics*, Phys. Usp. **66** (2023), 148–172. doi: 10.3367/UFNe.2021.08.039038.
- [63] L. Niemeyer, L. Pietronero, H. J. Wiesmann, *Fractal dimension of dielectric breakdown*, Phys. Rev. Lett. **52** (1984), 1033–1036. doi: 10.1103/PhysRevLett.52.1033.
- [64] T. B. Benjamin, J. E. Feir, *The disintegration of wave trains on deep water. Part 1. Theory*, J. Fluid Mech. **27** (1967), 417–430. doi: 10.1017/S002211206700045X.
- [65] D. H. Peregrine, *Water waves, nonlinear Schrödinger equations and their solutions*, J. Austral. Math. Soc. Ser. B **25** (1983), 16–43. doi: 10.1017/S0334270000003891.
- [66] N. N. Akhmediev, V. I. Korneev, *Modulation instability and periodic solutions of the nonlinear Schrödinger equation*, Theoret. Math. Phys. **69** (1986), 1089–1093. doi: 10.1007/BF01037866.
- [67] R. S. MacKay, S. Aubry, *Proof of existence of breathers for time-reversible or Hamiltonian networks of weakly coupled oscillators*, Nonlinearity **7** (1994), 1623–1643. doi: 10.1088/0951-7715/7/6/006.
- [68] M. Grillakis, J. Shatah, W. Strauss, *Stability theory of solitary waves in the presence of symmetry, I*, J. Funct. Anal. **74** (1987), 160–197. doi: 10.1016/0022-1236(87)90044-9.
- [69] N. G. Vakhitov, A. A. Kolokolov, *Stationary solutions of the wave equation in a medium with nonlinearity saturation*, Radiophys. Quantum Electron. **16** (1973), 783–789. doi: 10.1007/BF01031343.
- [70] M. I. Weinstein, *Excitation thresholds for nonlinear localized modes on lattices*, Nonlinearity **12** (1999), no. 3, 673–691. doi: 10.1088/0951-7715/12/3/314.
- [71] L. J. Landau, *Bessel functions: monotonicity and bounds*, J. London Math. Soc. (2) **61** (2000), no. 1, 197–215. doi: 10.1112/S0024610799008352.

- [72] C. Kharif, E. Pelinovsky, *Physical mechanisms of the rogue wave phenomenon*, Eur. J. Mech. B/Fluids **22** (2003), 603–634. doi:10.1016/j.euromechflu.2003.09.002.
- [73] D. R. Solli, C. Ropers, P. Koonath, B. Jalali, *Optical rogue waves*, Nature **450** (2007), 1054–1057. doi:10.1038/nature06402.
- [74] M. Onorato, S. Residori, U. Bortolozzo, A. Montina, F. T. Arecchi, *Rogue waves and their generating mechanisms in different physical contexts*, Phys. Rep. **528** (2013), 47–89. doi:10.1016/j.physrep.2013.03.001.